A GRADUATE LECTURE

Concept Bottleneck Models

From black-box predictions to interpretable, editable concepts

Machine Learning Concepts

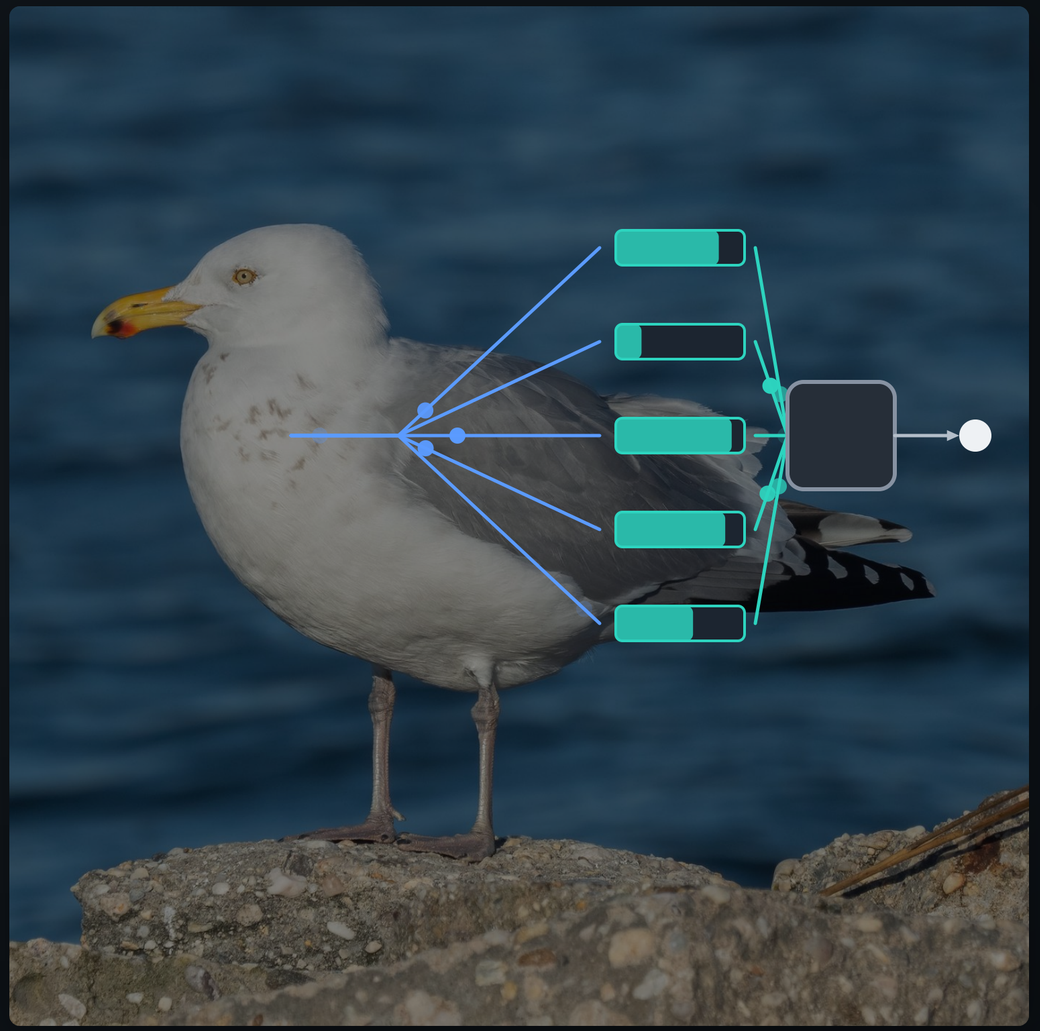

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

PART 1

Why make predictions interpretable?

What evidence did the model use?

PART 1

The black box

Can a person inspect, and correct, the evidence a network used?

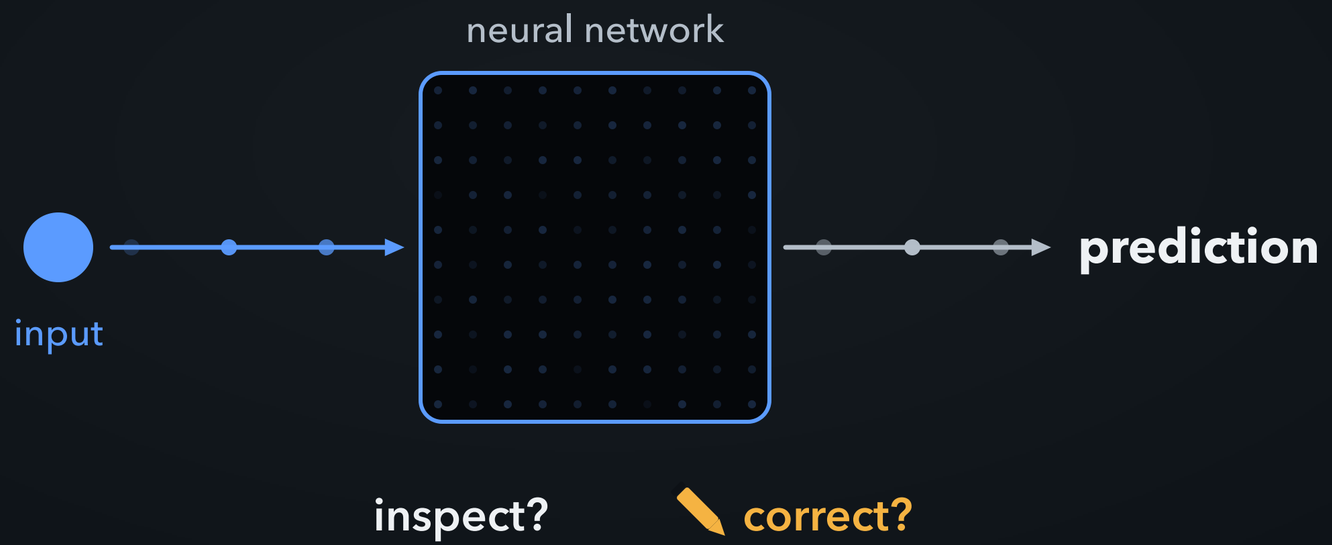

The black box

A real photograph and a real model, CLIP

Contrastive Language-Image Pre-training (CLIP)

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

a real, trained model

COMPUTED · REAL MODEL · CLIP ViT-B/32, zero-shot over 6 species names

The black box

CLIP compares the photo with one sentence per species

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

"a photo of a **{species}**, a type of bird."

a photo of a **American herring gull**, a type of bird.

a photo of a **glaucous-winged gull**, a type of bird.

a photo of a **ring-billed gull**, a type of bird.

a photo of a **great black-backed gull**, a type of bird.

a photo of a **laughing gull**, a type of bird.

a photo of a **glaucous gull**, a type of bird.

the image is compared with each sentence

PART 1

The black box

The photo goes in, a score for every species comes out

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

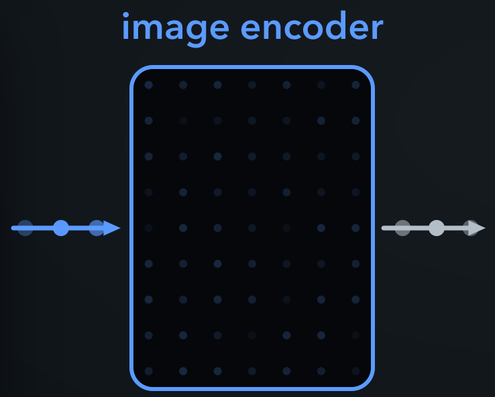

COMPUTED · REAL MODEL · CLIP ViT-B/32, softmax over 6 species names

PART 1

The black box

The answer is right. What made the model prefer it?

image encoder

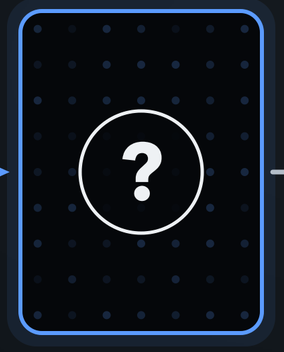

American herring gull	**0.54** ✓
Glaucous-winged gull	0.21
Great black-backed gull	0.11
Ring-billed gull	0.09
Glaucous gull	0.05
Laughing gull	0.00

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

COMPUTED · REAL MODEL · CLIP ViT-B/32, softmax over 6 species names

The black box

Occlusion: grey out one patch, measure the score drop

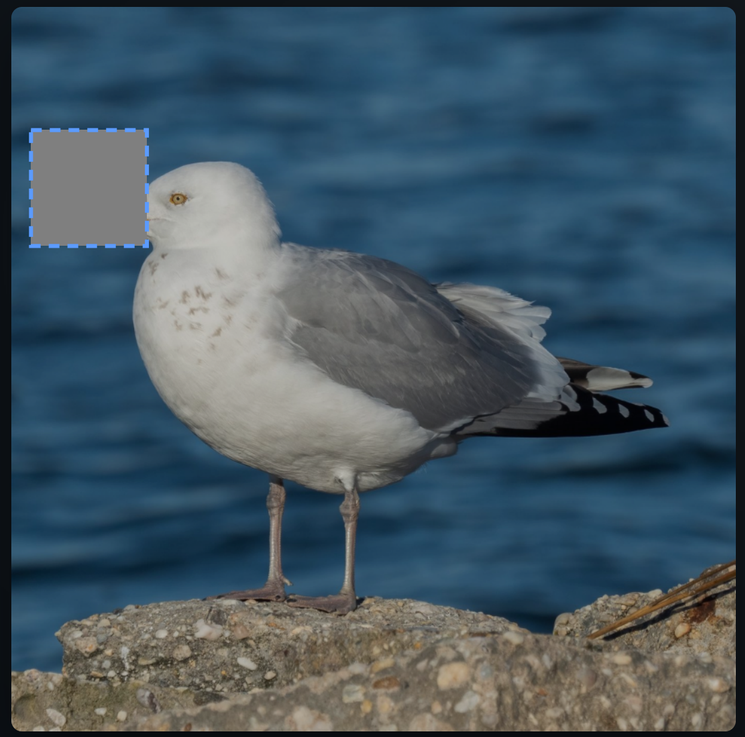

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

grey out one patch, run the model again

drop in the **herring gull** log-probability

0.92

COMPUTED · REAL MODEL · CLIP, one grey patch at each of 14 x 14 positions

The black box

The largest drops sit on the bill and the head

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

brighter: a larger drop in the **score**

`COMPUTED · REAL MODEL · occlusion sensitivity, 14 x 14 grey patches`

The black box

Which property did the model recognise there?

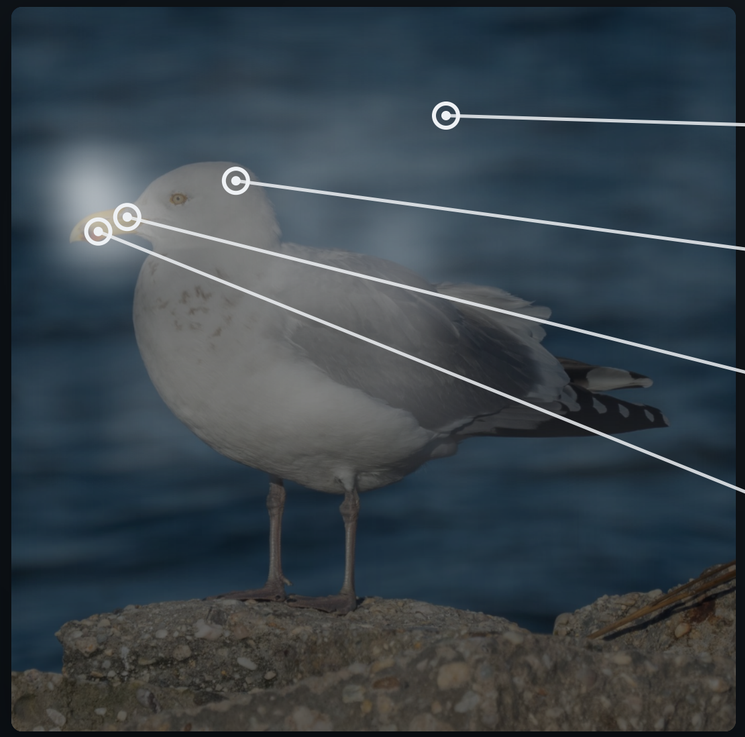

the water behind it?

the shape of the head?

the yellow bill?

the red spot near the tip?

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

The black box

The map shows where the score is sensitive, not which property

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

- ✓ where the **score** is sensitive
- ✗ which property the model detected
- ✗ what a person could correct

The black box

What we want: evidence stated as named properties

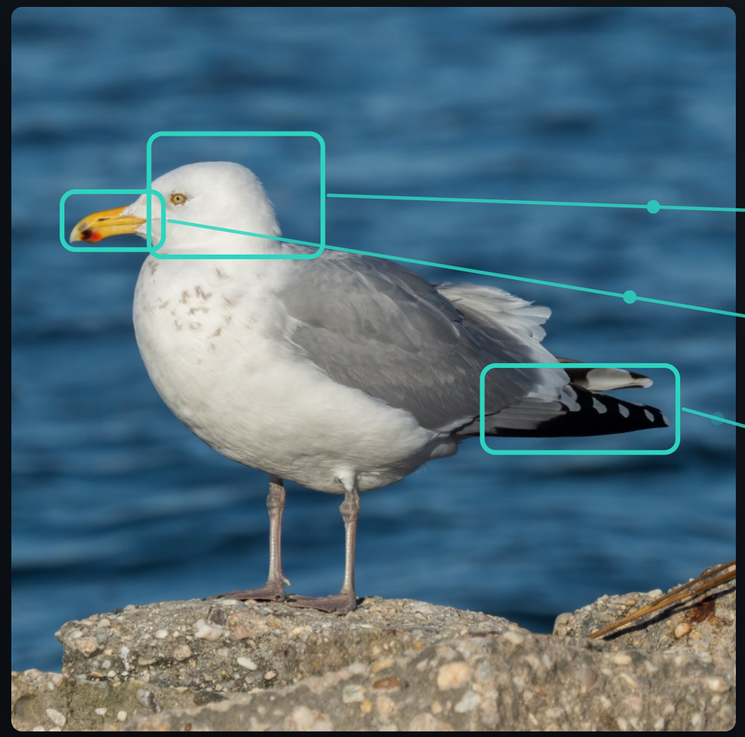

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

named properties

- head colour
- bill shape
- wingtip colour

each claim can be checked

Post-hoc or interpretable by design?

Two routes: explain a model afterwards, or build one you can read

post-hoc

explain it afterwards

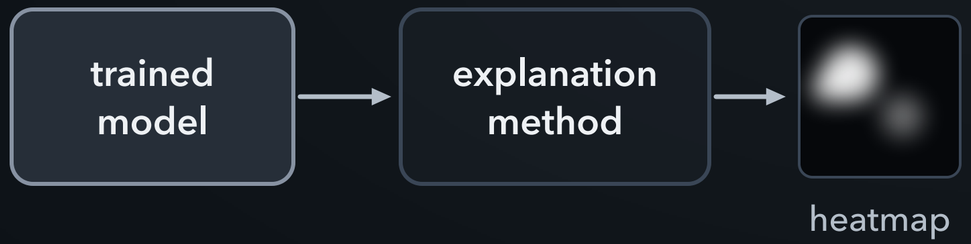

ante-hoc

interpretable by design

Post-hoc or interpretable by design?

Post-hoc: a method explains a trained model on one input

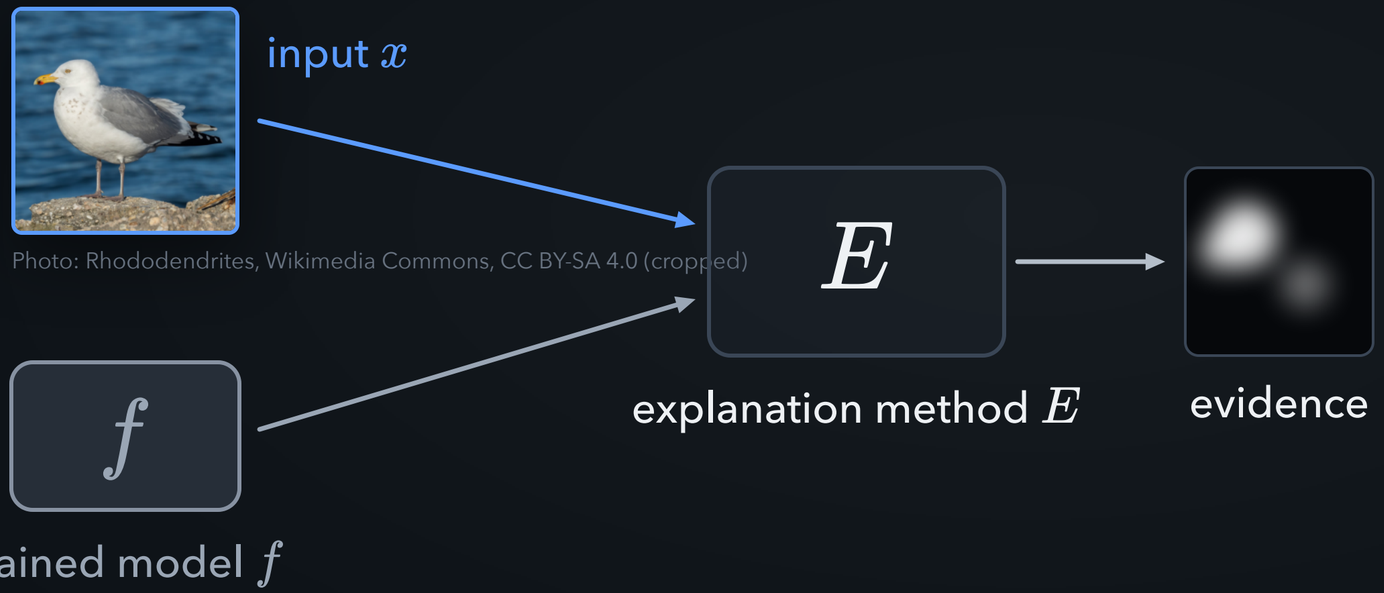

$$\text{evidence} = E(f, x)$$

Post-hoc or interpretable by design?

Saliency: nudge one pixel, watch one class score

input photo

Photo: Yathin S Krishnappa, Wikimedia Commons, CC BY-SA 3.0; centre crop

red value: +0.001

ResNet-50
`trained, weights fixed`

Residual Network with 50 layers

chosen score: **zebra**

before softmax

+0.00014

tiny change in the score

sensitivity = size of **score change** ÷ size of pixel change

$$S_i^{(k)} = \left| \frac{\partial s_k(x)}{\partial x_i} \right| \approx \left| \frac{\Delta s_k}{\Delta x_i} \right| \qquad \frac{0.00014}{0.001} \approx 0.14$$

`COMPUTED · REAL MODEL · real photo, fixed ImageNet model, one pixel channel changed`

Post-hoc or interpretable by design?

Same photo, same model, a different class to explain

same input photo

same ResNet-50, fixed weights

zebra score

probability: 37.7% (top class)

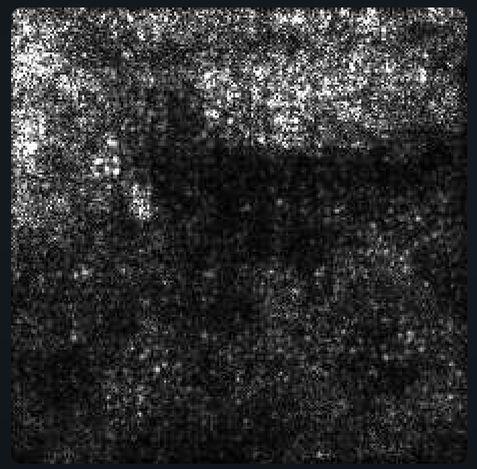

saliency for the zebra score

tiger score

probability: 0.08% (not predicted)

saliency for the tiger score

Same brightness scale: brighter = more sensitive

Low class probability does not mean zero pixel sensitivity

COMPUTED · REAL MODEL · raw score gradients; max over RGB; shared scale, brightest 1% clipped

Post-hoc or interpretable by design?

Occlusion: grey out a patch and measure how far the score drops

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

COMPUTED · REAL MODEL · CLIP occlusion map of the gull, Part 1

score on the photo minus **score** with patch p greyed

$$\Delta_p = \hat{y}_k(x) - \hat{y}_k(x \text{ with } p \text{ greyed})$$

Post-hoc or interpretable by design?

From an input photo and a class score to a heatmap

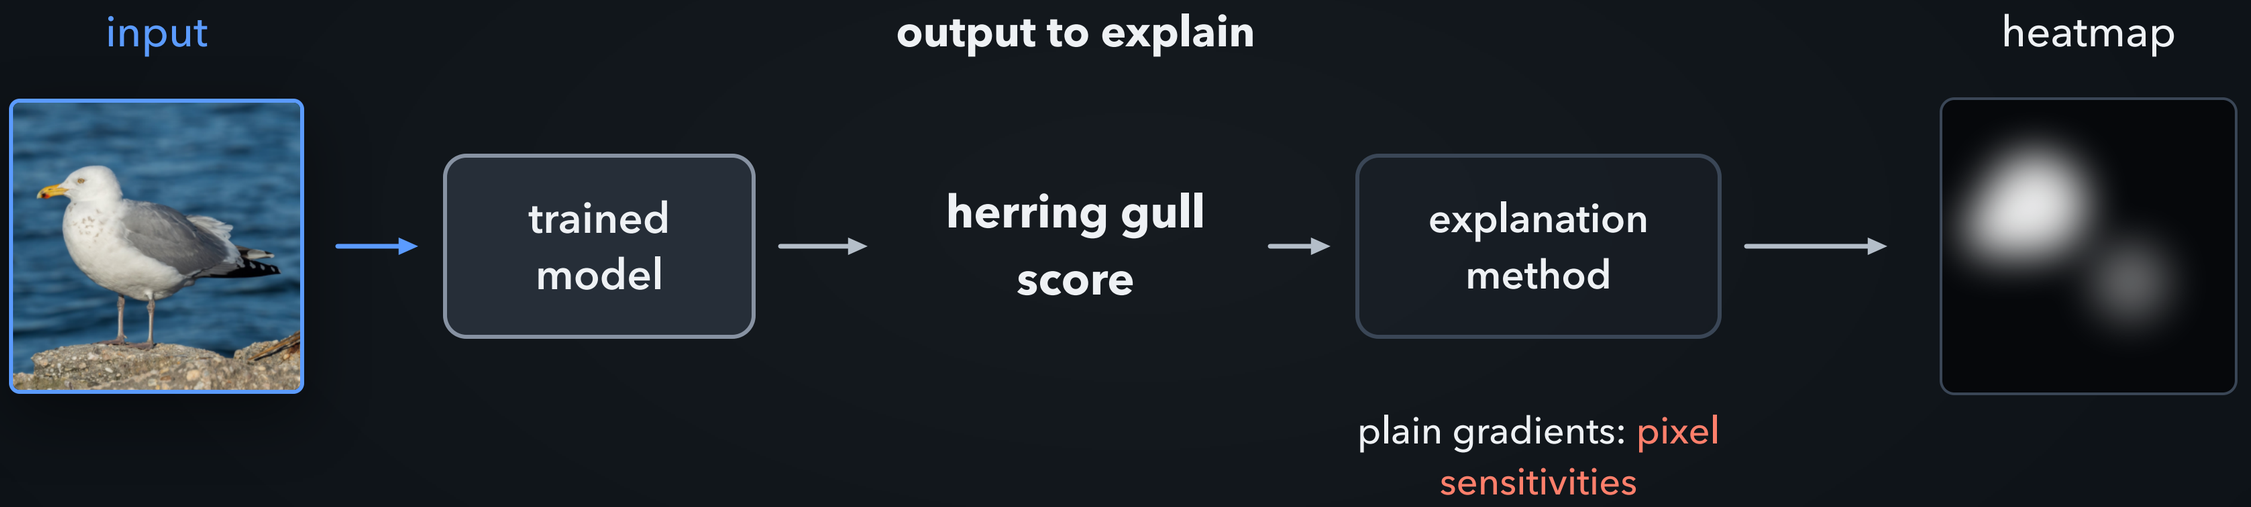

The method uses this photo, the model and the chosen **class**

Brighter pixels: the chosen score is more sensitive there

`SCHEMATIC · workflow and heatmap illustration, not computed for this photo`

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

Post-hoc or interpretable by design?

Does the heatmap depend on what the model learned?

Same photo, same **herring gull class**, same method

input photo

trained weights

replace learned weights

almost unchanged

random weights

reported example: guided backpropagation

Before trusting a heatmap, check that it depends on the learned weights

`SCHEMATIC · drawn maps illustrating a reported failure, Adebayo et al., NeurIPS 2018`

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

Post-hoc or interpretable by design?

The setup: learn pneumonia from two hospital datasets

TRAIN: use examples from both hospitals

Mount Sinai	34.2%	
X-rays + pneumonia labels	positive labels	
NIH	1.2%	
X-rays + pneumonia labels	positive labels	

NIH: National Institutes of Health Clinical Center

→ **learn a pneumonia predictor**

Labels come from radiology reports

positive = pneumonia reported

TEST: predict for a new image

one new X-ray → the trained predictor → **pneumonia score**

label hidden

Risk: the hospital is already a clue to the label

REPORTED · dataset label rates; Zech et al., PLOS Medicine 2018, Table 1

Post-hoc or interpretable by design?

The network could tell which hospital an X-ray came from

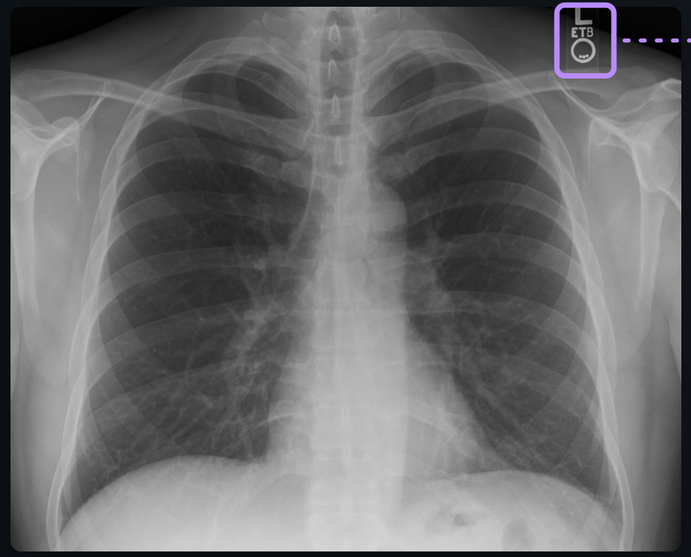

Chest X-ray: Stillwaterising, Wikimedia Commons, CC0

REPORTED · Zech et al., PLOS Medicine 2018

which hospital?

hospital named correctly for 99.95% of NIH X-rays

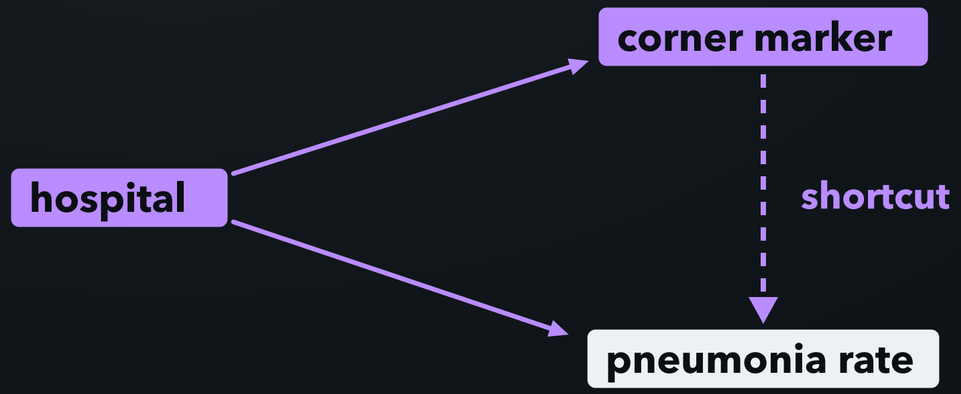

it can predict pneumonia without looking at the lungs

Post-hoc or interpretable by design?

An explanation can expose the shortcut, but the fix is new data

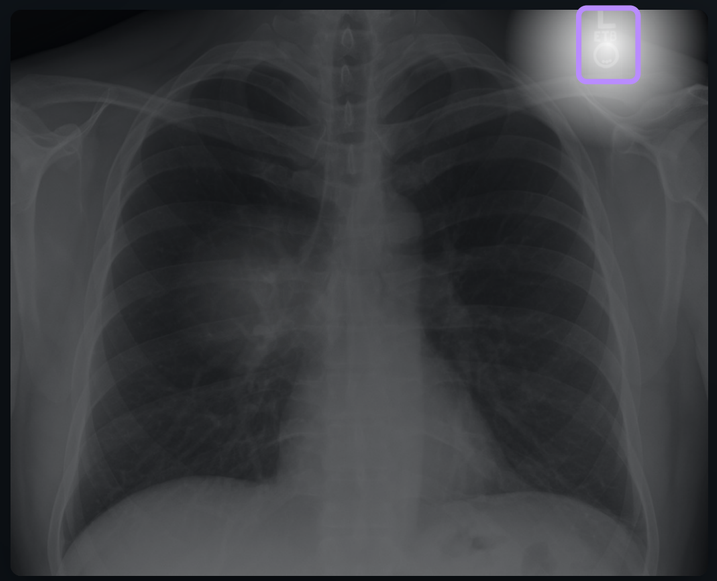

Chest X-ray: Stillwaterising, Wikimedia Commons, CC0

REPORTED · Zech et al., PLOS Medicine 2018

✓ exposes the shortcut

✗ cannot remove it

pneumonia AUC

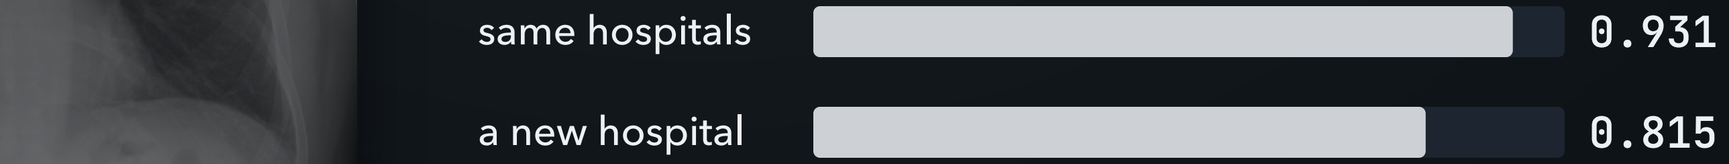

fix: **new data** and retraining

Post-hoc or interpretable by design?

Limit 3: a heatmap says where, not what

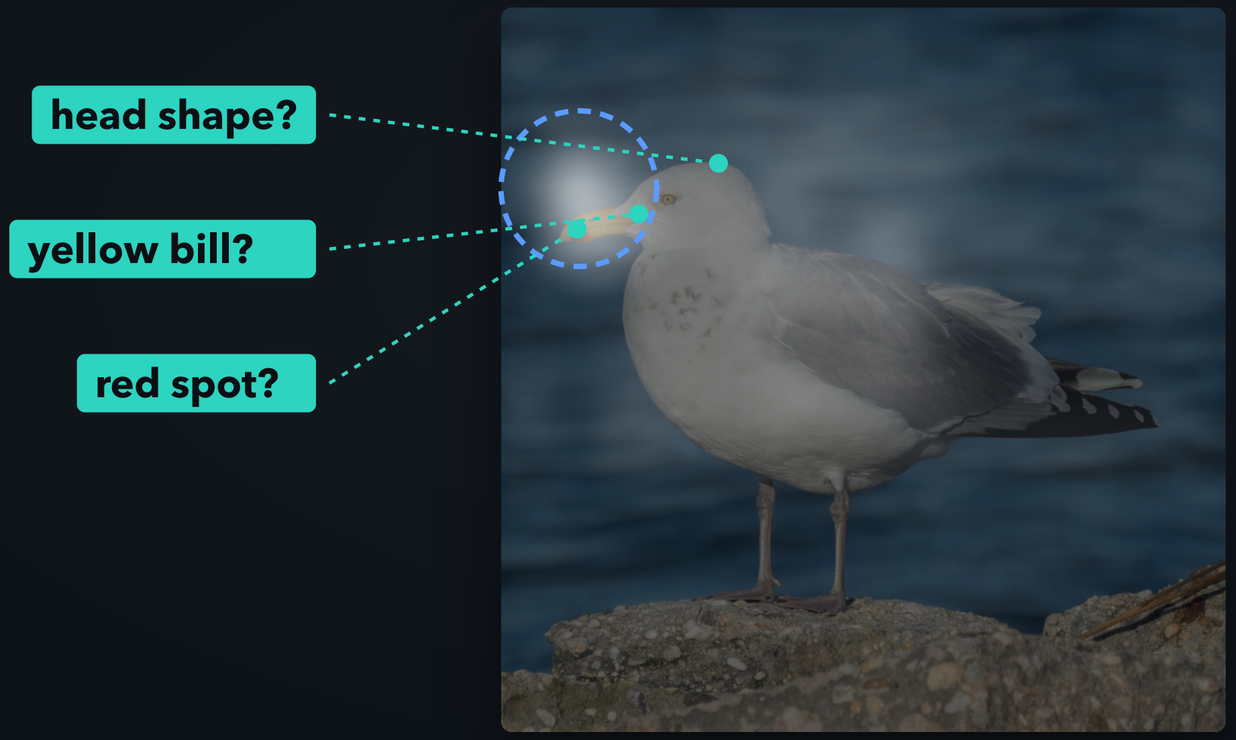

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

COMPUTED · REAL MODEL · CLIP occlusion map of the gull, Part 1

where: the score was sensitive here

what the model saw there: not said

Post-hoc or interpretable by design?

Experts speak in findings, and the model has no knob for them

findings

- narrowed joint space
- bone spur
- black wingtips

expert

black-box model

wingtip colour

no named part to correct

next: a concept bottleneck adds that knob (Part 2)

Post-hoc or interpretable by design?

Ante-hoc: the prediction is computed from quantities a person can read

sparse linear model

$$\hat{y} = 2.1\,\text{wingtip black} - 1.3\,\text{bill hooked} + \dots$$

short decision tree

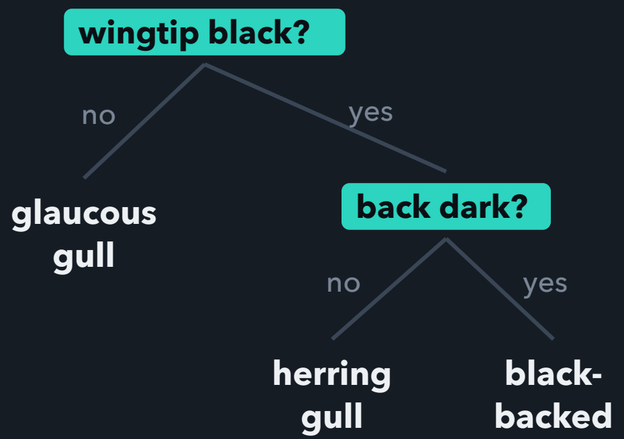

prototype network

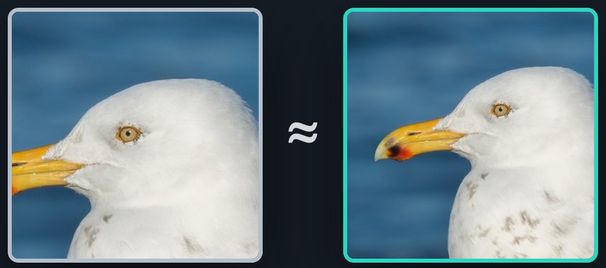

this part ≈ looks like that

Chen et al. 2019

ILLUSTRATIVE VALUES · hand-written examples

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

Post-hoc or interpretable by design?

The case for models that are interpretable by design

> For high-stakes decisions, prefer models people can read over black boxes explained afterwards.
>
> Rudin, Nature Machine Intelligence 2019 (paraphrase)

Post-hoc or interpretable by design?

A concept bottleneck predicts concepts first, then the answer

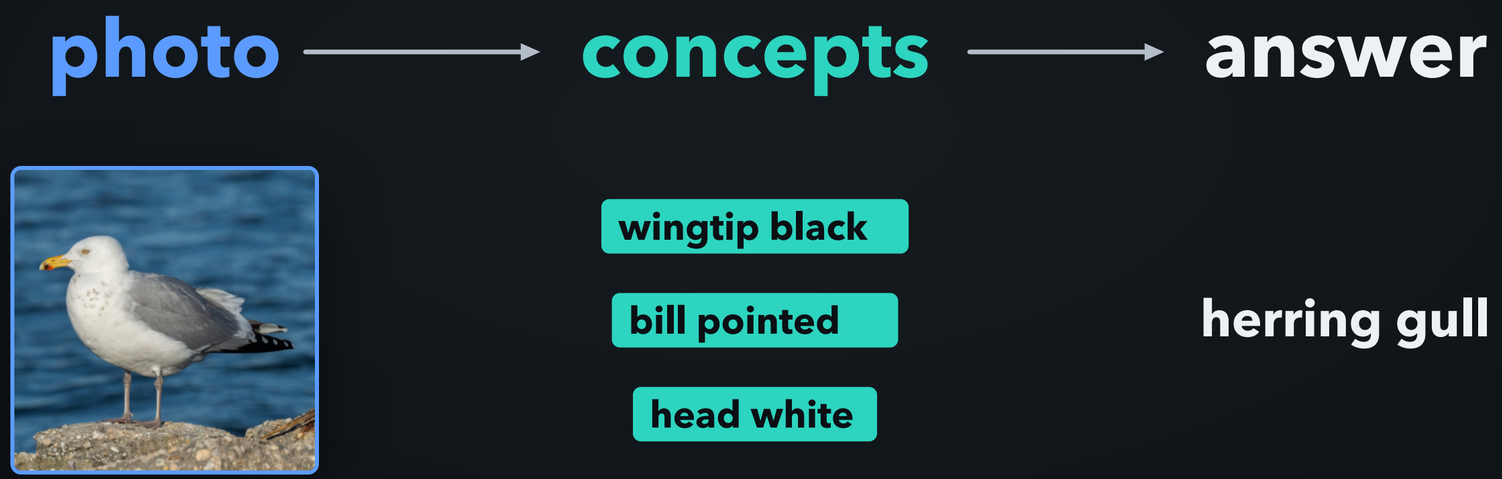

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

the **answer** is computed from the concepts alone

Post-hoc or interpretable by design?

The same path in symbols

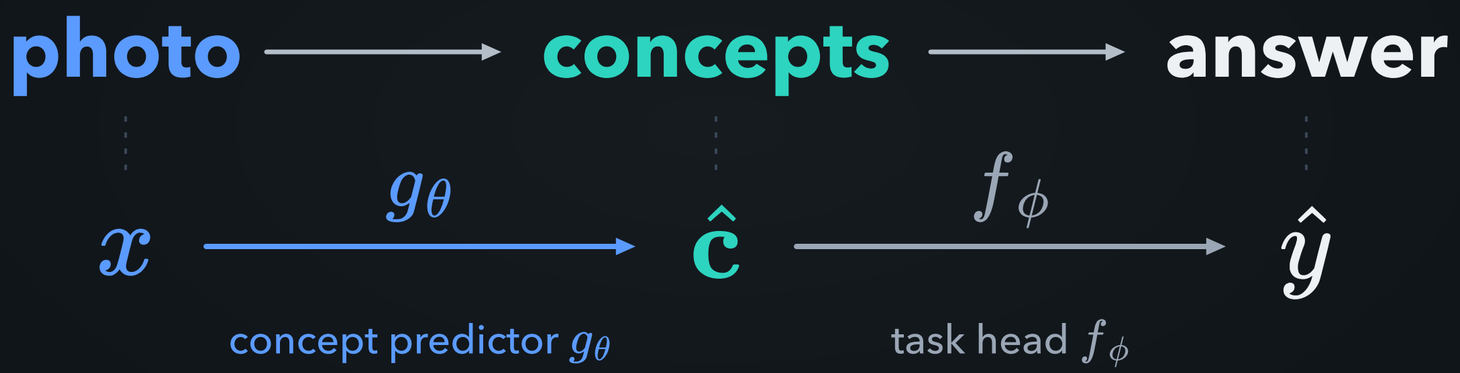

$\hat{\mathbf{c}}$: one predicted value per concept

PART 2

The concept bottleneck

Make the prediction pass through named concepts

Introducing the concept bottleneck

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Here is that idea on a real photograph

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

a layer of named variables

Introducing the concept bottleneck

A concept predictor fills the layer, a task head reads it

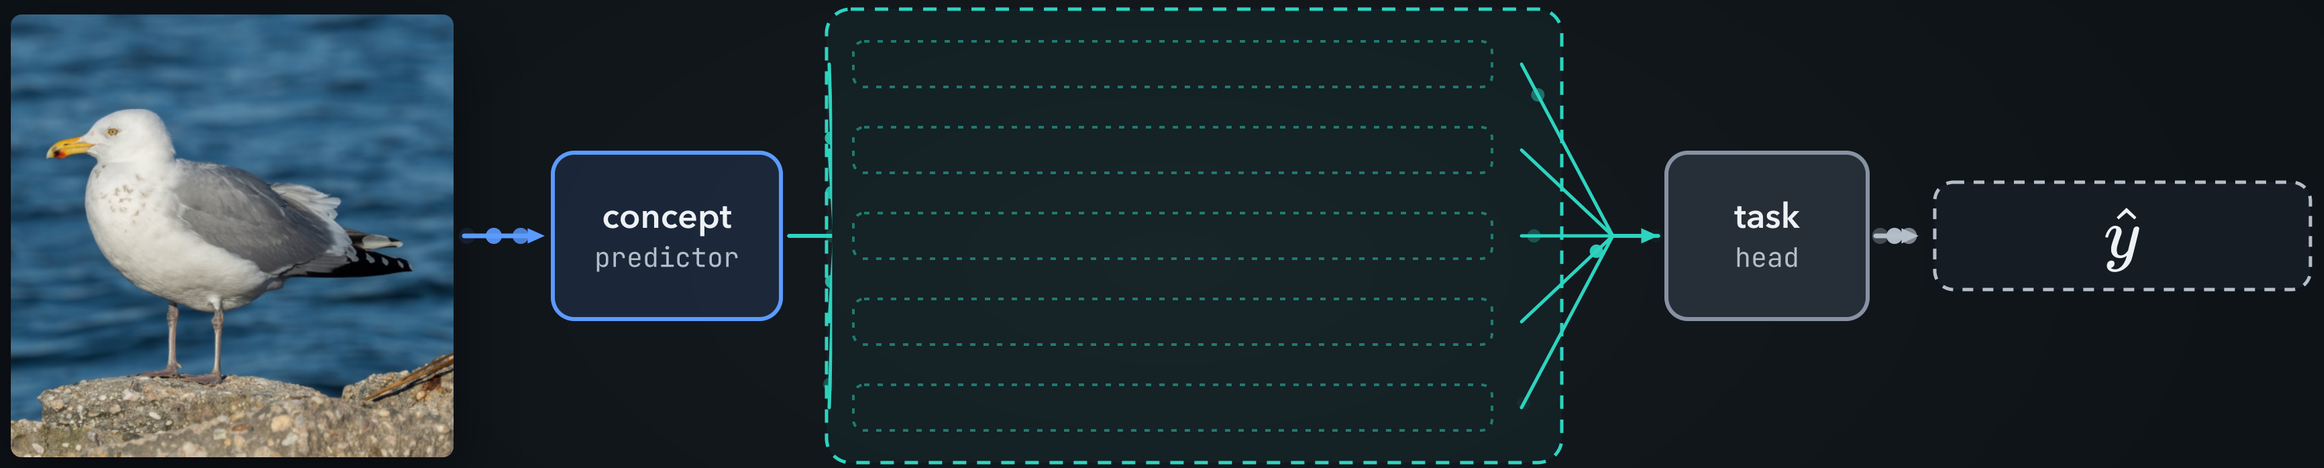

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

$$x \xrightarrow{g_\theta} \hat{\mathbf{c}} \xrightarrow{f_\phi} \hat{y}$$

Introducing the concept bottleneck

The concept vector has five named entries

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

g_θ

$$\hat{\mathbf{c}} = \begin{bmatrix} 0.80 \\ 0.20 \\ 0.90 \\ 0.85 \\ 0.60 \end{bmatrix}$$

- 0.80 — wingtip grey
- 0.20 — wingtip black
- 0.90 — bill pointed
- 0.85 — head white
- 0.60 — tail rounded

each entry: a value between 0 and 1

ILLUSTRATIVE VALUES · hand-written concept values, for teaching

Introducing the concept bottleneck

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

The rule: the task head reads only these five numbers

concept predictor

wingtip grey	0.80
wingtip black	0.20
bill pointed	0.90
head white	0.85
tail rounded	0.60

task head

$\hat{y}$

the bottleneck: the only input to the task head

no other path from the image

ILLUSTRATIVE VALUES · hand-written concept values and head, for teaching

Introducing the concept bottleneck

A strict CBM predicts the concepts first, then the answer from them alone

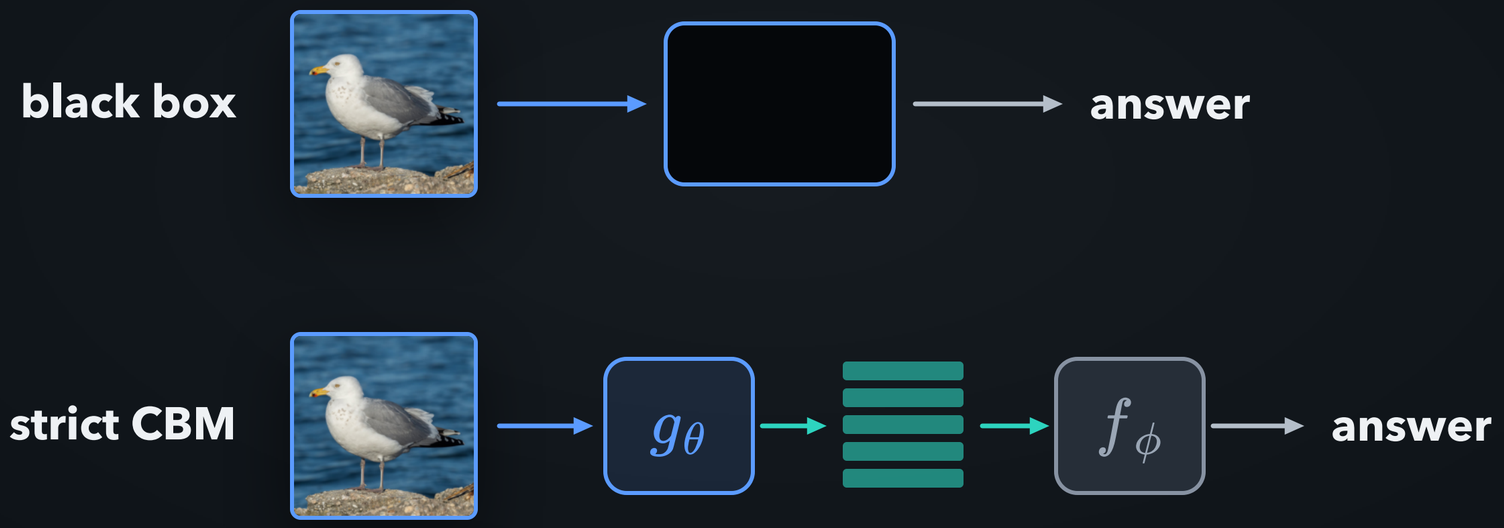

the **answer** from the concepts alone

Introducing the concept bottleneck

The values and the head on these slides are written by hand

	weight toward American herring gull
wingtip grey	-2.0
wingtip black	+2.0
bill pointed	+0.4
head white	-0.2
tail rounded	+0.2

the head adds up weight × value, then turns the total into two **shares**

ILLUSTRATIVE VALUES · hand-written concept values and head, for teaching

Introducing the concept bottleneck

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

The model believes the wingtips are grey, so it says glaucous-winged gull

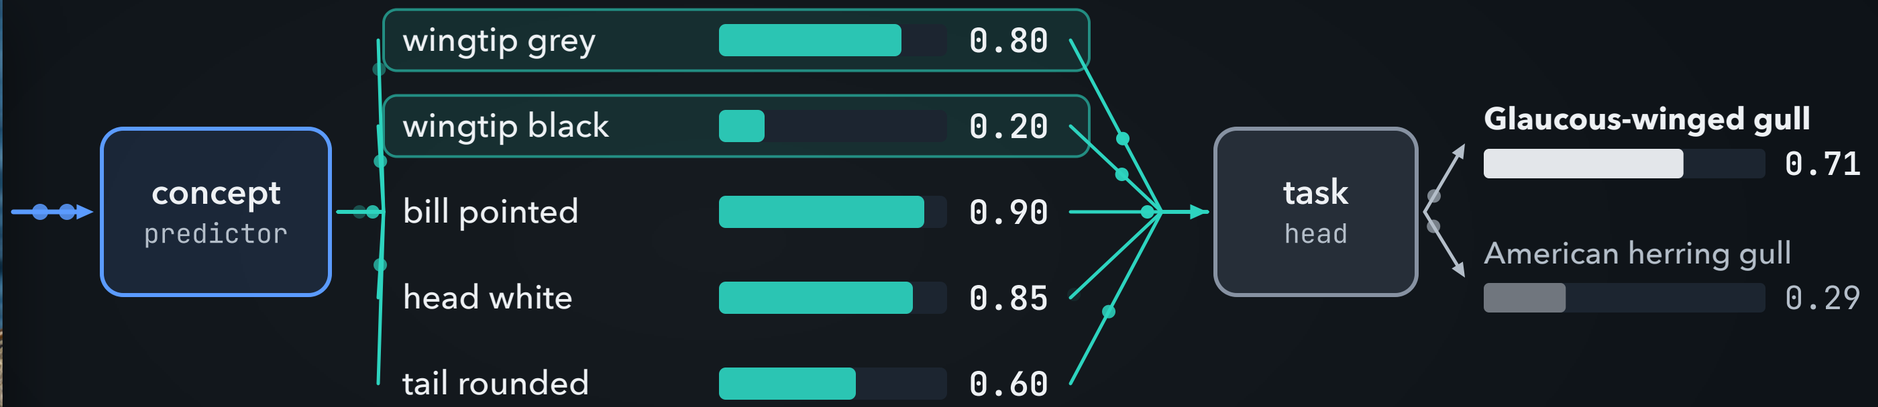

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

grey high, black low: the head favours **glaucous-winged gull**

ILLUSTRATIVE VALUES · hand-written concept values and head, for teaching

Introducing the concept bottleneck

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

The wingtips are black: an observer corrects the wing colour

the wingtips are black

concept predictor

wingtip grey	0.00
wingtip black	1.00
bill pointed	0.90
head white	0.85
tail rounded	0.60

task head

Glaucous-winged gull 0.71

American herring gull 0.29

wing colour: corrected as one group

ILLUSTRATIVE VALUES · hand-written concept values and head, for teaching

Introducing the concept bottleneck

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

The same task head recomputes, and herring gull now leads

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

unchanged

concept predictor

- wingtip grey 0.00
- wingtip black 1.00
- bill pointed 0.90
- head white 0.85
- tail rounded 0.60

task head

unchanged

- Glaucous-winged gull 0.09
- **American herring gull** 0.91

same image, same head, edited concepts: a new **answer**

ILLUSTRATIVE VALUES · hand-written concept values and head, for teaching

Introducing the concept bottleneck

An edited belief, a new prediction

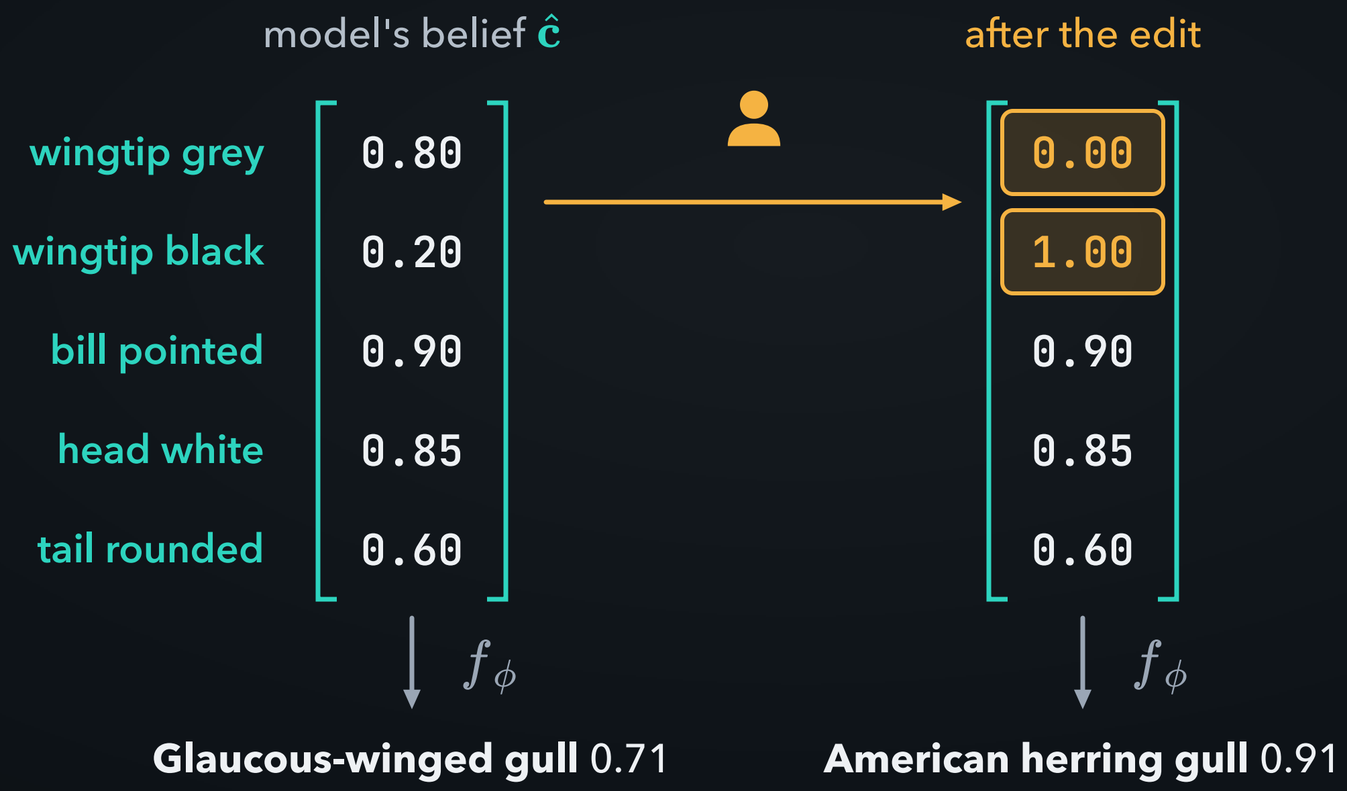

ILLUSTRATIVE VALUES · hand-written concept values and head, for teaching

Introducing the concept bottleneck

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Multi-task learning puts a concept head beside the class head

multi-task learning

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

shared encoder

concept head

- wingtip grey 0.80
- wingtip black 0.20
- bill pointed 0.90
- head white 0.85
- tail rounded 0.60

class head

- **Glaucous-winged gull**
- American herring gull

Introducing the concept bottleneck

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

There, an edit to the concepts changes nothing

multi-task learning

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

shared encoder

concept head

wingtip grey	0.00
wingtip black	1.00
bill pointed	0.90
head white	0.85
tail rounded	0.60

class head

Glaucous-winged gull

American herring gull

no change

the class head never reads the concept head

Introducing the concept bottleneck

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

In a CBM, the concepts lie on the path to the prediction

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

every **prediction** passes through the concepts

ILLUSTRATIVE VALUES · hand-written concept values and head, for teaching

The CUB benchmark

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

The standard benchmark for concept bottlenecks is a bird dataset

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

our photo, not from CUB

Caltech-UCSD Birds-200-2011

CUB

Wah et al., Caltech 2011

The CUB benchmark

Many photographs of many similar species

11,788

photographs

200

species

8 of the species are gulls

California · Glaucous-winged · Heermann · Herring · Ivory · Ring-billed · Slaty-backed · Western

DATASET · CUB-200-2011, counted

The CUB benchmark

Every photograph is described by the same attributes

312 yes-or-no attributes, in 28 groups

back colour	back pattern	belly colour	belly pattern
bill colour	bill length	**bill shape**	breast colour
breast pattern	crown colour	**eye colour**	forehead colour
head pattern	leg colour	nape colour	primary colour
shape	size	tail pattern	tail shape
throat colour	under tail colour	underparts colour	upper tail colour
upperparts colour	**wing colour**	**wing pattern**	wing shape

The CUB benchmark

Each attribute is a yes-or-no question

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

Does it have a hooked bill?	yes →	1
Is its wing grey?	yes →	1
Is its wing yellow?	no →	0

the answer to question j: yes is 1, no is 0

$$c_j \in \{0, 1\}$$

The CUB benchmark

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Annotators also recorded how sure they were

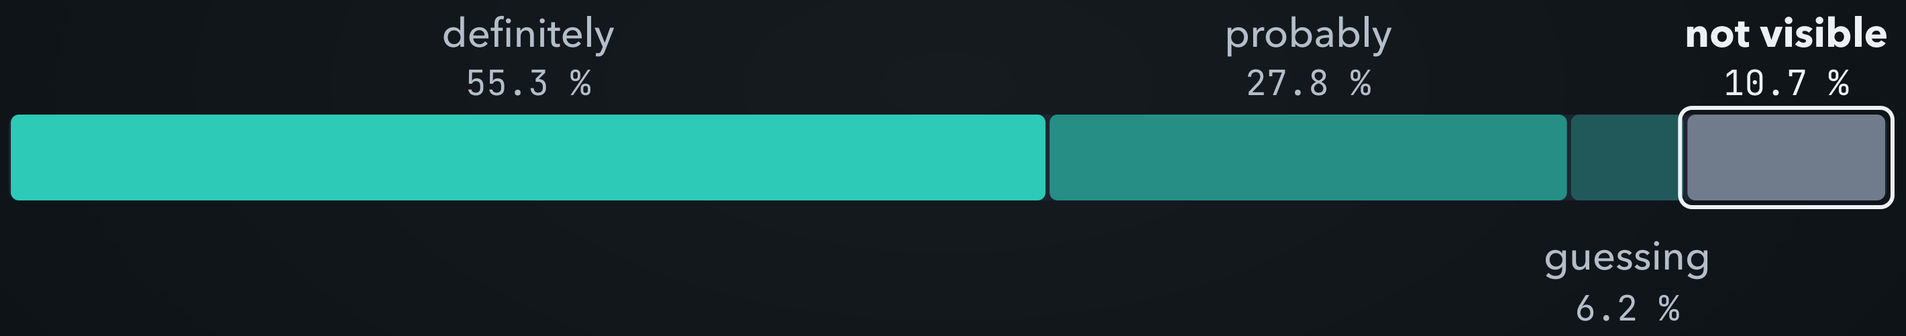

about one label in ten: not visible

DATASET · CUB annotations, counted

The CUB benchmark

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Single photos are noisy: take a vote within each species

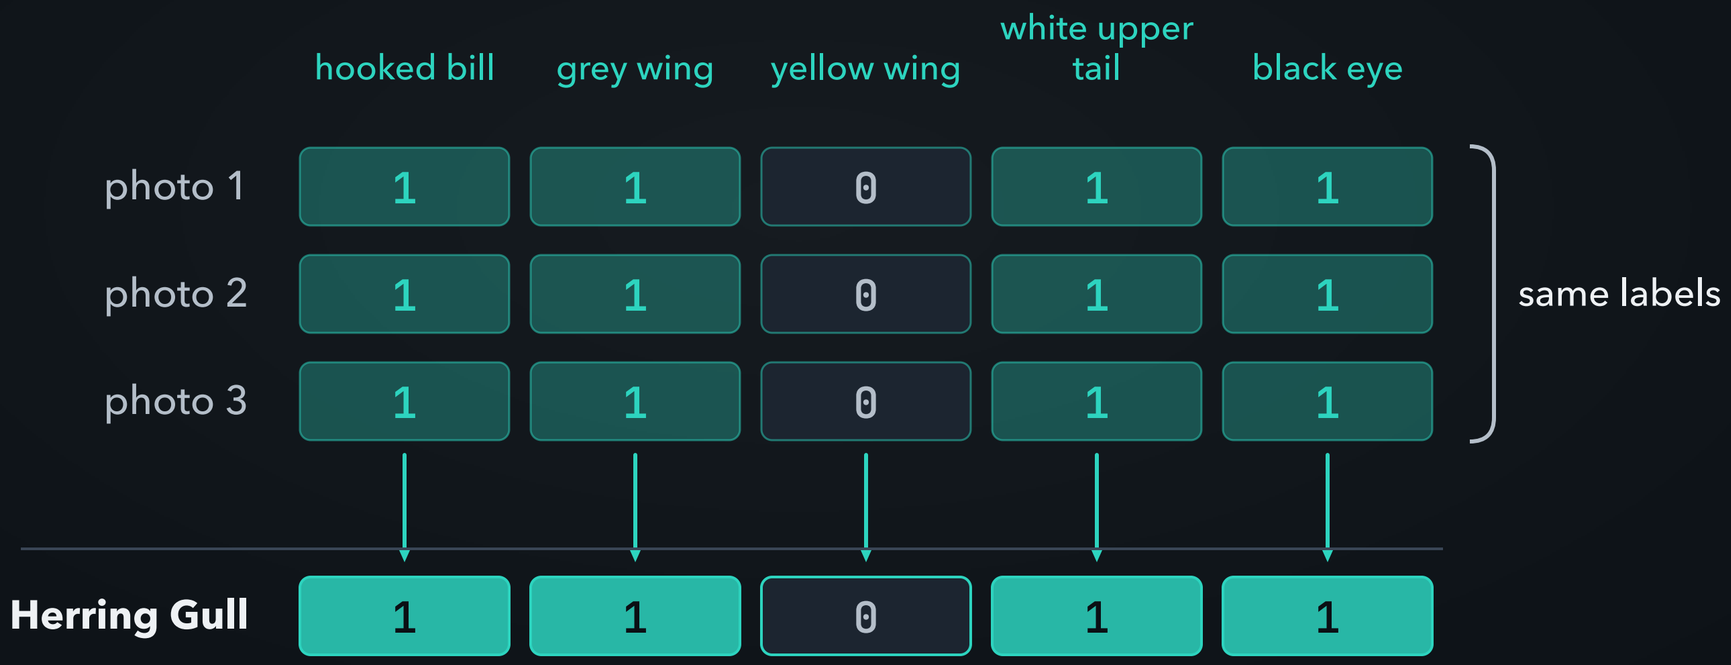

Koh et al., 2020

ILLUSTRATIVE VALUES · photo rows; the voted row is Herring Gull's in CUB

The CUB benchmark

Keep only the attributes that many species share

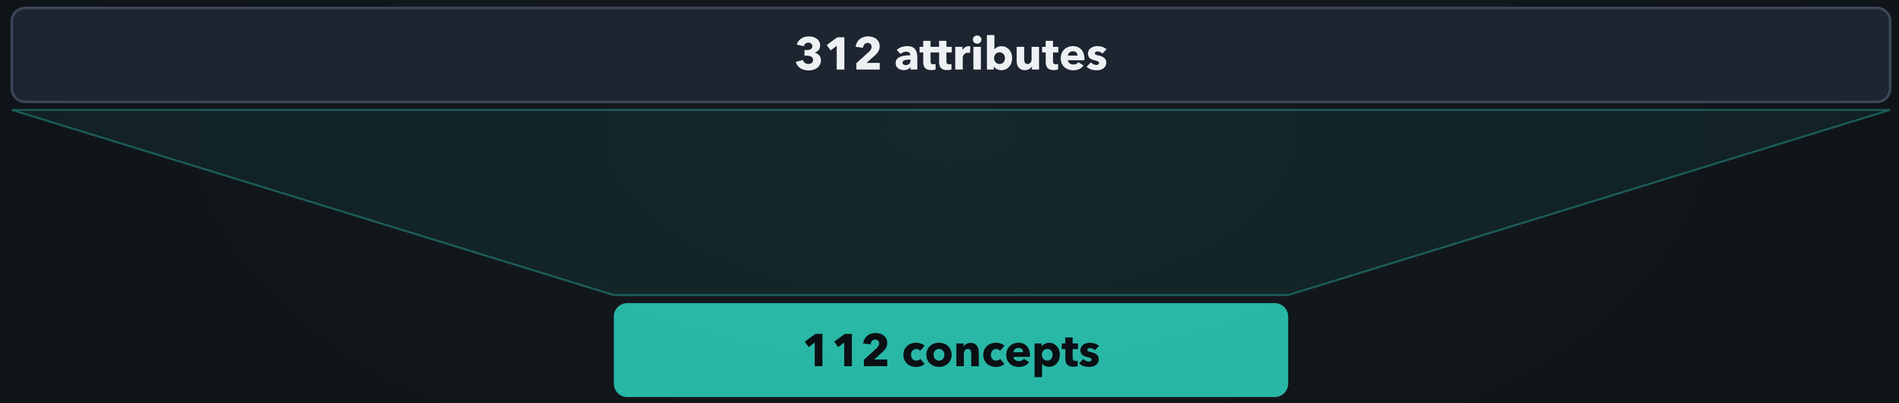

present in at least 10 species

DATASET · CUB, filtered as in Koh et al. 2020

The CUB benchmark

One training record has three parts

photo

Photo: Rhododendrites, CC BY-SA 4.0

112 concept labels, 0 or 1

0 1 0 0 0 1 0 0 1 0 0 1 0 0 1 0 0 0 0 0 …

has_bill_shape::hooked_seabird = 1

has_wing_color::grey = 1

has_upperparts_color::grey = 1

species

Herring Gull

one of 200

DATASET · Herring Gull's labels in CUB, first 20 of 112; our photo

The CUB benchmark

The same record in symbols

concept labels, each 0 or 1

$$\mathbf{c} \in \{0,1\}^{112}$$

species label

$$y \in \{1, \ldots, 200\}$$

predicted concepts — concept predictor — one value per concept

$$\hat{\mathbf{c}} = g_\theta(x) \in [0,1]^{112}$$

PART 3

What is a concept?

Choose what the model should describe

What is a concept?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

A concept is a claim that a person can check

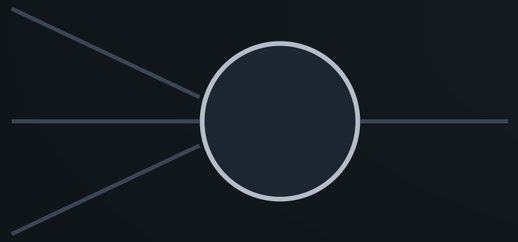

a name tag on a neuron

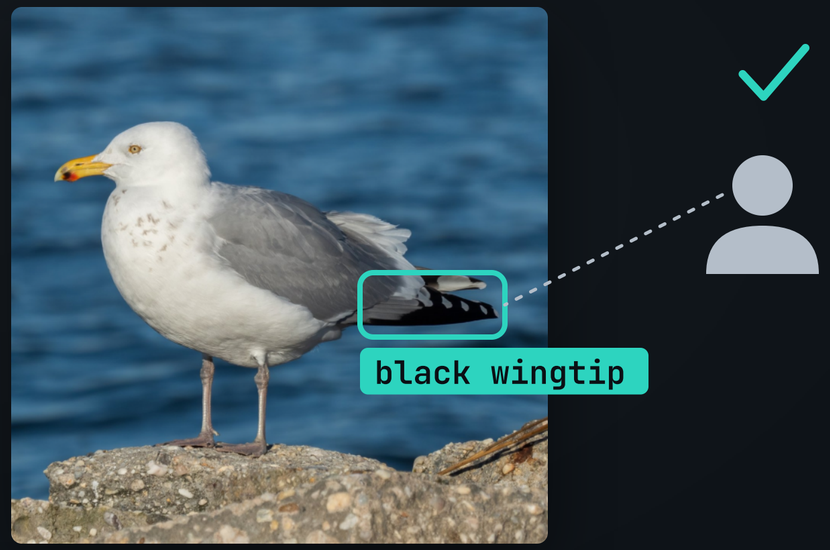

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

a claim a person checks in the input

What is a concept?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Meaning: a definition that an expert can apply

1 **meaning** 2 grounding 3 representation

expert

`grey wingtip`

the tips of the folded flight feathers are **grey**, not black

What is a concept?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Grounding: the evidence sits in one region of the photo

1 meaning 2 **grounding** 3 representation

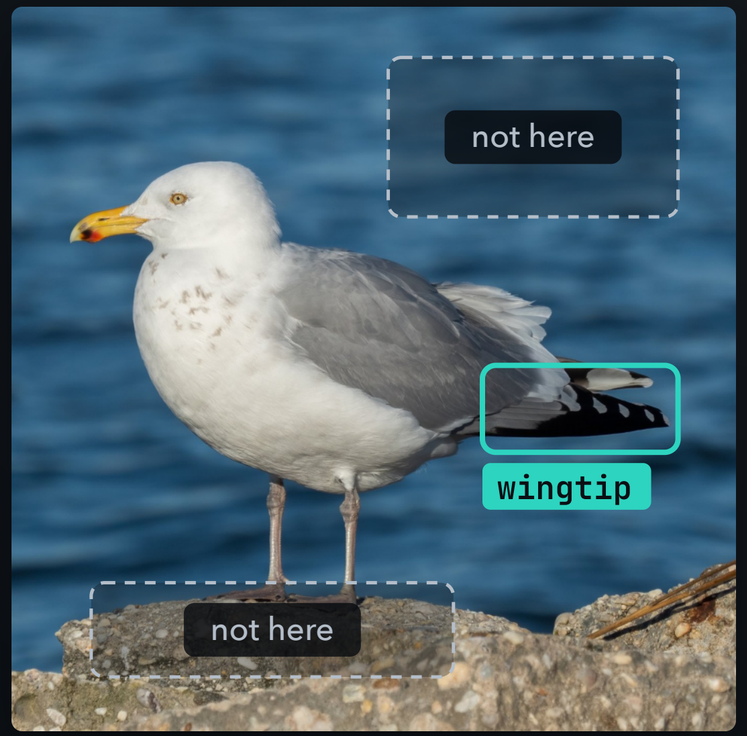

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

the wingtip colour estimate should depend on these feathers

not on the water or the rock

What is a concept?

Representation: the simplest concept is binary

1 meaning 2 grounding 3 **representation**

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

the value of concept j: 0 or 1

$$c_j \in \{0, 1\}$$

What is a concept?

A soft concept turns a logit into a value between 0 and 1

representation ● ● ○ ○ ○ ○

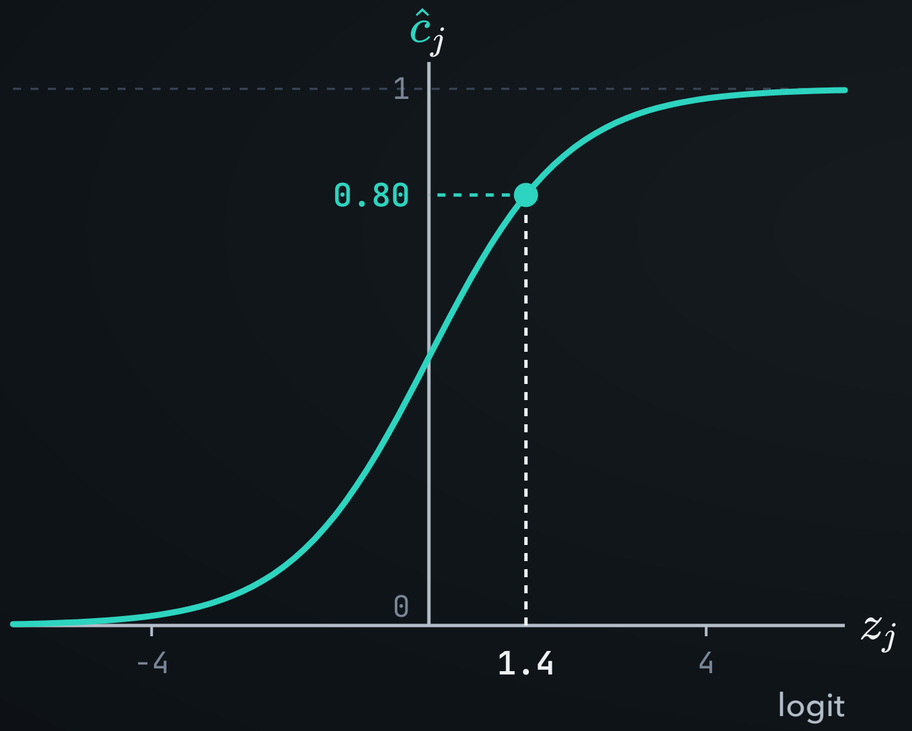

the soft score is the sigmoid of a logit

$$\hat{c}_j = \sigma(z_j) = \frac{1}{1 + e^{-z_j}}$$

example: $z_j = 1.4$ gives $\hat{c}_j = 0.80$

ILLUSTRATIVE VALUES · a worked logit

What is a concept?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

A soft score is not automatically a probability

representation ●●○○○○

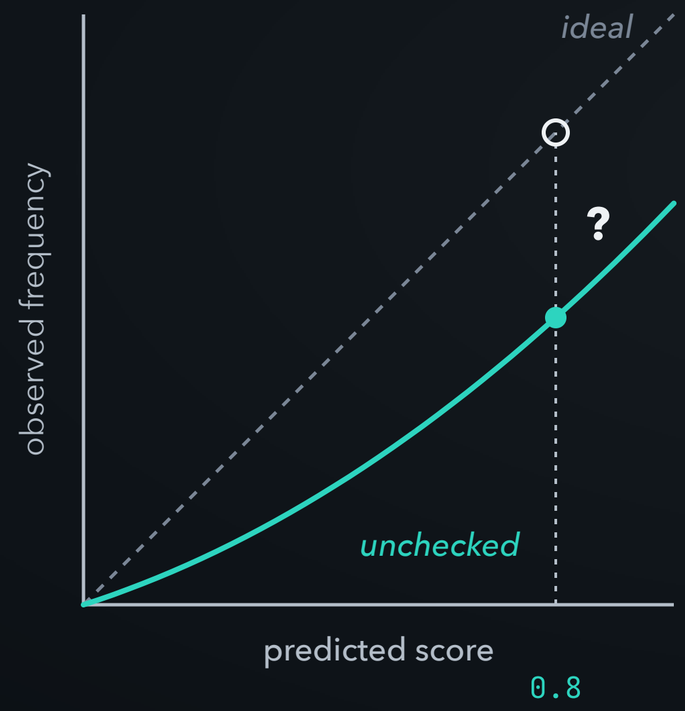

0.8 is a score, not automatically a probability

check on data: do 80 % of the cases scored 0.8 have grey wingtips?

SCHEMATIC · a sketch, no data

What is a concept?

Categorical: one value from a list, with scores that sum to 1

representation ●●●○○○

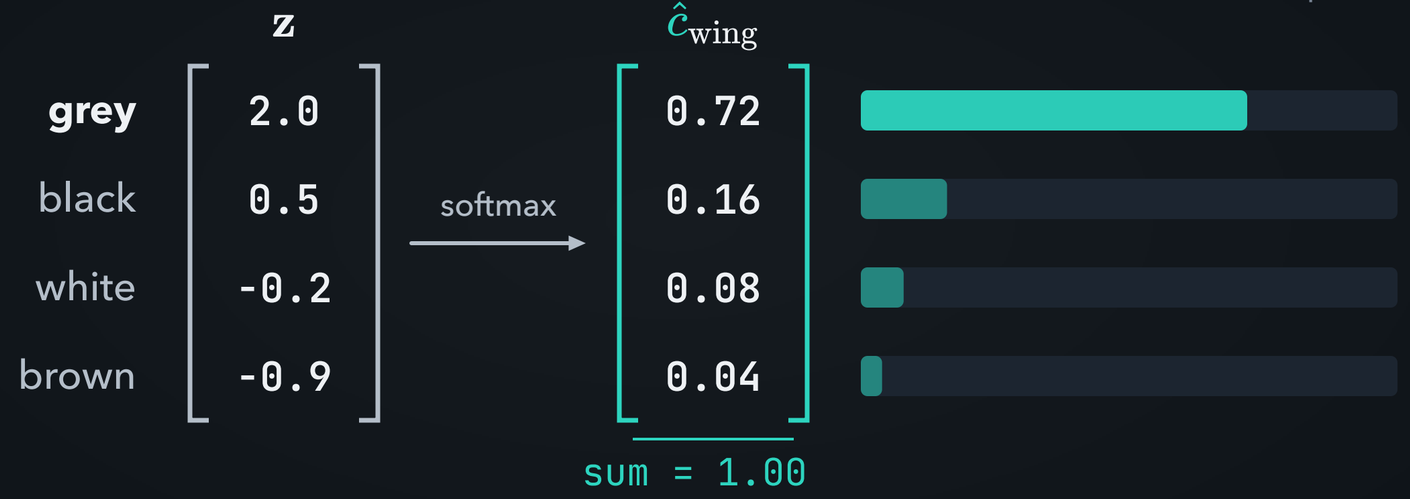

the wing colour scores are a softmax of the logits, and they sum to 1

$$\hat{c}_{\text{wing}} = \text{softmax}(\mathbf{z})$$

ILLUSTRATIVE VALUES · logits made up for the example

What is a concept?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Continuous: a measurement, such as bill length

representation ●●●●○○

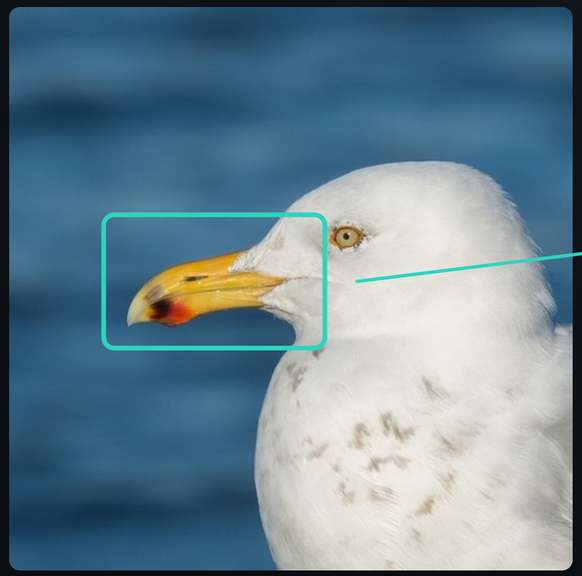

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

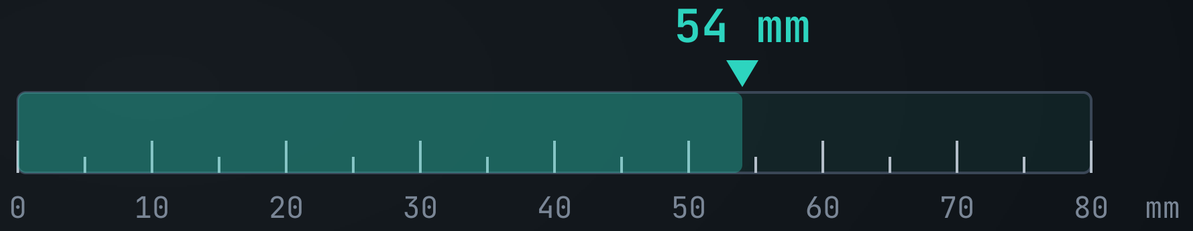

the bill length is a real number, in mm

$$\hat{c}_{\text{bill}} \in \mathbb{R}$$

predicted by regression

ILLUSTRATIVE VALUES · value

What is a concept?

Richer concepts: a vector or a whole distribution

representation ●●●●●○

a vector (embedding)

$$\begin{bmatrix} 0.31 \\ -1.20 \\ 0.85 \\ 0.02 \\ -0.47 \end{bmatrix}$$

$$\hat{c}_j \in \mathbb{R}^m$$

a distribution

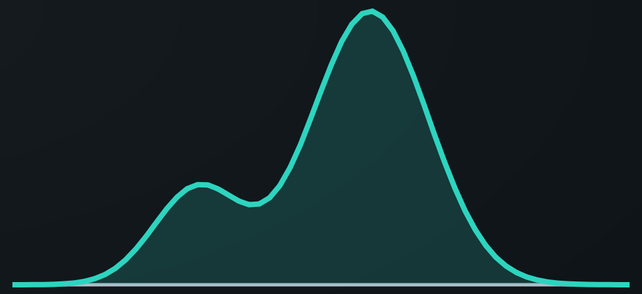

possible values of the concept

$$p(c_j \mid x)$$

or

we meet these in Part 10

ILLUSTRATIVE VALUES · values

What is a concept?

Unknown: not visible is different from absent

CONSTRUCTED EXAMPLE · wingtip covered for the example

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

representation

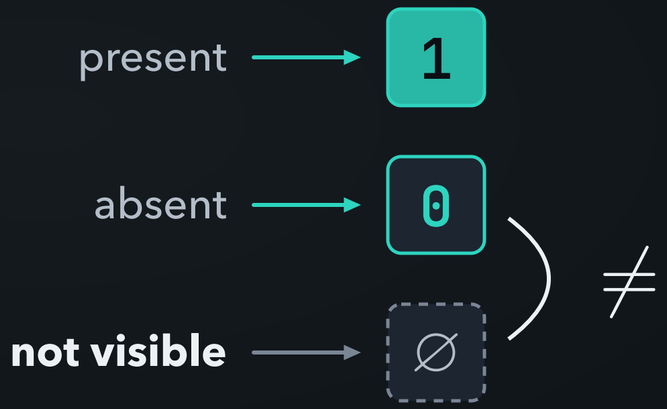

CUB label not visible: unknown, not absent

$$c_j = \varnothing \neq 0$$

What is a concept?

Six ways to store a concept as a number

type	in symbols	example
binary	$c_j \in \{0,\ 1\}$	grey wingtip: 1
soft	$\hat{c}_j = \sigma(z_j)$	logit 1.4 gives 0.80
categorical	$\hat{c}_{\text{wing}} = \text{softmax}(\mathbf{z})$	grey 0.72, scores sum to 1
continuous	$\hat{c}_{\text{bill}} \in \mathbb{R}$	bill length 54 mm
vector or distribution	$\hat{\boldsymbol{c}}_j \in \mathbb{R}^m,\ \ p(c_j \mid \boldsymbol{x})$	more in Part 10
unknown	$c_j = \varnothing$	not visible, not absent

What is a concept?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

The same pattern in a sleep EEG recording

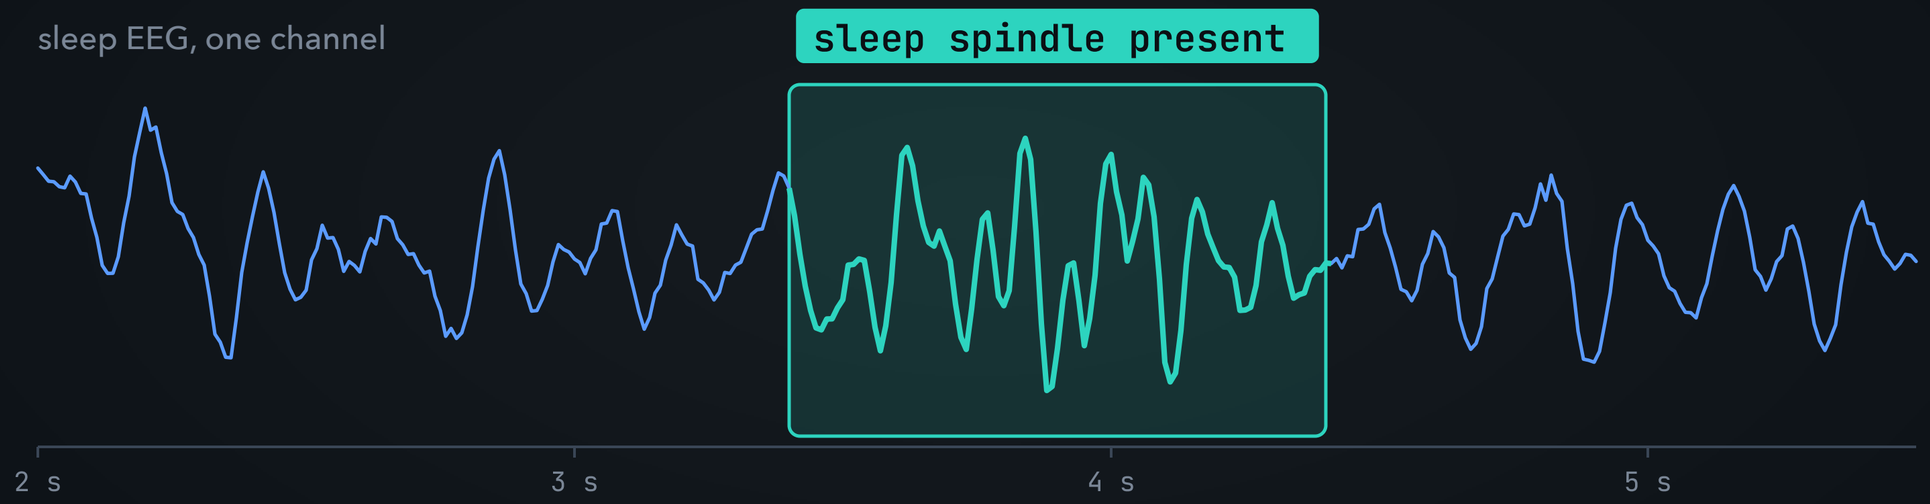

EEG trace → spindle present → sleep stage

SCHEMATIC · seeded sinusoids, not a recording

What is a concept?

A spindle is a brief burst of brain waves

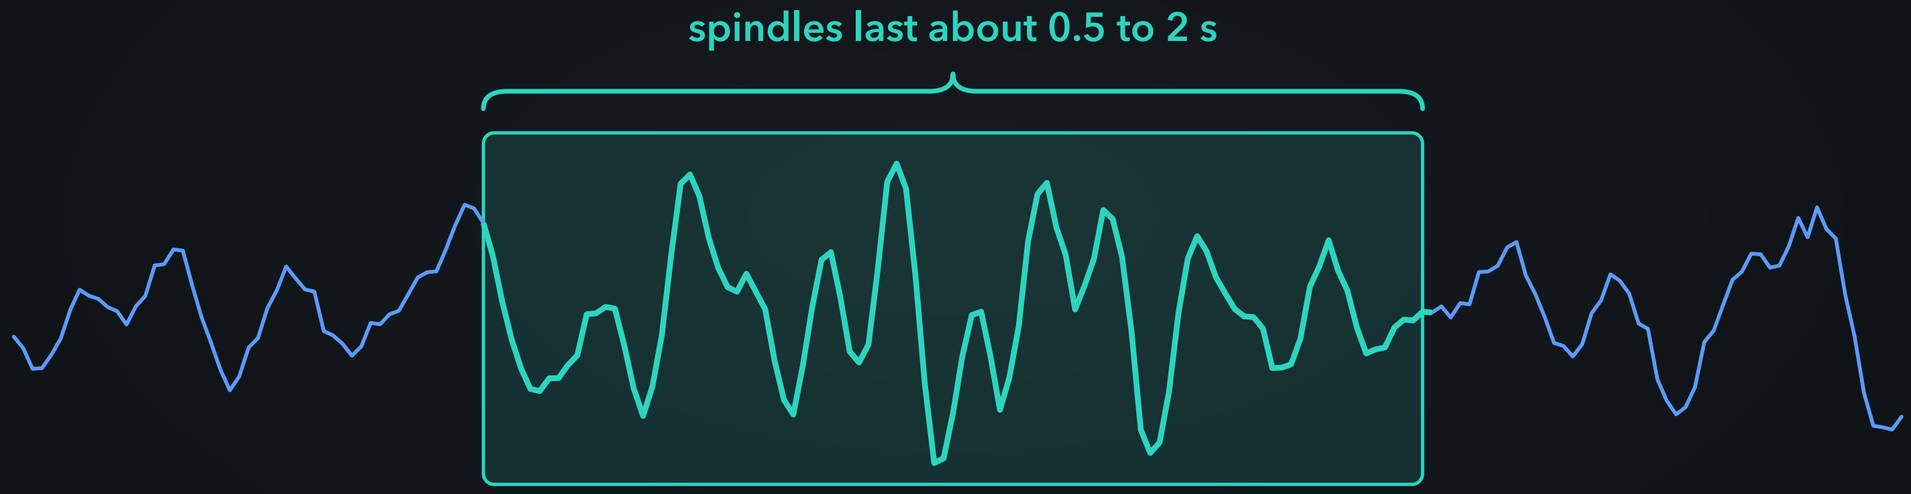

a burst of 11 to 16 cycles per second (Hz)

sleep scorers: a sign of **stage 2 sleep**

SCHEMATIC · this spindle: 13 Hz, 1 s; not a recording

What is a concept?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

The same pattern in a customer message

CUSTOMER MESSAGE

The app works, but I was charged twice. Please refund the extra payment.

↓

`duplicate charge reported`

the customer says one purchase was billed twice

message → duplicate charge → support queue

ILLUSTRATIVE VALUES · an example message

What is a concept?

Each concept is a named claim that a specialist can check

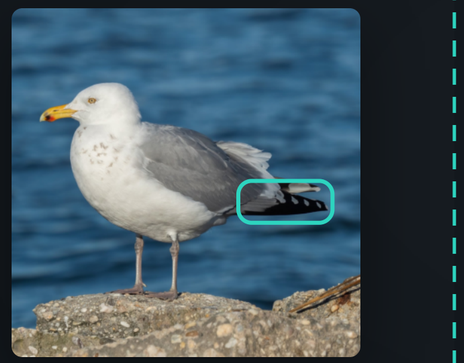

`black wingtip`

bird expert

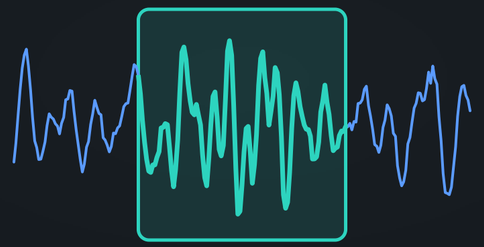

`spindle present`

sleep scorer

"…I was charged twice…"

`duplicate charge`

support agent

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

PART 4

From concepts to a prediction

Follow one image through the task head

From concepts to predictions: the arithmetic

Follow the image all the way to class probabilities

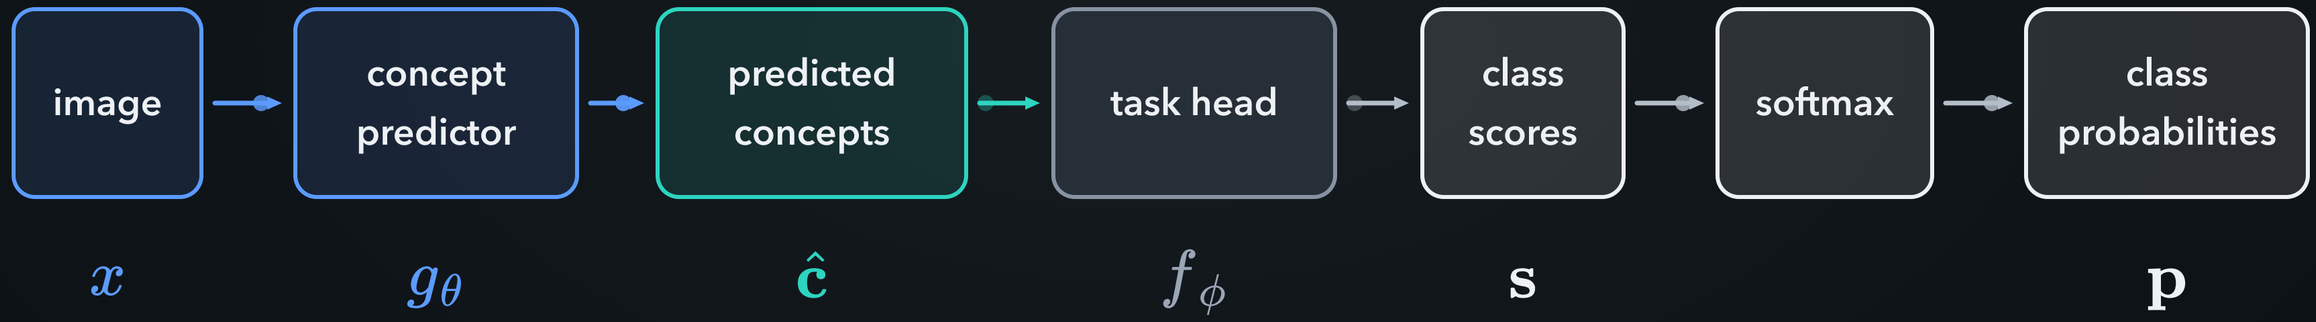

task head: a linear layer or a multilayer perceptron (MLP)

our worked example uses one linear layer

From concepts to predictions: the arithmetic

Three classes and three concepts

classes

zebra

horse

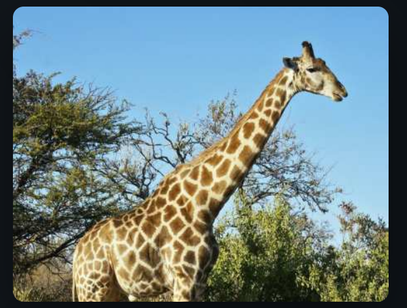

giraffe

concepts

stripes

horse-like shape

long neck

Photo: Yathin S Krishnappa, Wikimedia Commons, CC BY-SA 3.0
Photo: Dietmar Rabich, Wikimedia Commons, CC BY-SA 4.0
Photo: Hans Hillewaert, Wikimedia Commons, CC BY-SA 4.0

ILLUSTRATIVE VALUES · hand-chosen concept outputs and weights; arithmetic computed from these values

From concepts to predictions: the arithmetic

The concept predictor turns this image into three estimates

input image

concept predictor

g_θ

$\hat{\mathbf{c}}$

stripes	0.1
horse-like shape	0.9
long neck	0.1

assumed outputs for this photo

$$\hat{\mathbf{c}} = g_\theta(x)$$

the hat means predicted

Photo: Yathin S Krishnappa, Wikimedia Commons, CC BY-SA 3.0

ILLUSTRATIVE VALUES · hand-chosen concept outputs and weights; arithmetic computed from these values

From concepts to predictions: the arithmetic

The task head turns predicted concepts into class scores

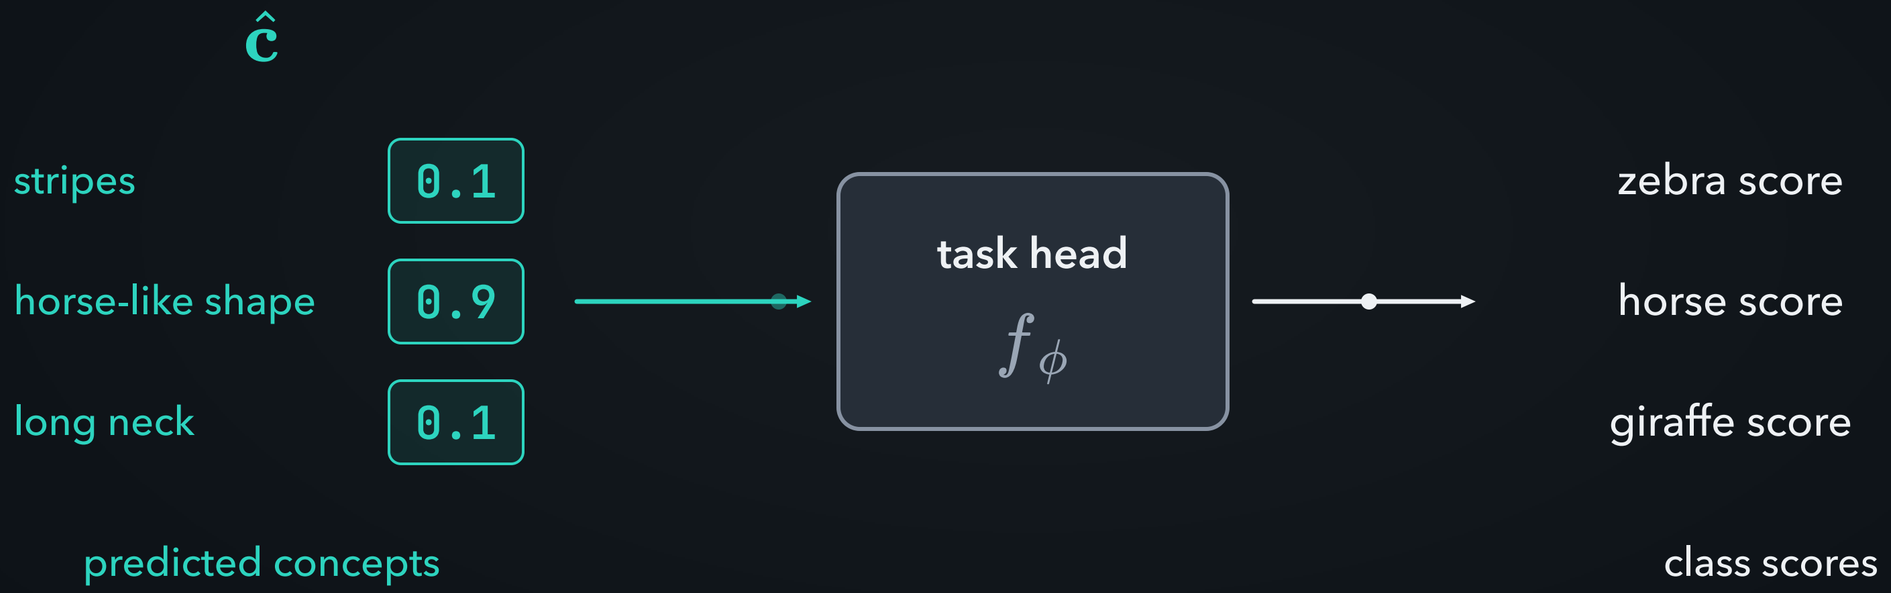

The task head reads only the concepts, not the image

ILLUSTRATIVE VALUES · hand-chosen concept outputs and weights; arithmetic computed from these values

From concepts to predictions: the arithmetic

The concept values describe this image; the weights belong to the layer

values for this image

stripes	0.1
horse-like shape	0.9
long neck	0.1

$\hat{c}$

change with the image

weights inside the linear layer

	stripes	horse-like shape	long neck
zebra	3	1	0
horse	-1	2	-1
giraffe	-1	0	3

W

fixed during prediction

this example has no bias

ILLUSTRATIVE VALUES · hand-chosen concept outputs and weights; arithmetic computed from these values

From concepts to predictions: the arithmetic

Inside the layer: multiply each concept value by its weight, then add

	stripes		horse-like shape		long neck		
zebra weights	3	×	1	×	0		
this image's values	0.1		0.9		0.1		
contributions	0.3	+	0.9	+	0	=	1.2
							zebra score

$$\text{zebra score} = 3 \times 0.1 + 1 \times 0.9 + 0 \times 0.1 = 1.2$$

ILLUSTRATIVE VALUES · hand-chosen concept outputs and weights; arithmetic computed from these values

From concepts to predictions: the arithmetic

The same layer produces all three class scores

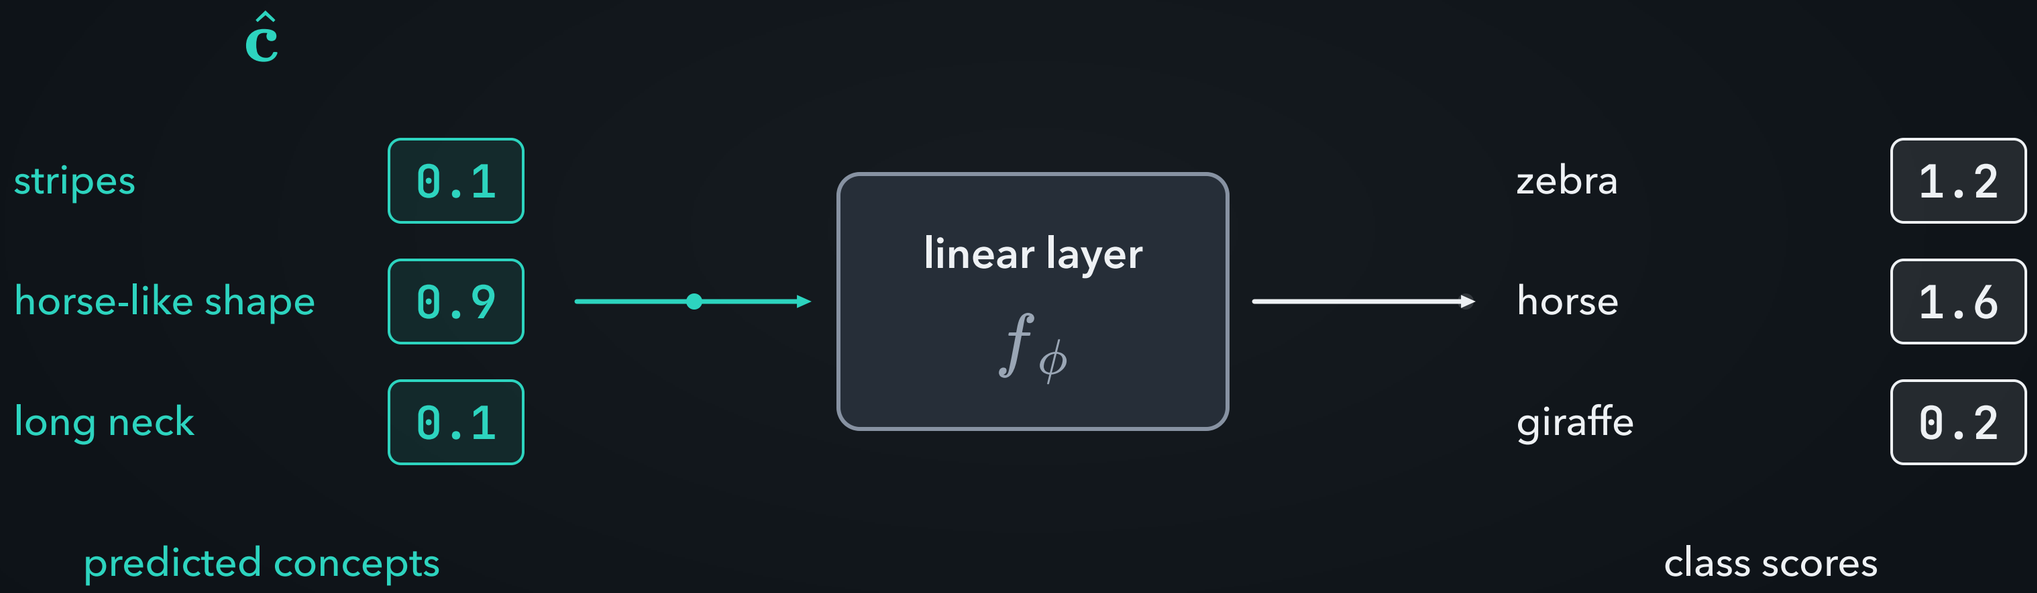

scores = layer weights × predicted concept values

$$\mathbf{s} = f_\phi(\hat{\mathbf{c}}) = W\hat{\mathbf{c}}$$

ILLUSTRATIVE VALUES · hand-chosen concept outputs and weights; arithmetic computed from these values

From concepts to predictions: the arithmetic

Softmax turns the scores into class probabilities

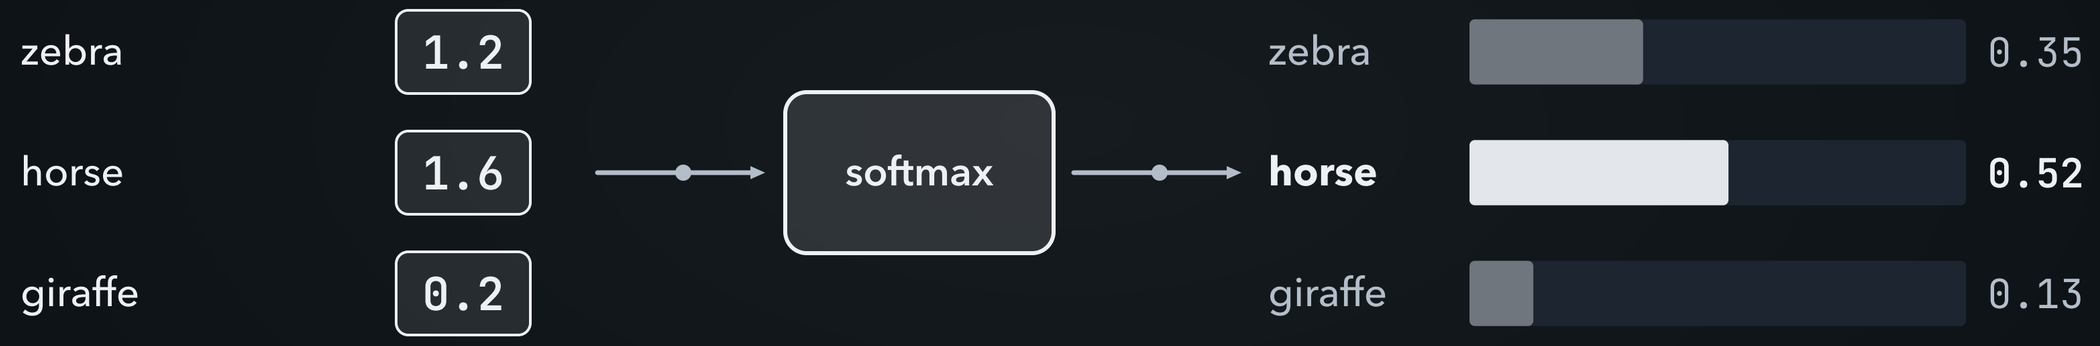

$$p_k = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

probabilities sum to one

ILLUSTRATIVE VALUES · hand-chosen concept outputs and weights; arithmetic computed from these values

From concepts to predictions: the arithmetic

x g_θ $\hat{c}$ f_ϕ $\hat{y}$

A person corrects one predicted concept

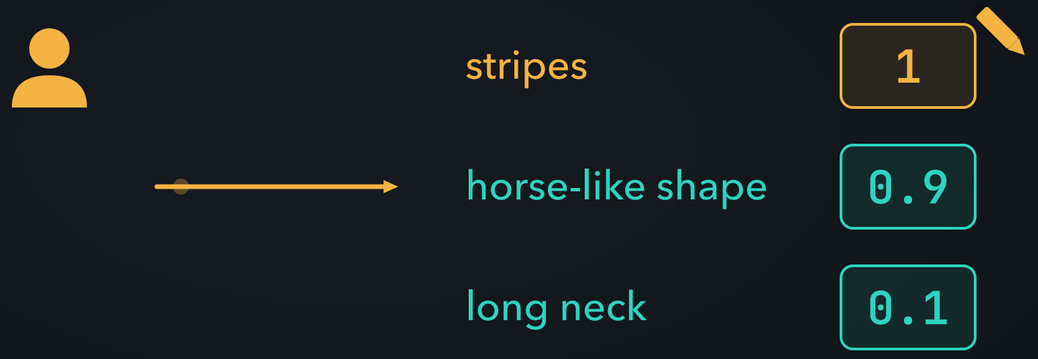

$$\text{stripes: } 0.1 \rightarrow 1$$

Photo: Yathin S Krishnappa, Wikimedia Commons, CC BY-SA 3.0

ILLUSTRATIVE VALUES · hand-chosen concept outputs and weights; arithmetic computed from these values

From concepts to predictions: the arithmetic

Pass the corrected concepts through the same layer

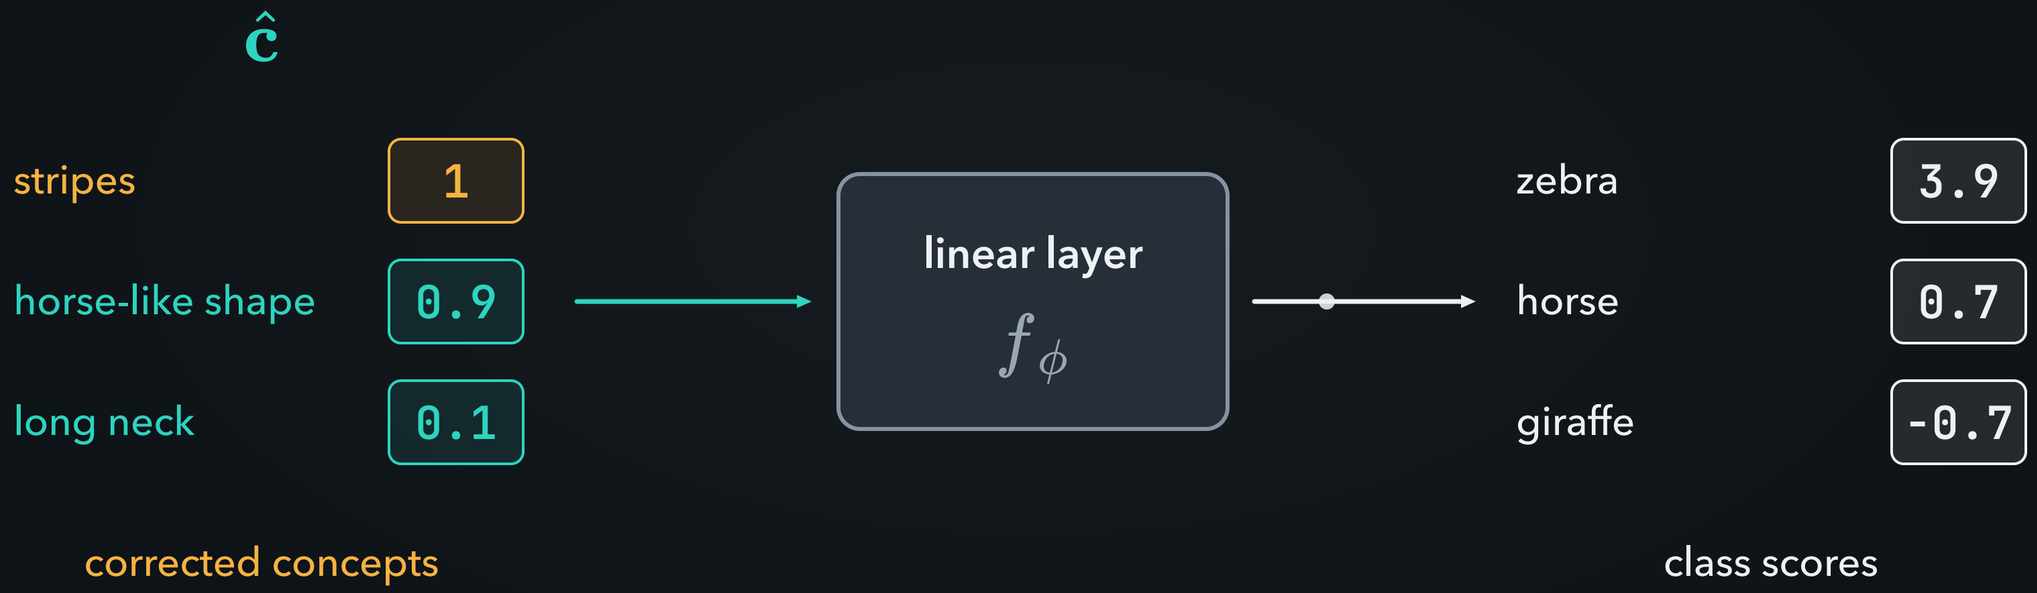

same image, same predictor, same weights; one corrected concept

ILLUSTRATIVE VALUES · hand-chosen concept outputs and weights; arithmetic computed from these values

From concepts to predictions: the arithmetic

x g_θ $\hat{c}$ f_ϕ $\hat{y}$

After softmax, zebra has the highest probability

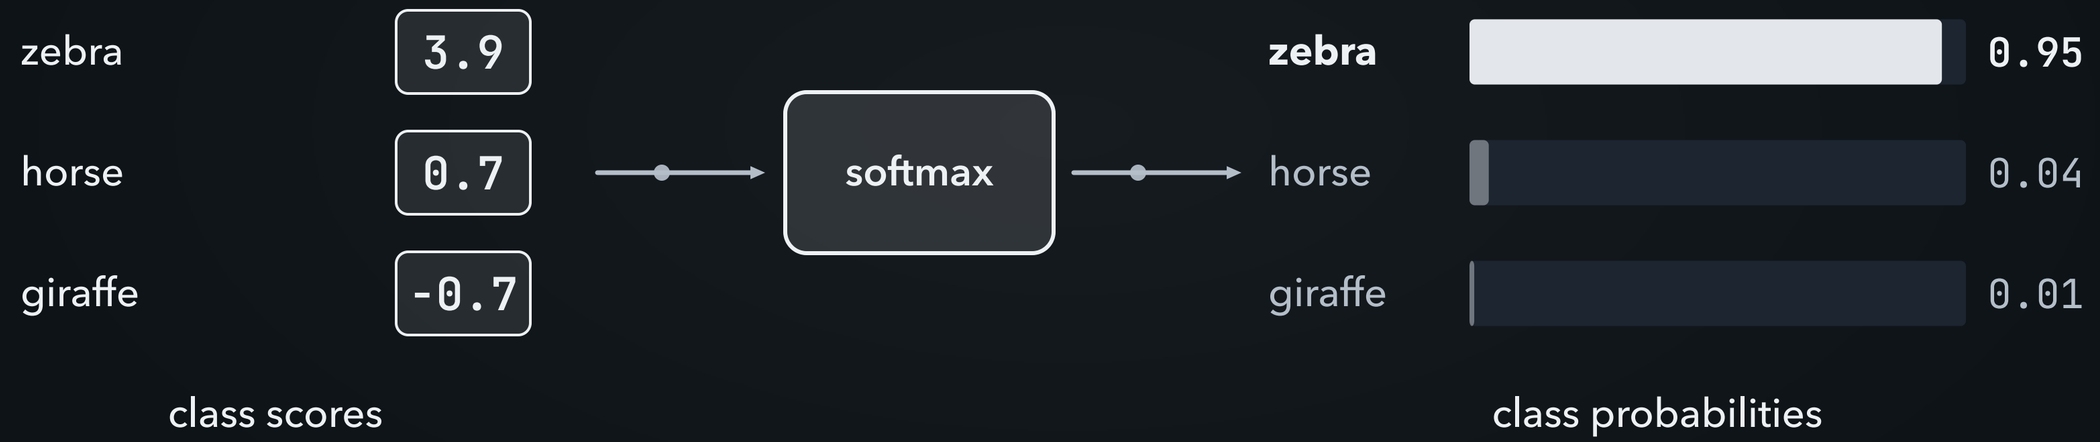

before: horse 0.52

after: zebra 0.95

ILLUSTRATIVE VALUES · hand-chosen concept outputs and weights; arithmetic computed from these values

From concepts to predictions: the arithmetic

Why did the zebra score rise?

change in stripes		zebra's stripes weight		change in zebra score
0.9	×	3	=	$+2.7$

$$\text{zebra score: } 1.2 + 2.7 = 3.9$$

all other contributions stay the same

ILLUSTRATIVE VALUES · hand-chosen concept outputs and weights; arithmetic computed from these values

From concepts to predictions: the arithmetic

x g_θ $\hat{c}$ f_ϕ $\hat{y}$

The concept vector is the point where a person can intervene

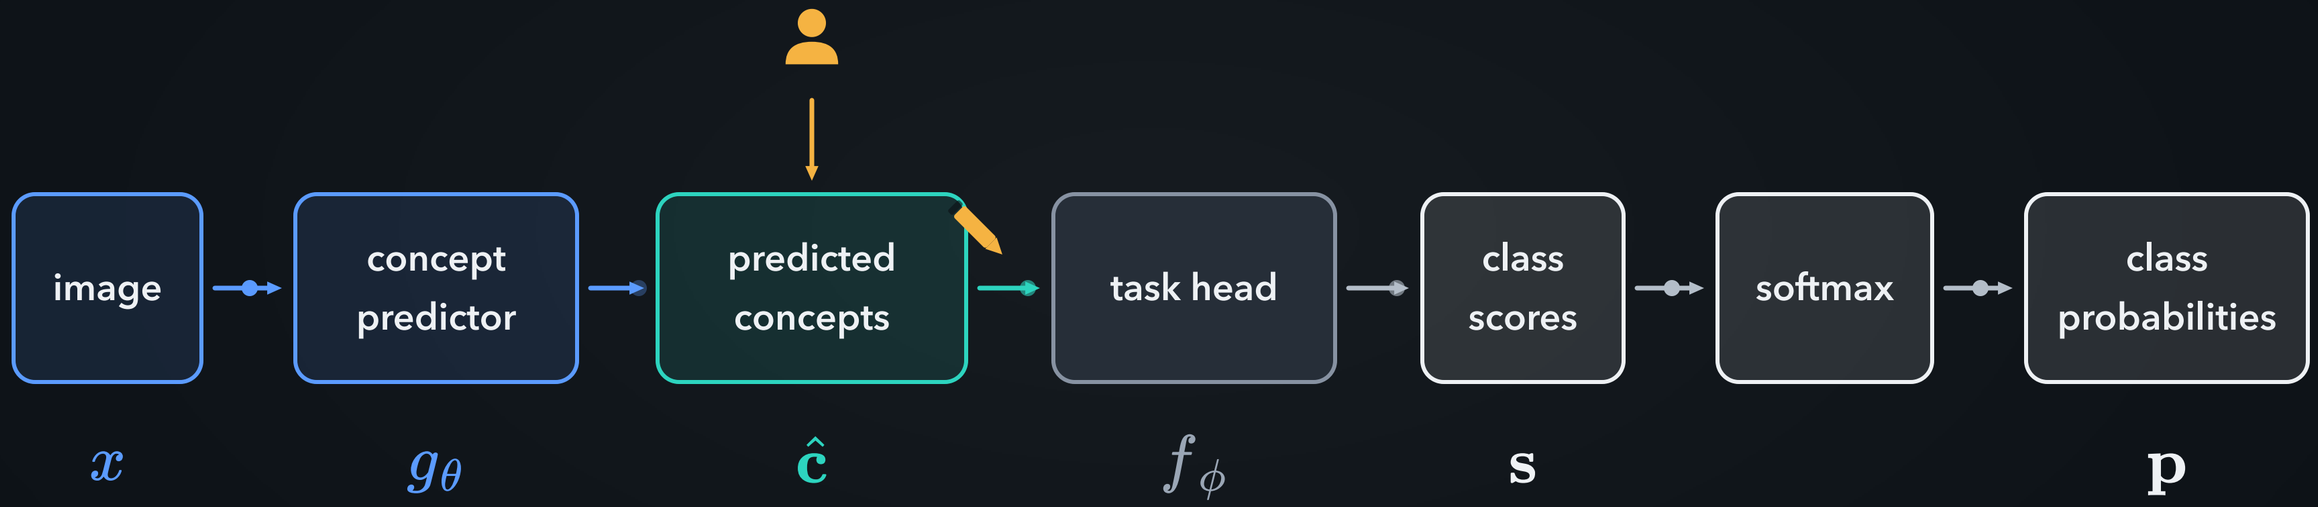

correct a concept, then run the same task head again

no retraining

PART 5

Learning the concepts

What do the labels teach each part of the model?

How does the model learn the concepts?

x g_θ $\hat{c}$ f_ϕ $\hat{y}$

Where do the concept values come from?

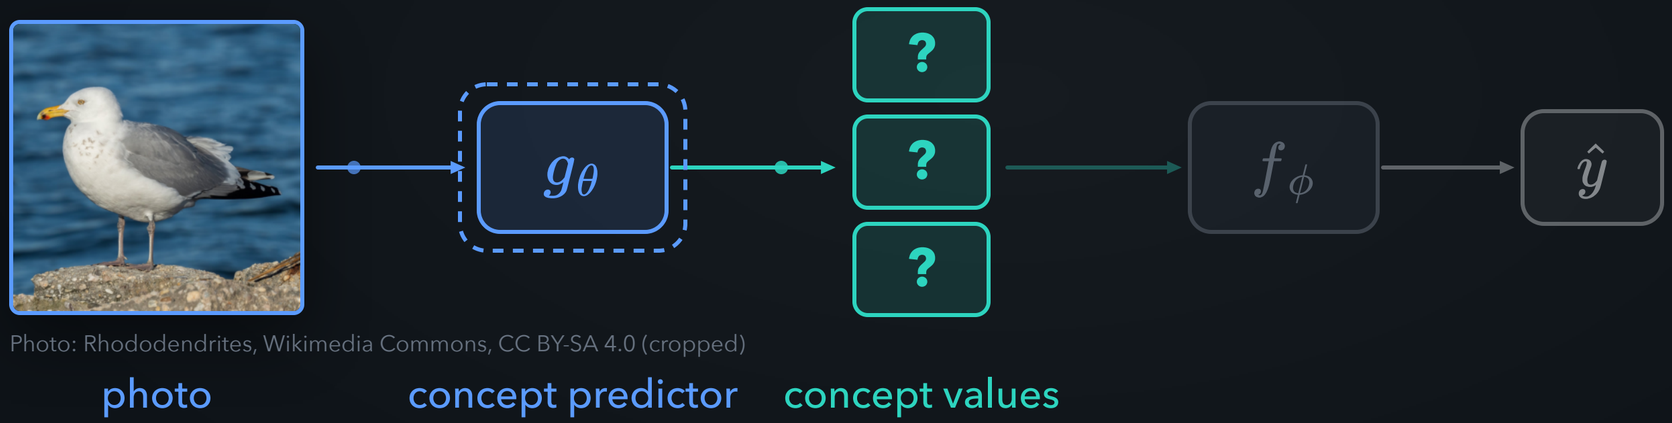

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

photo concept predictor concept values

learned from labelled examples

How does the model learn the concepts?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Each training example carries three things

hooked bill	1
grey wing	1
yellow wing	0
white throat	1

... 112 labels in all

Herring Gull

input x_i

concept labels, 0 or 1 $\mathbf{c}_i$

task label y_i

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

DATASET · Herring Gull's labels in CUB, 4 of 112; our photo

How does the model learn the concepts?

The concept loss compares each predicted concept with its label

for each concept, compare the prediction $\hat{c}_k$ with the label c_k

loss for concept k — label 1: a low prediction costs — label 0: a high prediction costs

$$\ell_k = -c_k \log \hat{c}_k - (1 - c_k) \log(1 - \hat{c}_k)$$

binary cross-entropy, averaged over the concepts: $L_c = \frac{1}{K} \sum_k \ell_k$

How does the model learn the concepts?

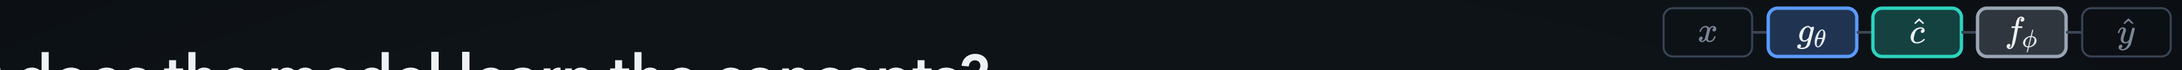

One term of the concept loss, worked out

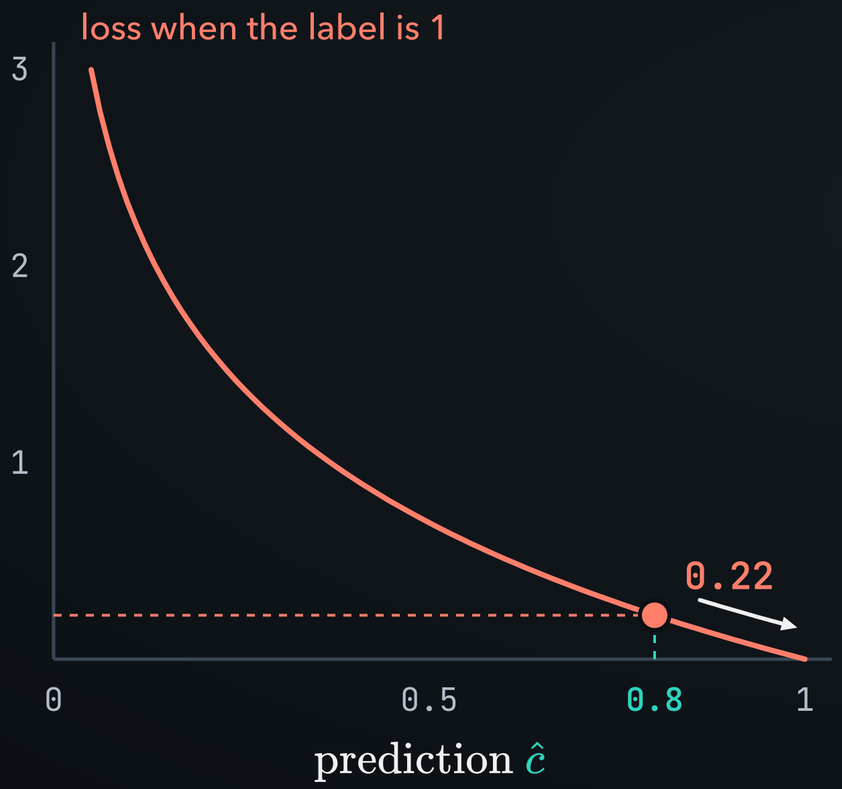

grey wing: label 1, prediction 0.8

$$\ell = -\left[1 \cdot \log 0.8 + 0 \cdot \log(1 - 0.8)\right]$$

$$= -\log 0.8 \approx 0.22$$

closer to 1, smaller loss

ILLUSTRATIVE VALUES · a hand-chosen prediction

How does the model learn the concepts?

The task loss compares the prediction with the task label

task loss: minus the log of the share given to the correct **class**

$$L_y = -\log \hat{y}_{y_i}$$

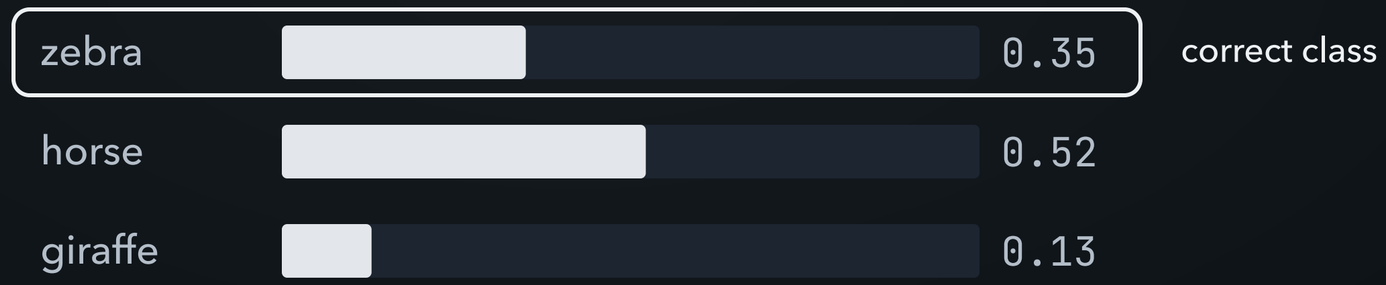

the zebra photo of Part 4: $L_y = -\log 0.35 \approx 1.05$

ILLUSTRATIVE VALUES · shares from the Part 4 worksheet

PART 5

How does the model learn the concepts?

A weighted sum combines the two losses

total loss — weight × concept loss — weight × task loss

$$L = \lambda_c L_c + \lambda_y L_y$$

λ_c and λ_y set how much each objective counts

How does the model learn the concepts?

Three ways to train the concept predictor and the task head

INDEPENDENT

each part on its own labels

SEQUENTIAL

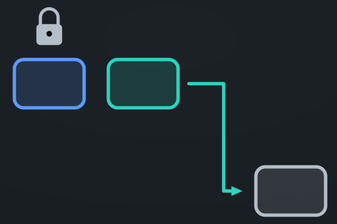

predictor first, then the head

JOINT

everything at once

Koh et al., ICML 2020

How does the model learn the concepts?

Independent: each part is trained on its own labels

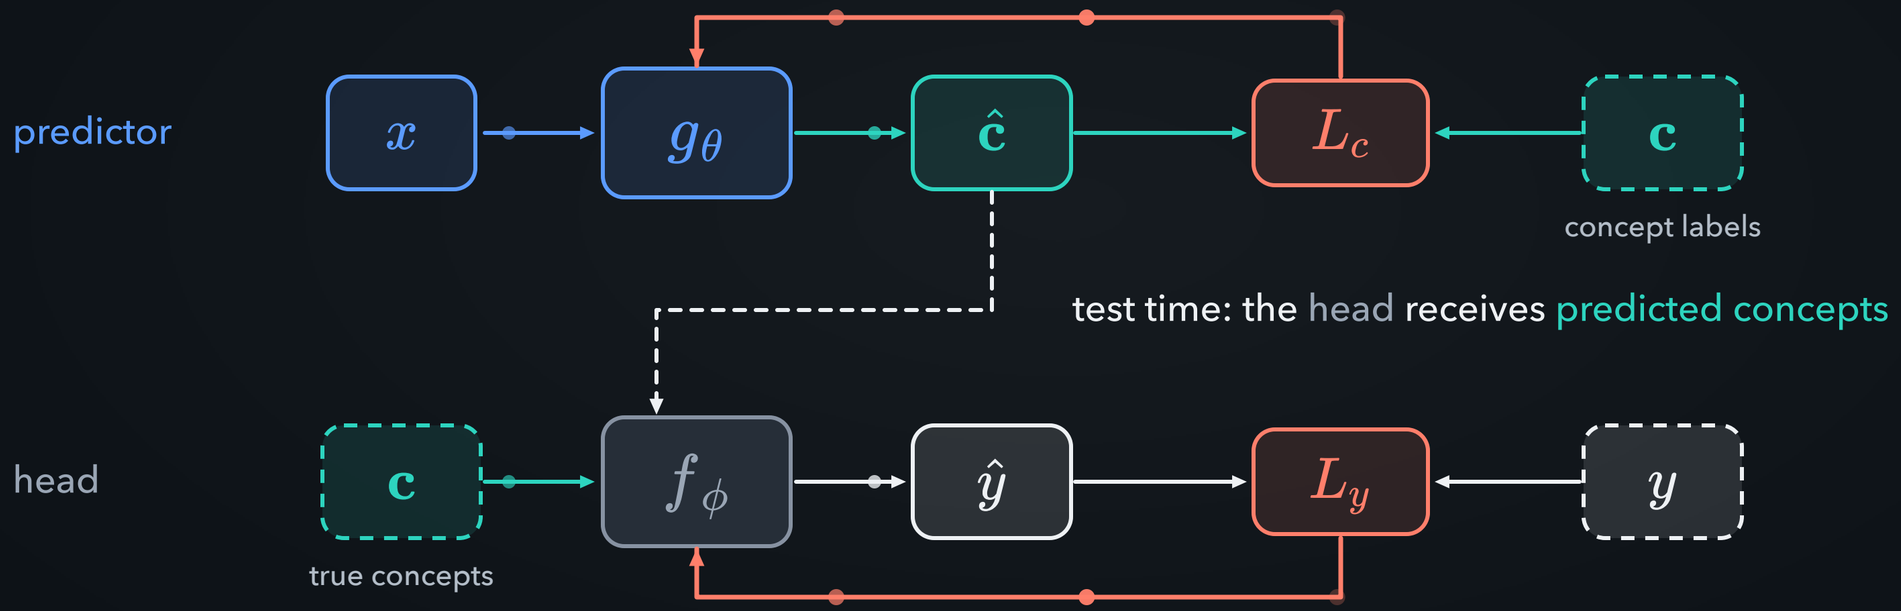

coral: the gradient of each loss, flowing back to what it trains

How does the model learn the concepts?

Sequential: freeze the predictor, then train the head on its outputs

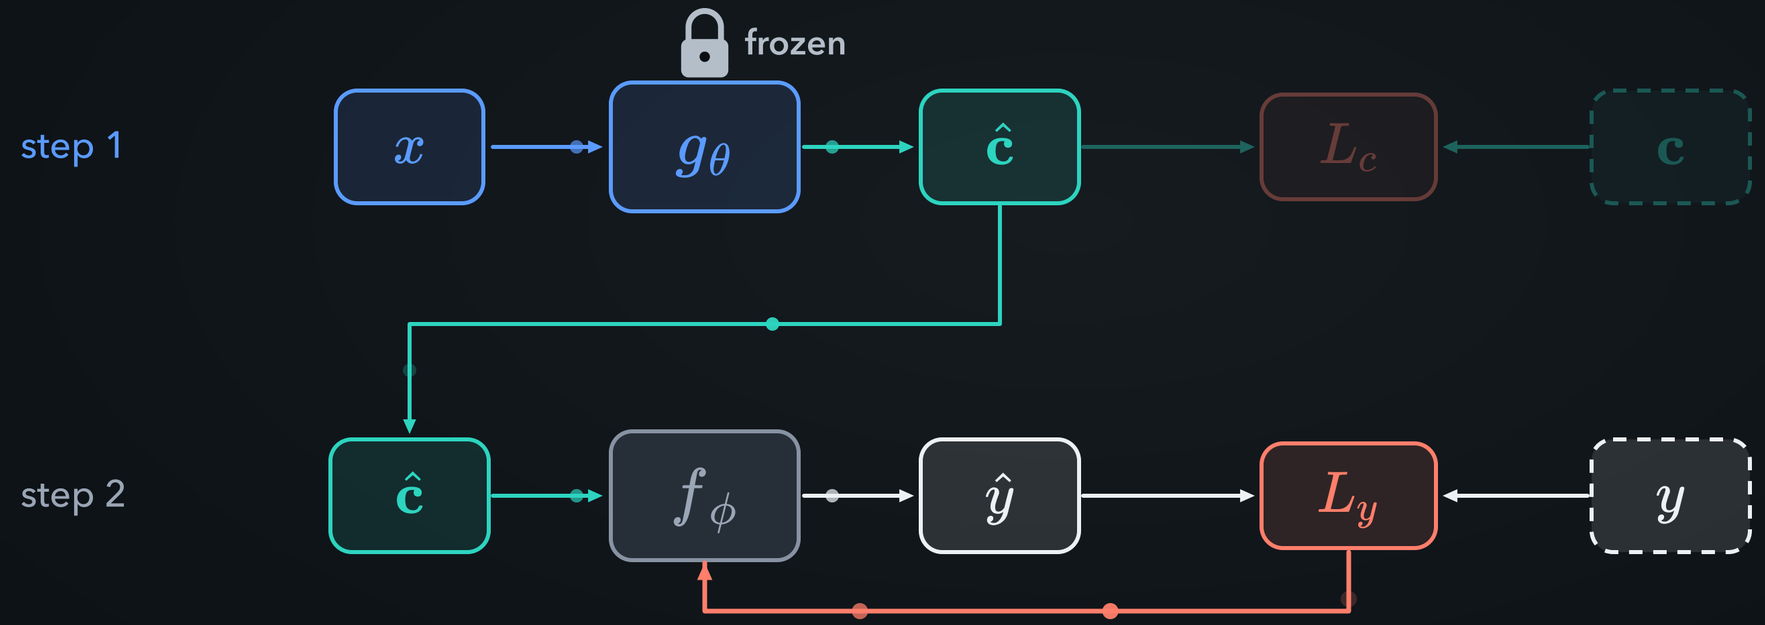

the head learns from the predictor's actual mistakes

How does the model learn the concepts?

Joint: train the whole system at once, with both losses

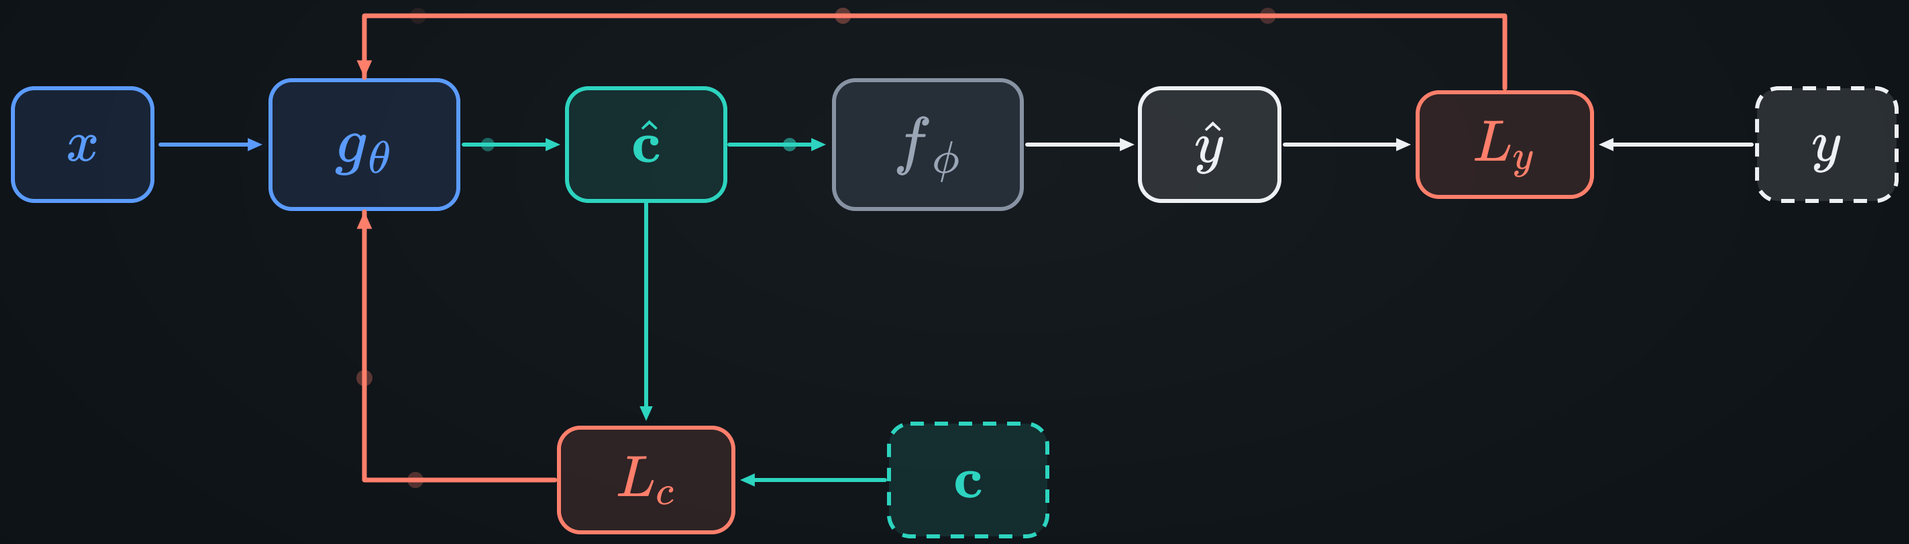

both losses train the whole system

How does the model learn the concepts?

In joint training, the task loss also shapes the concept predictor

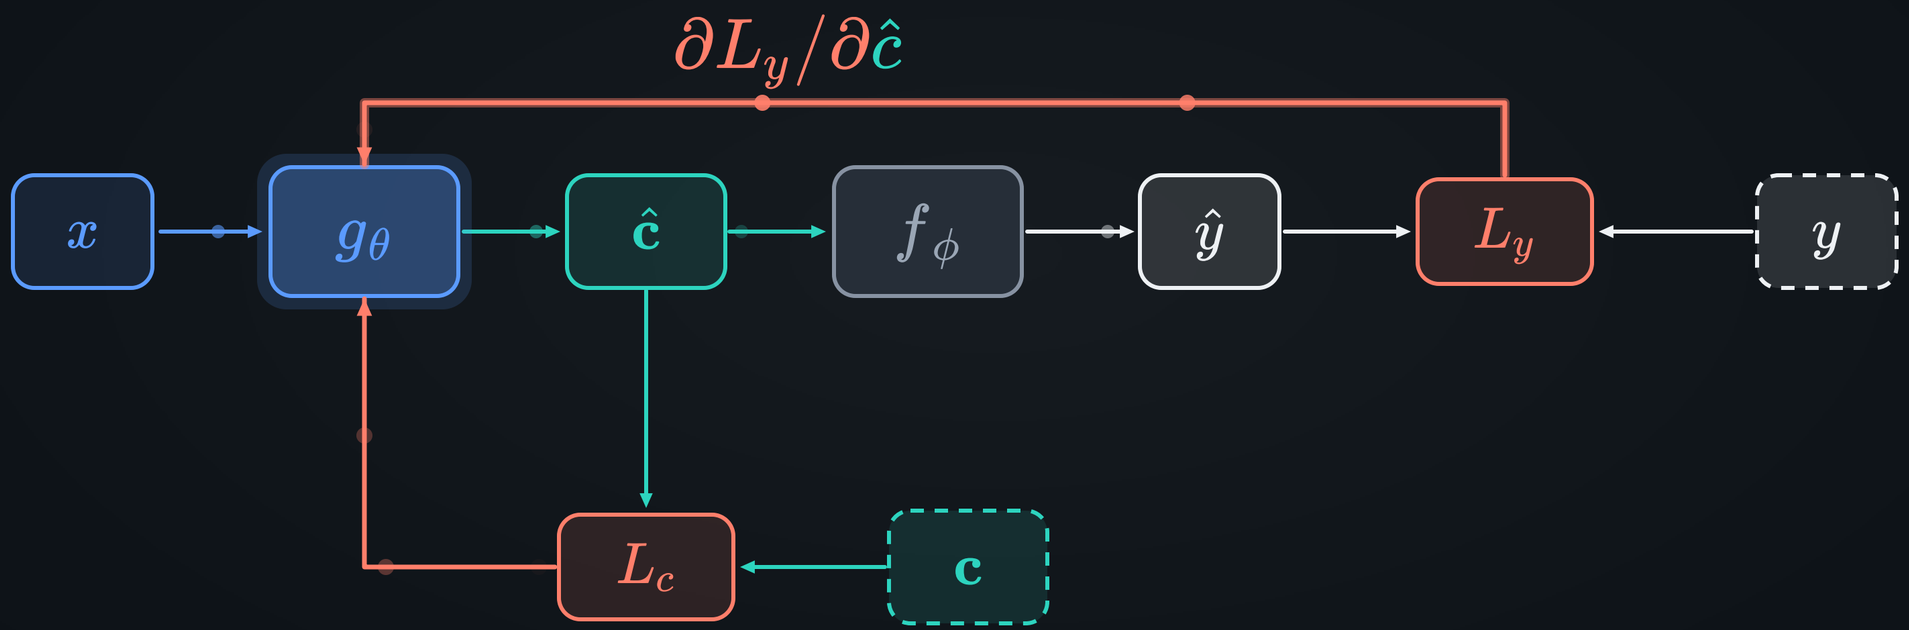

the concept predictor is shaped by both losses

How does the model learn the concepts?

A concept score can be pulled two ways

can improve **task accuracy**, can hurt concept fidelity

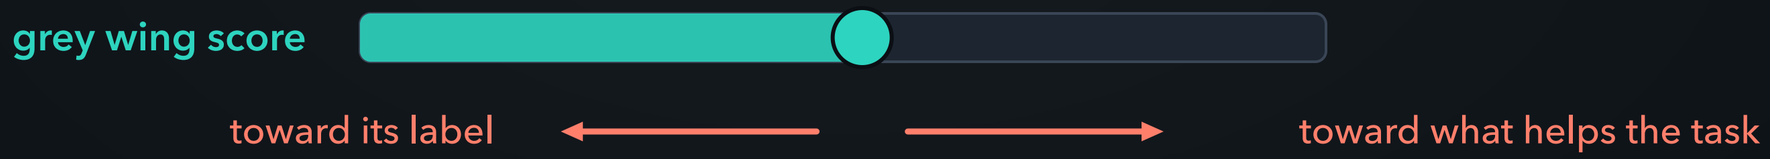

whether it happens has to be measured

How does the model learn the concepts?

Our measured run: joint training, $\lambda_c = 2$, $\lambda_y = 1$

concept loss, weight 2 task loss, weight 1

$$L = 2\,L_c + 1\,L_y$$

joint training, in the experiment of Part 11

MEASURED · settings of the trained shape model, Part 11

The three schemes on CUB

Does the training scheme matter on real photographs?

independent

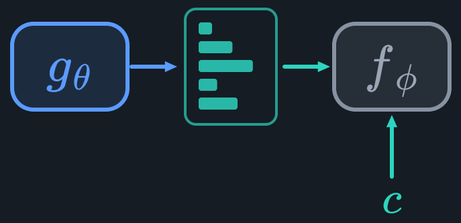

head fitted on true concepts

sequential

predictor frozen, head fitted on its predictions

joint

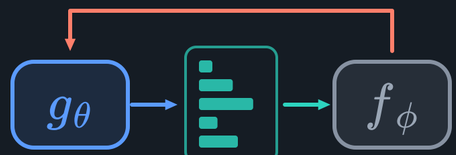

both trained together on both losses

our small run on CUB: 5,994 training photos, 5,794 test photos

The three schemes on CUB

x g_θ $\hat{c}$ f_ϕ $\hat{y}$

Our setup: frozen CLIP features → 112 concepts → 200 species

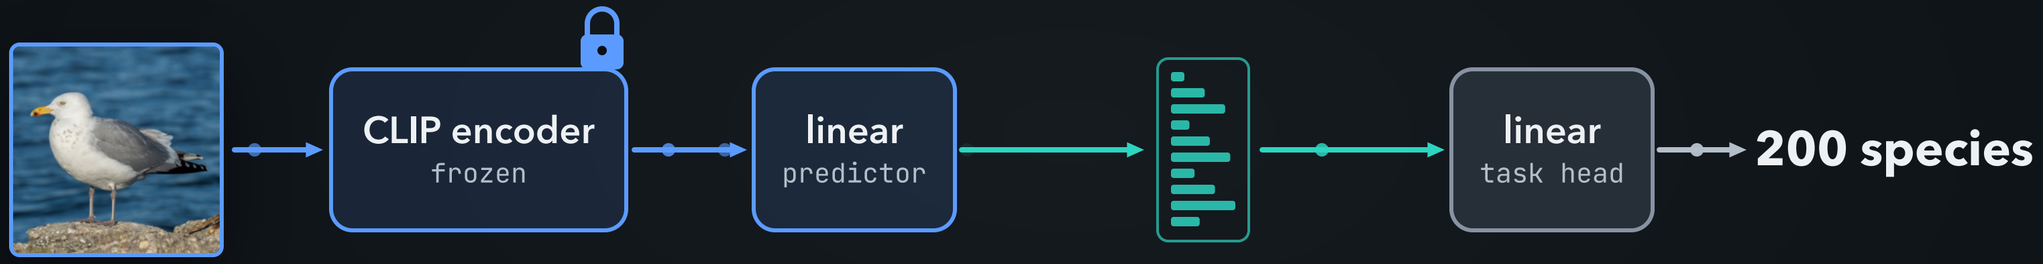

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

x $e(x) \in \mathbb{R}^{512}$ g_θ $\hat{\mathbf{c}} \in [0,1]^{112}$ f_ϕ $\hat{y}$

CUB official split: 5,994 training photos, 5,794 test photos

MEASURED · our run: frozen CLIP ViT-B/32 + linear models, CUB official split, one run

The three schemes on CUB

One photo through the model: features, then concepts, then species

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

COMPUTED · REAL MODEL · our gull photo: CLIP ViT-B/32 feature, our CUB model

PART 5

The three schemes on CUB

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Joint training is the most accurate, independent the least

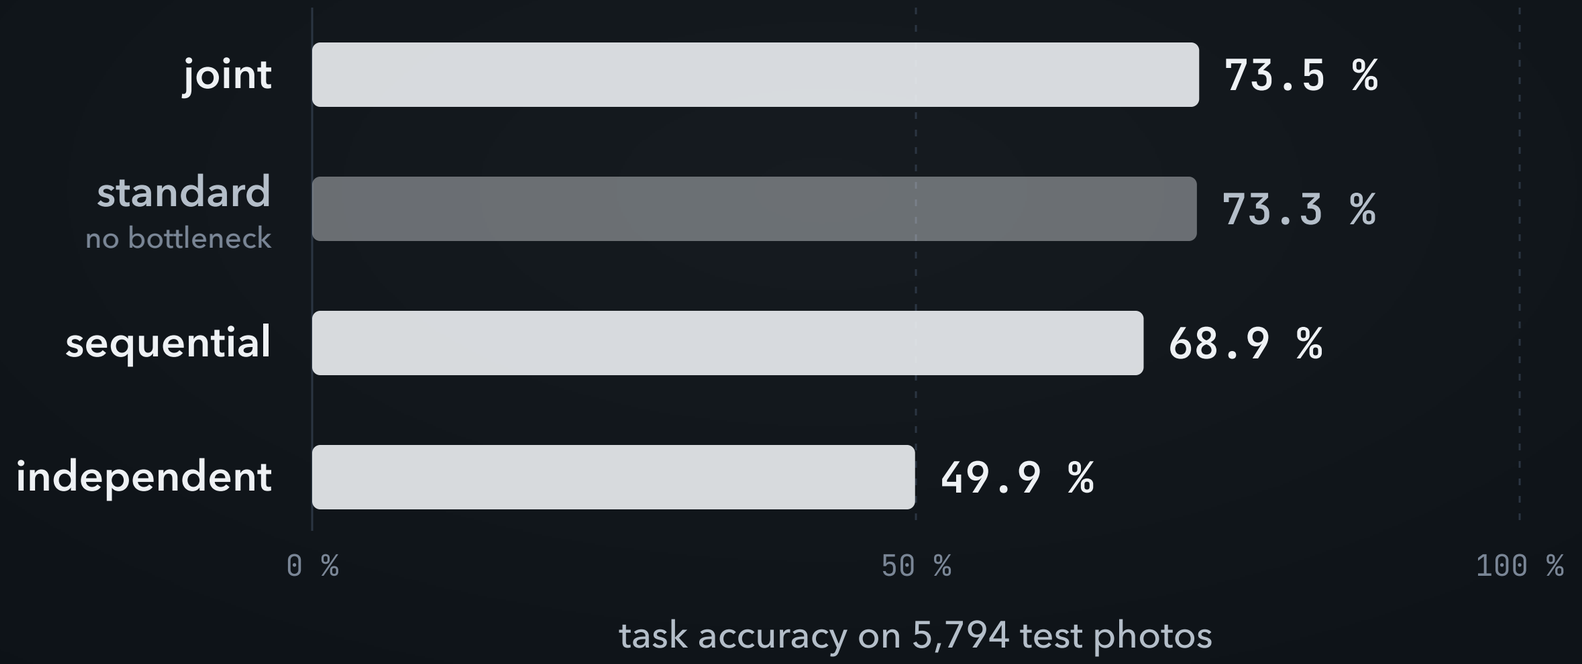

MEASURED · our CUB run, one run

The three schemes on CUB

An intervention replaces predicted values with true values

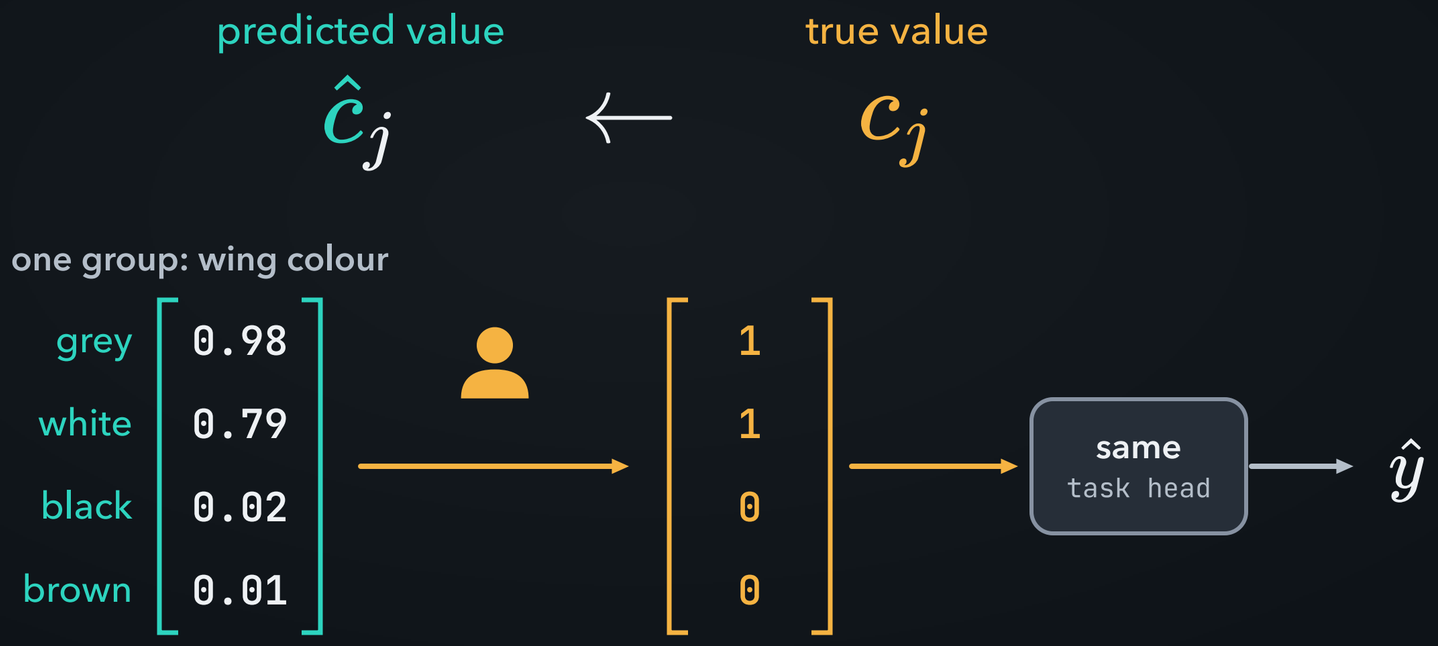

COMPUTED · REAL MODEL · our model on our gull photo; true values: Herring Gull's CUB labels

The three schemes on CUB

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

As more groups are corrected, the independent model gains most

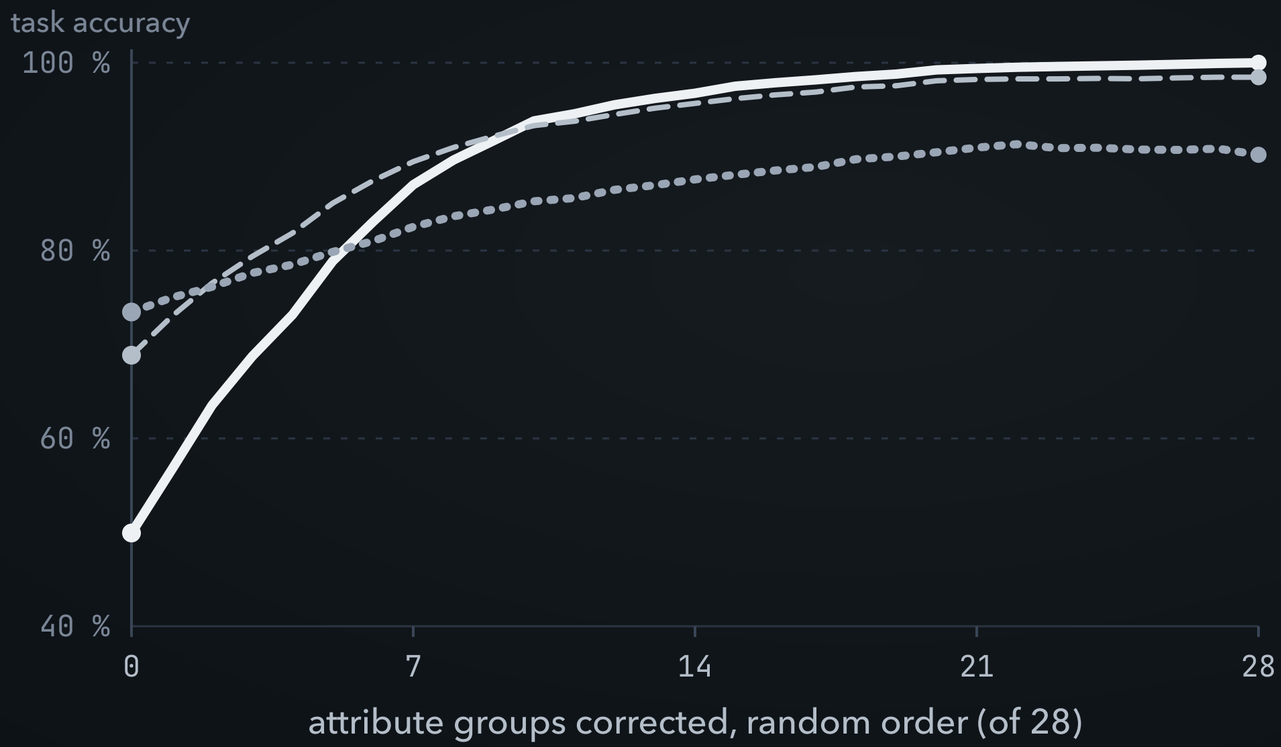

independent	100.0 %
sequential	98.5 %
joint	90.2 %

all 28 groups true: every species has its own pattern, so **independent** reaches 100.0 %

MEASURED · our CUB run, one run, averaged over random orders

PART 6

Human interventions

What changes when a person corrects a concept?

Human interventions

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

A person can fix a concept the model got wrong

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

wingtip black 0.20

the photo and the head stay the same; only a value changes

ILLUSTRATIVE VALUES · hand-written concept value (Part 2)

Human interventions

An intervention has four steps

1. predict concept values
2. a person checks one and corrects it
3. the fixed head recomputes
4. a new **prediction**

$$\hat{\mathbf{c}} = g_\theta(x) \qquad \hat{c}_j \leftarrow c_j \qquad f_\phi \text{ unchanged} \qquad \hat{y}_{\text{new}} = f_\phi(\hat{\mathbf{c}}_{\text{edited}})$$

Human interventions

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Wing colour is one group: correct it together

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

wing colour is one categorical concept: grey → 0, black → 1, together

ILLUSTRATIVE VALUES · hand-written concept values and head

Human interventions

Changing only one value could leave an impossible state

a wingtip has one colour, so the group changes as one

Human interventions

Same head, new values: herring gull rises

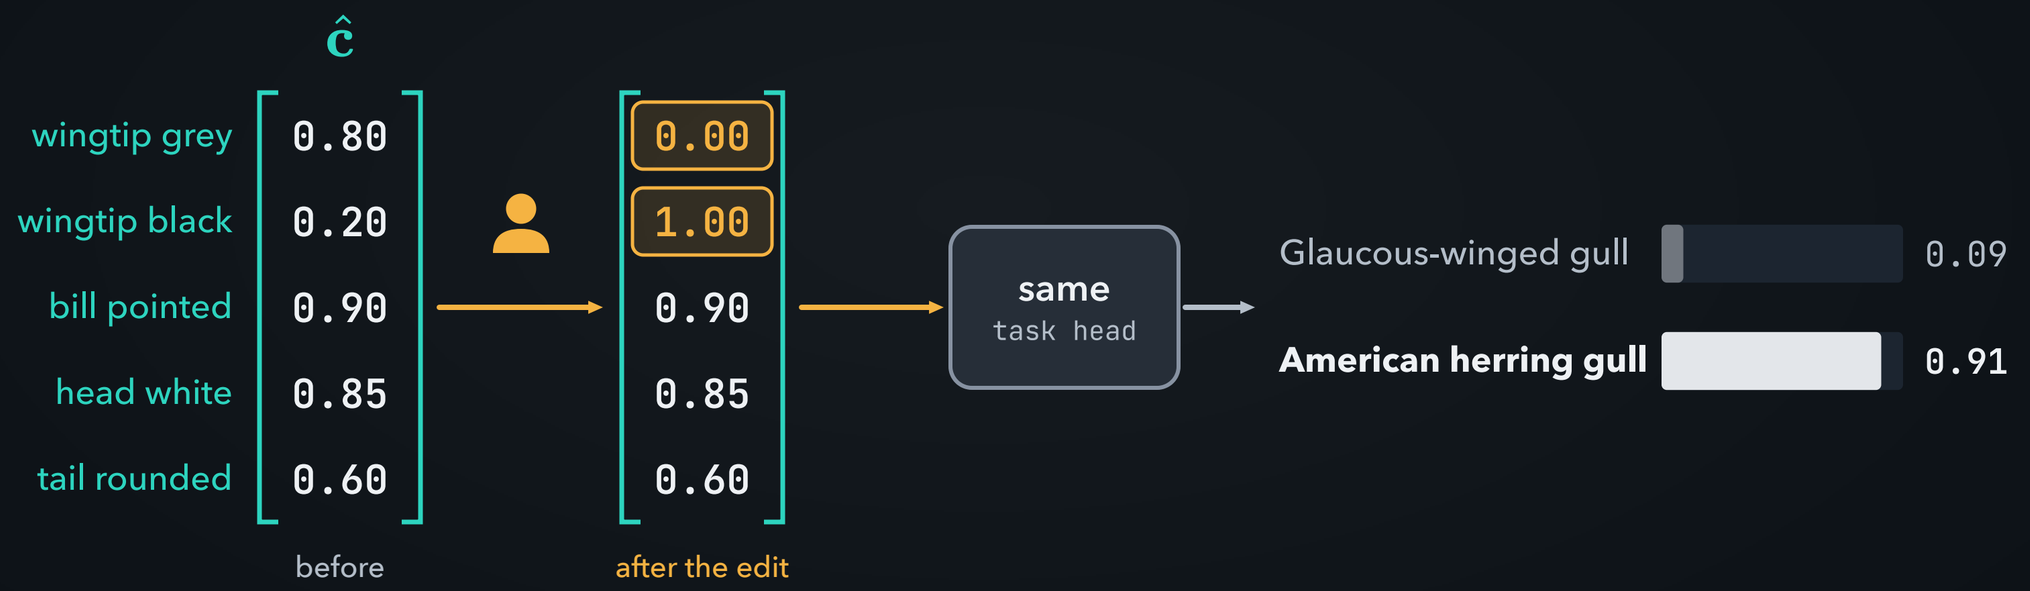

ILLUSTRATIVE VALUES · hand-written concept values and head

Human interventions

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

The same mechanism works on a signal

SCHEMATIC · authored trace, not a recording

Human interventions

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

The concept predictor gives spindle a low value

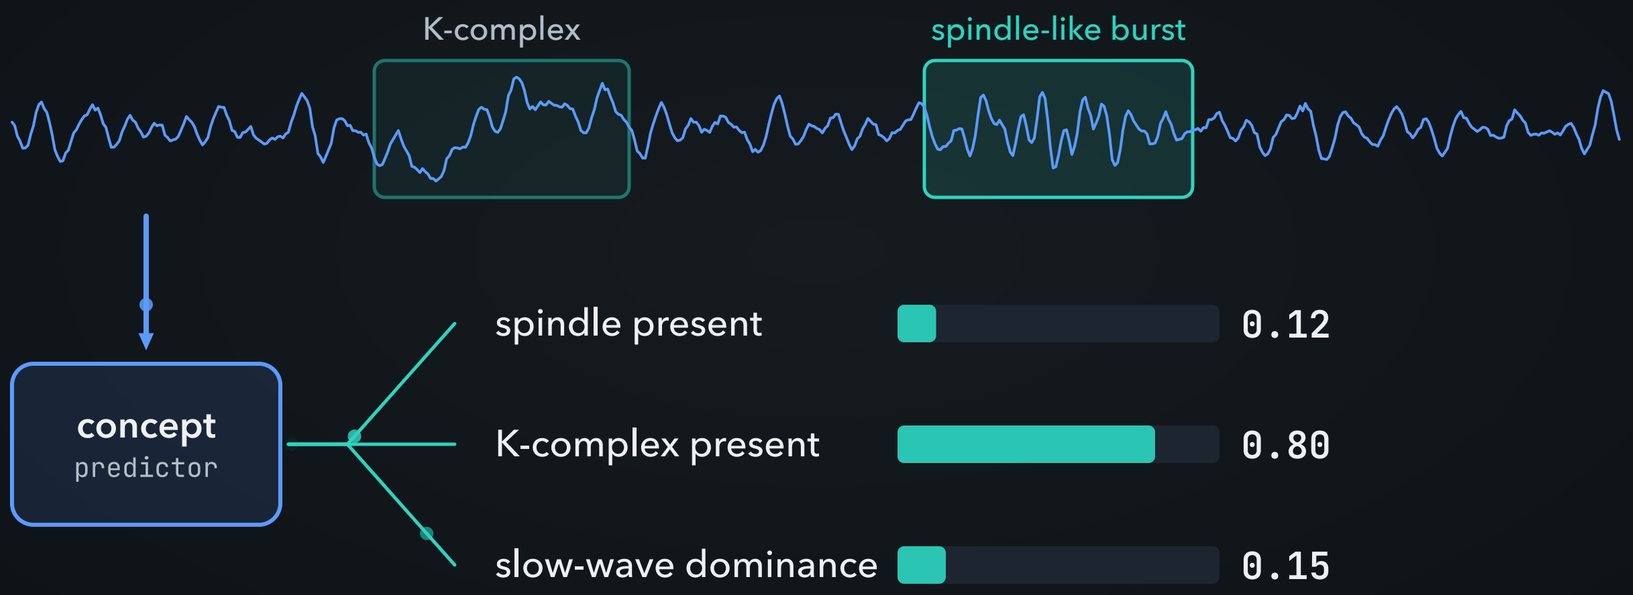

ILLUSTRATIVE VALUES · hand-written concept values; schematic trace

Human interventions

An illustrative stage head favours other stage

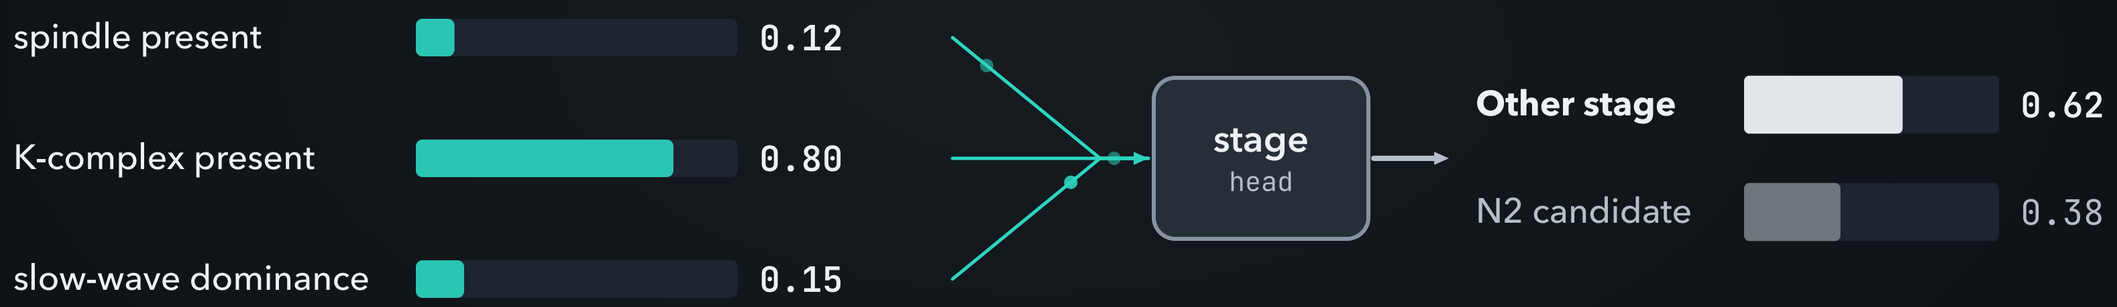

ILLUSTRATIVE VALUES · hand-written stage head

Human interventions

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

A sleep expert marks the spindle as present

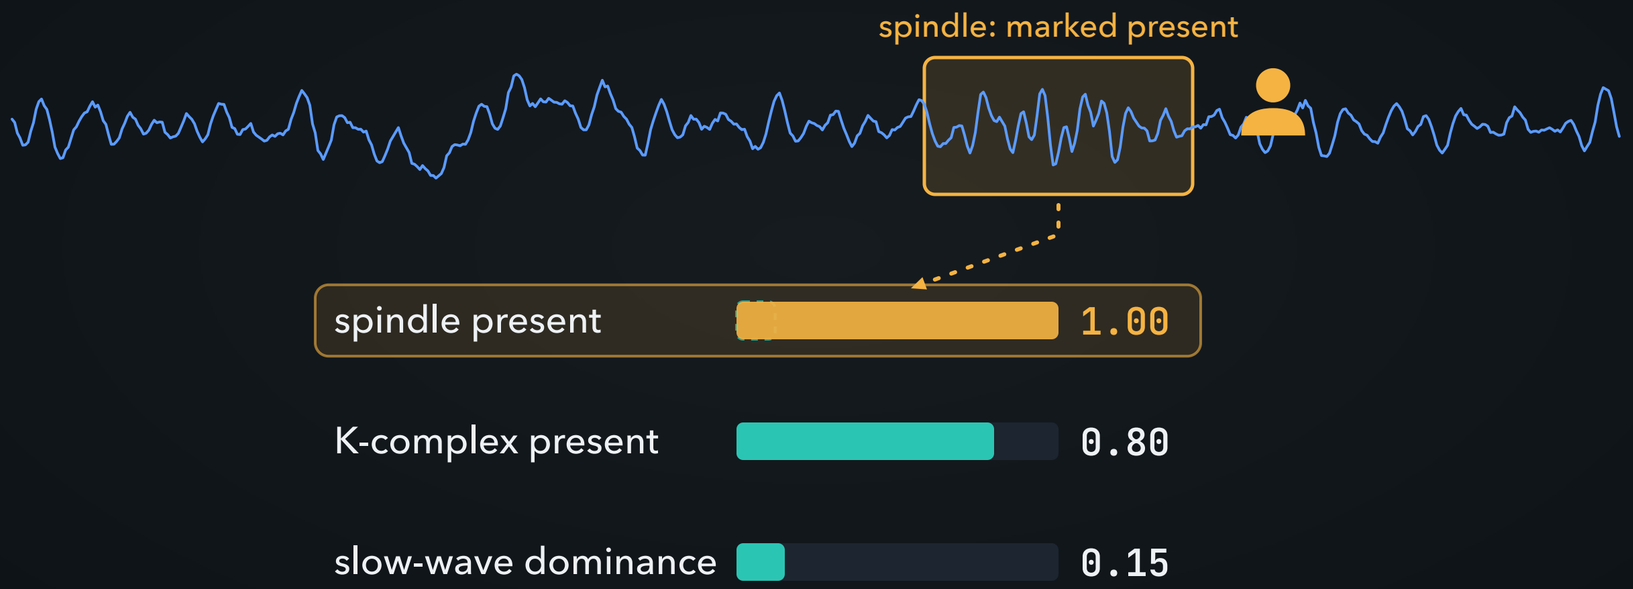

ILLUSTRATIVE VALUES · hand-written concept values; schematic trace

Human interventions

Same head, new input: N2 candidate rises

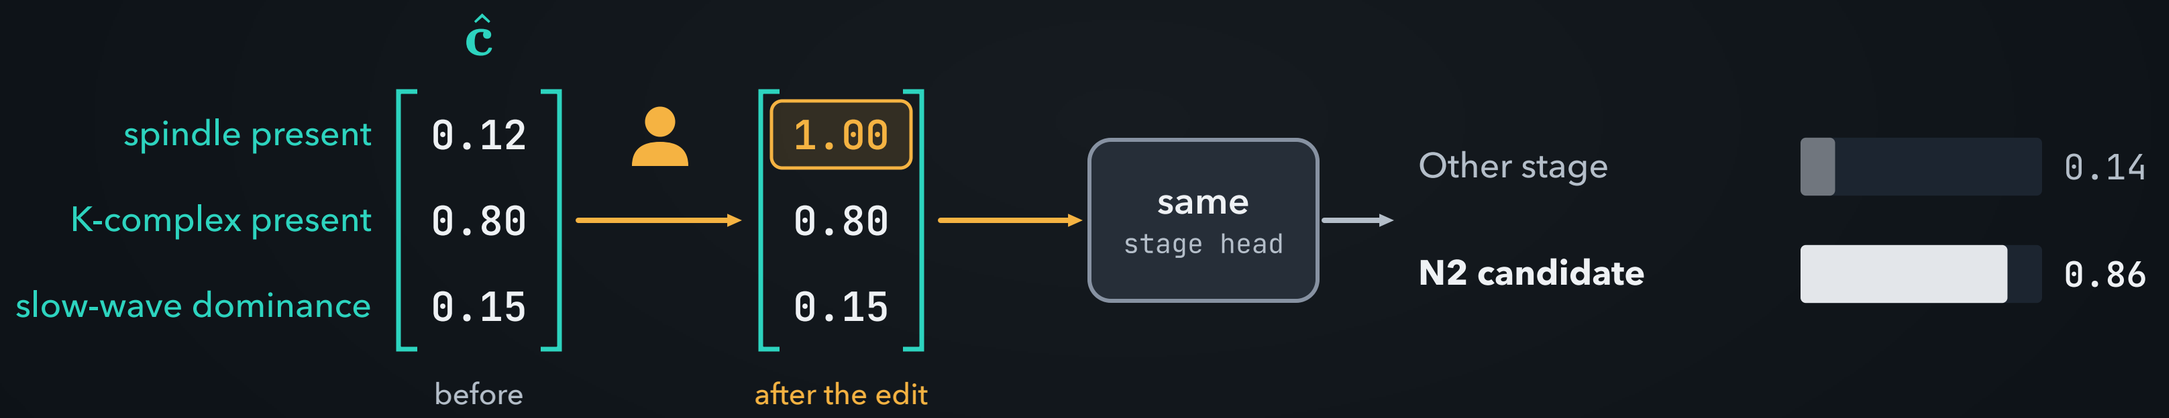

ILLUSTRATIVE VALUES · hand-written stage head

Human interventions

x g_θ $\hat{c}$ f_ϕ $\hat{y}$

A demonstration of the mechanism, not a clinical result

real staging: scoring rules applied to 30 s epochs

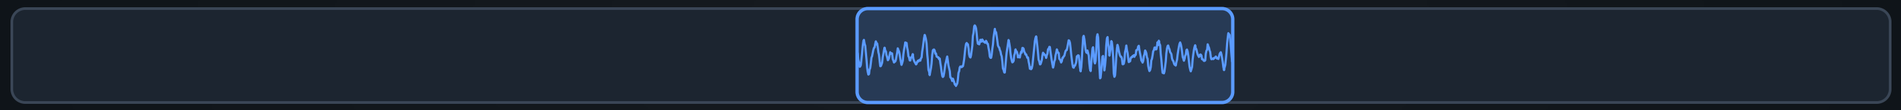

0 s — our 6 s trace — 30 s

this stage head was written by hand

SCHEMATIC · authored EEG trace

ILLUSTRATIVE VALUES · hand-written stage head

PART 7

When a correct edit hurts

A better concept value need not give a better answer

PART 7

A correct edit can still hurt

Does a correct edit always help?

gull, Part 2

herring gull 0.29 → 0.91

✓

zebra, Part 4

zebra 0.35 → 0.95

✓

EEG, Part 6

N2 candidate 0.38 → 0.86

✓

a correct edit

?

✗

so far, every correction helped

PART 7

A correct edit can still hurt

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

A support message that belongs in billing

SUPPORT MESSAGE

The app works, but I was charged twice.
Please refund the extra payment.

BILLING QUEUE

where it belongs

CONSTRUCTED EXAMPLE · an invented message

PART 7

A correct edit can still hurt

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

The concept predictor reads two concepts from the message

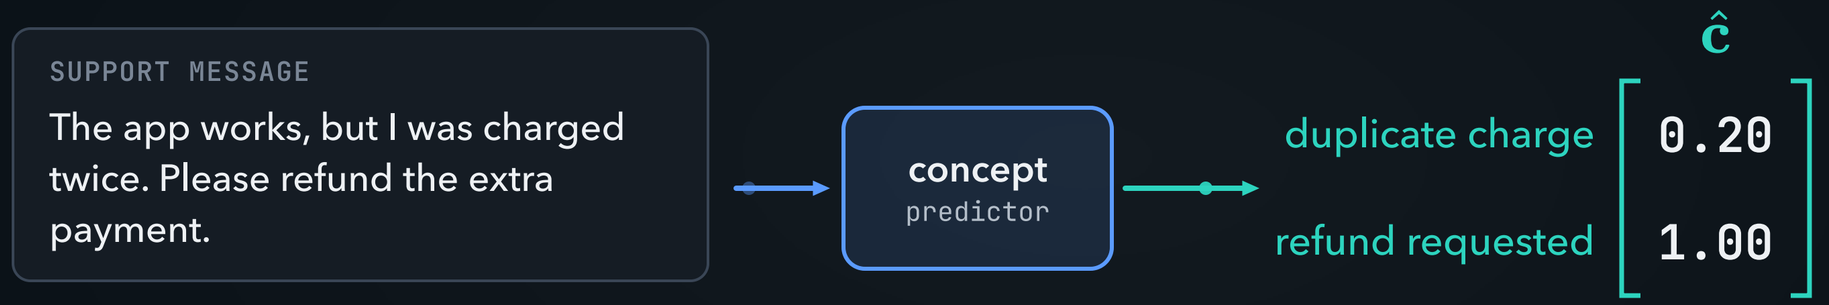

ILLUSTRATIVE VALUES · hand-written concept values

PART 7

A correct edit can still hurt

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

The task head, in words

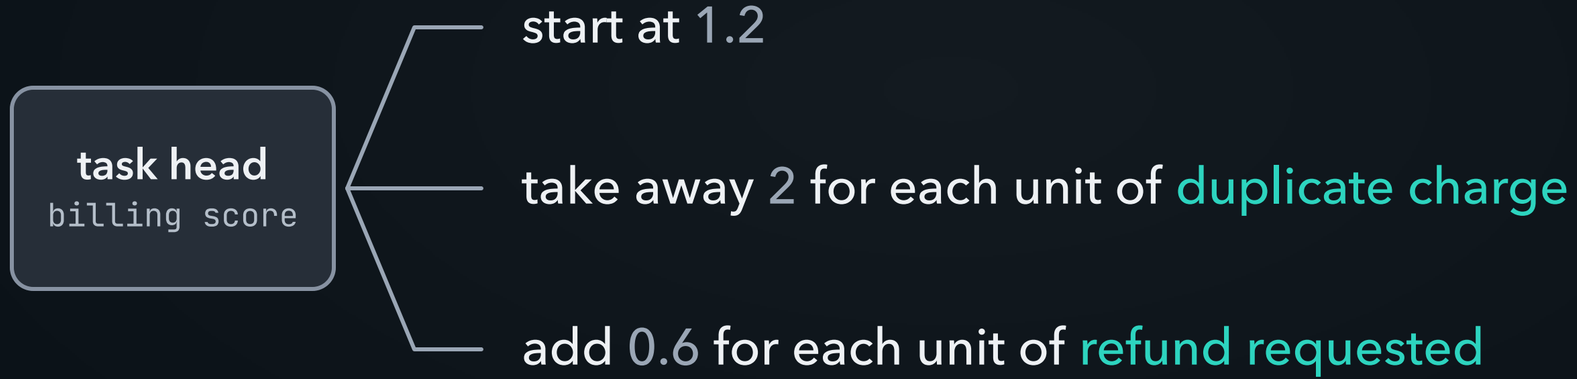

CONSTRUCTED EXAMPLE · a hand-written head

A correct edit can still hurt

The same head in symbols

score start duplicate charge refund requested

$$s = 1.2 - 2\,\hat{c}_{\text{dup}} + 0.6\,\hat{c}_{\text{refund}}$$

a sigmoid turns the score s into the **billing share**

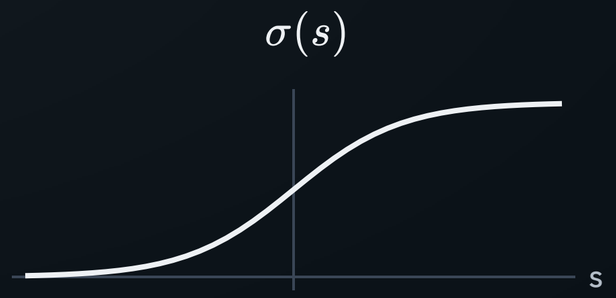

PART 7

A correct edit can still hurt

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Before the edit, the message goes to billing

$$s = 1.2 - 0.4 + 0.6 = 1.4$$

-2×0.20

0.6×1.00

BILLING QUEUE ✓

GENERAL SUPPORT

$\text{billing share} = \sigma(1.4)$

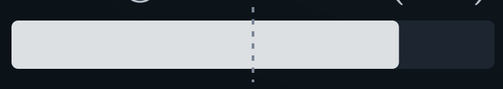

0.80

CONSTRUCTED EXAMPLE · hand-written head

PART 7

A correct edit can still hurt

A reviewer **correctly** sets duplicate charge to 1

SUPPORT MESSAGE

The app works, but I was charged twice. Please refund the extra payment.

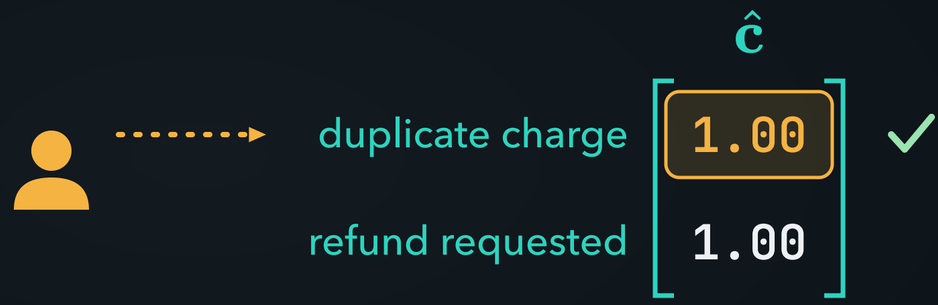

matches the message: charged twice

PART 7

A correct edit can still hurt

After the edit, the message leaves billing

$$s = 1.2 \quad -2 \quad +0.6 \quad = -0.2$$

-2×1.00

0.6×1.00

$$\text{billing share} = \sigma(-0.2)$$

0.45

BILLING QUEUE

GENERAL SUPPORT

CONSTRUCTED EXAMPLE · hand-written head

A correct edit can still hurt

The fault is in the head: a weight with the wrong sign

$$s = 1.2 \; \mathbf{-2}\,\hat{c}_{\text{dup}} + 0.6\,\hat{c}_{\text{refund}}$$

pushes **billing** down as duplicate charge rises

✓ the edit: correct

✗ the weight: wrong sign

CONSTRUCTED EXAMPLE · deliberately faulty head

A correct edit can still hurt

Correcting a concept does not repair the decision rule

measure the effect of an intervention; do not assume it

A correct edit can still hurt

A constructed case: the failure can happen

$$s = 1.2 - 2\,\hat{c}_{\text{dup}} + 0.6\,\hat{c}_{\text{refund}}$$

built faulty on purpose: the failure **can** happen

it says nothing about how often it does

CONSTRUCTED EXAMPLE · deliberately faulty head: existence, not frequency

PART 8

An incomplete concept vocabulary

What if a useful distinction has no name?

What if the concepts are incomplete?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

The bottleneck can pass on only what the vocabulary names

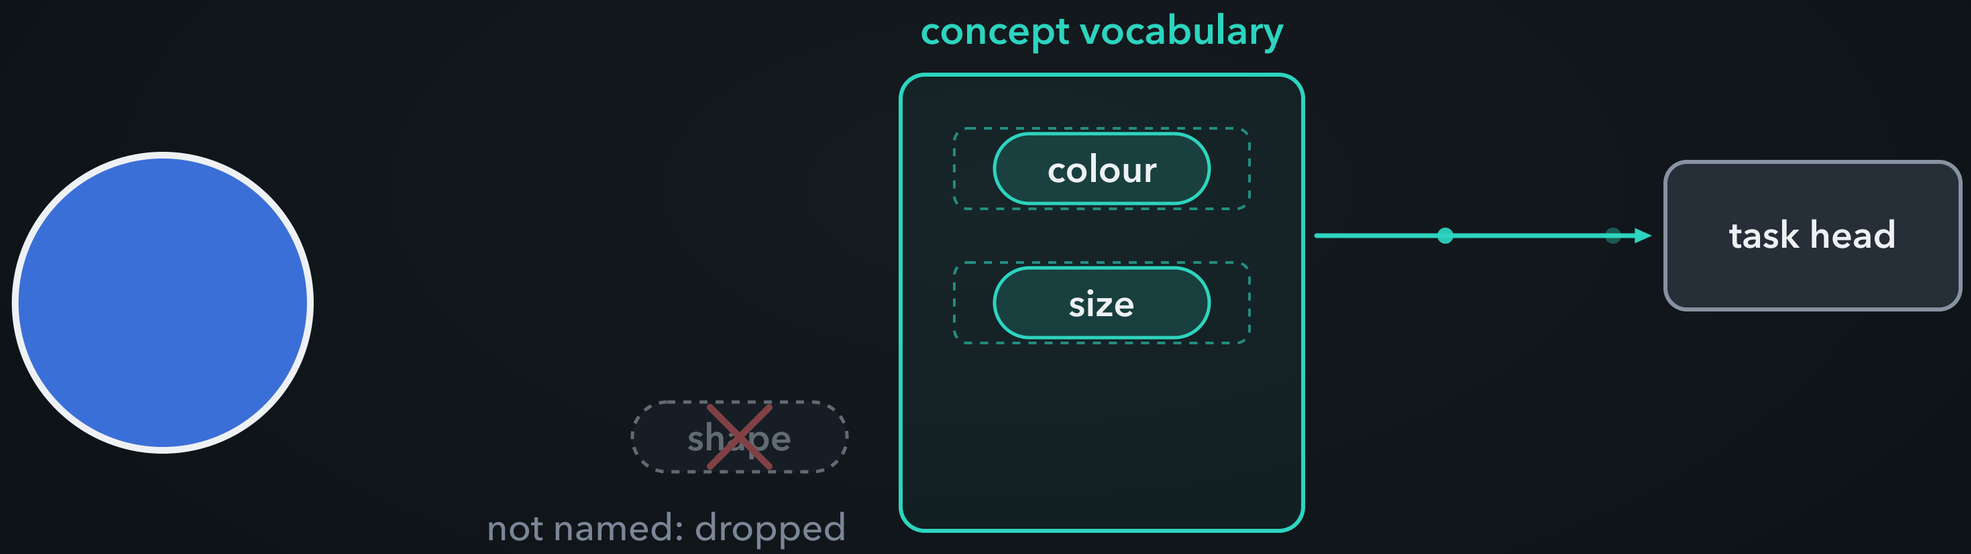

What if the concepts are incomplete?

Two objects from two different classes

fill: the object's own colour

CONSTRUCTED EXAMPLE · two-object example

What if the concepts are incomplete?

The vocabulary names only colour and size

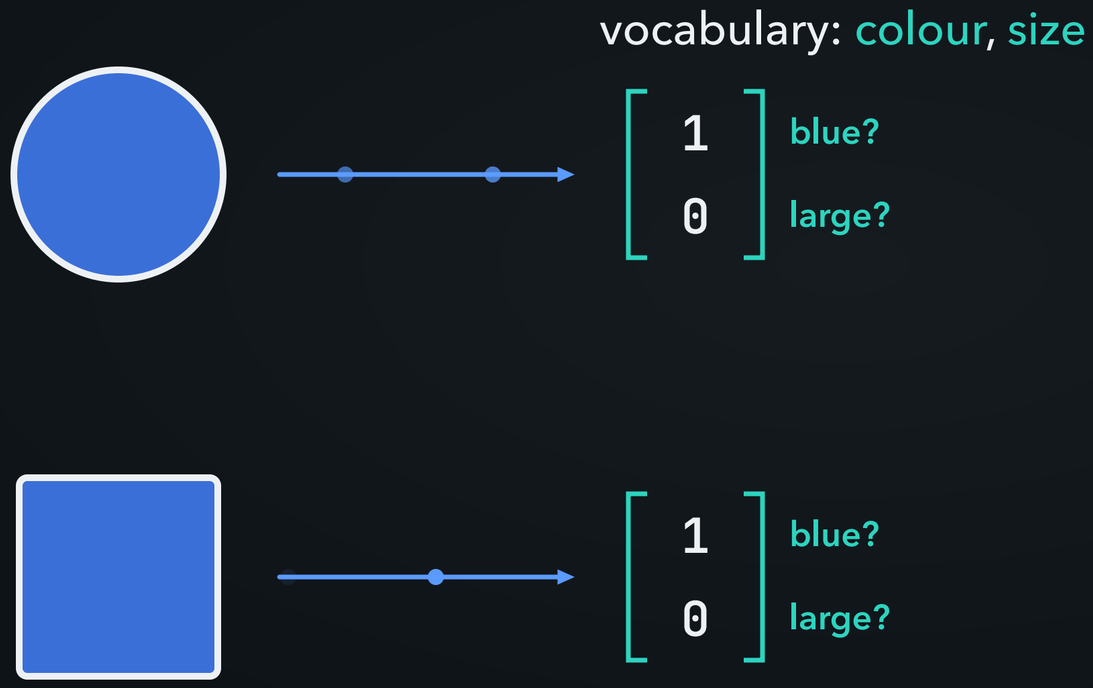

CONSTRUCTED EXAMPLE · two-object example

What if the concepts are incomplete?

Both objects get the same concept vector

$$c(x_1) = c(x_2), \quad y_1 \neq y_2$$

What if the concepts are incomplete?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Shape never reaches the task head

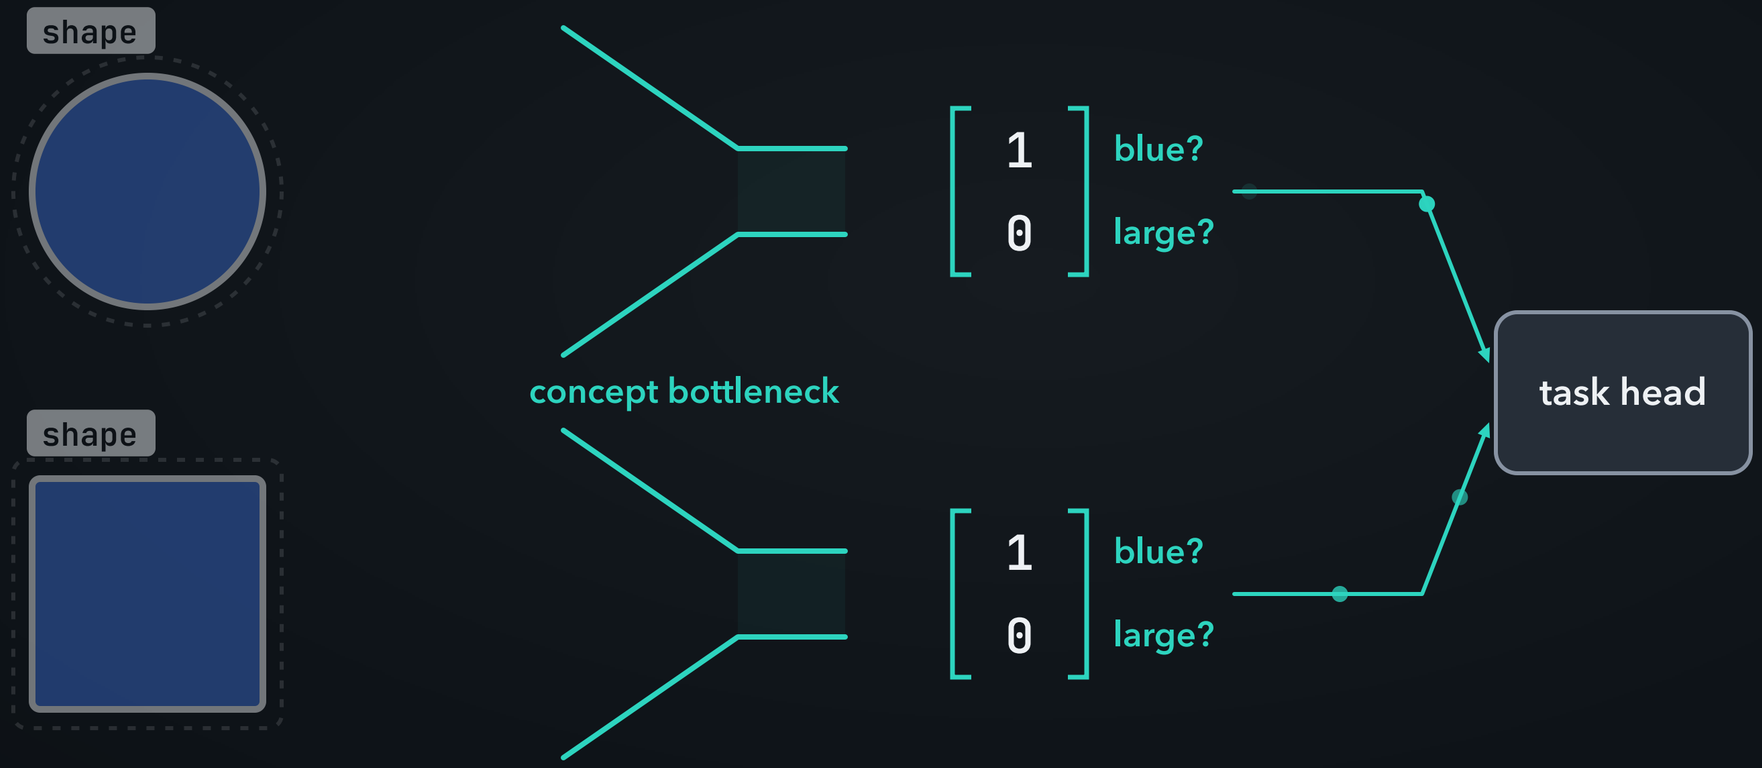

shape: in the image, not in the vector

What if the concepts are incomplete?

The same input to the head gives the same output

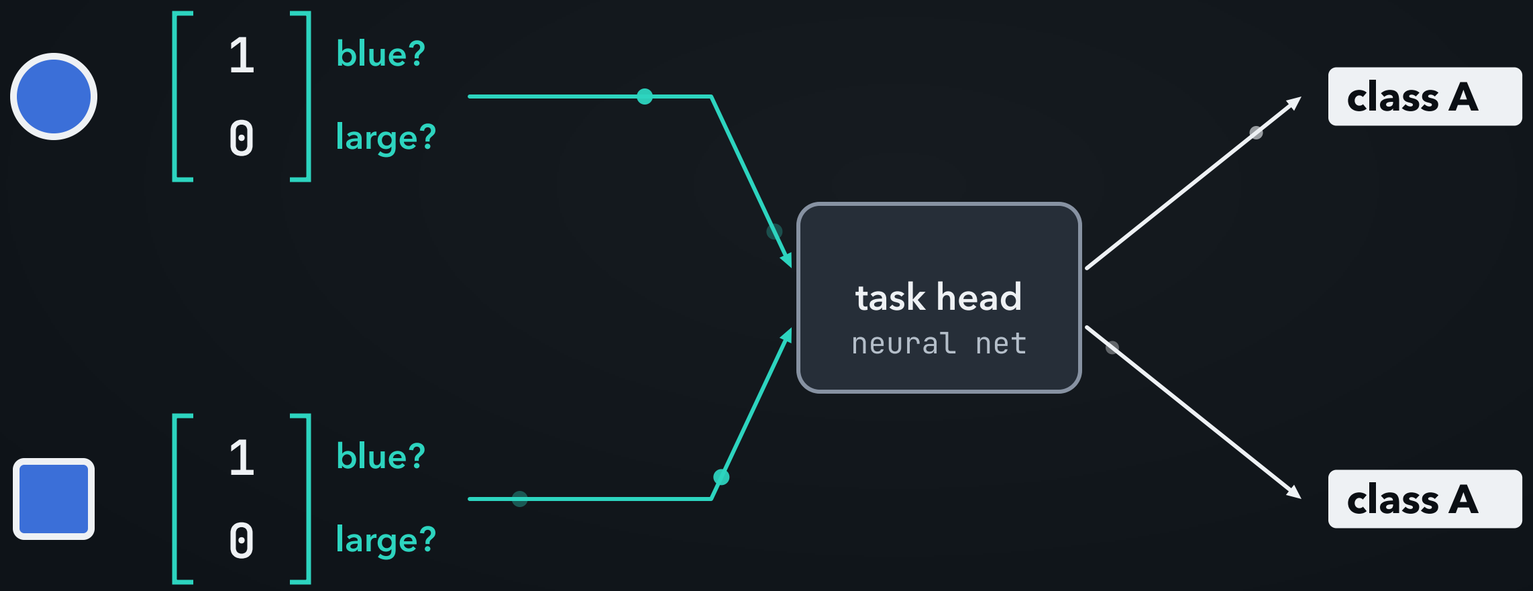

no head can recover what never reached it

What if the concepts are incomplete?

The best possible accuracy is one half

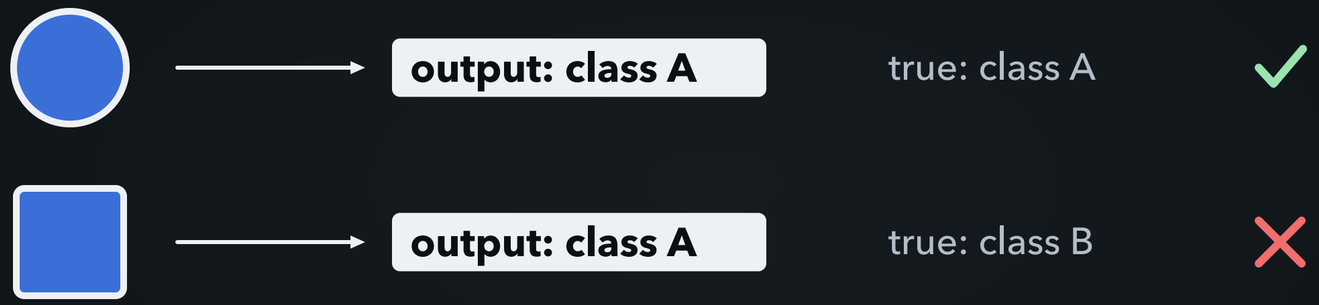

best possible accuracy

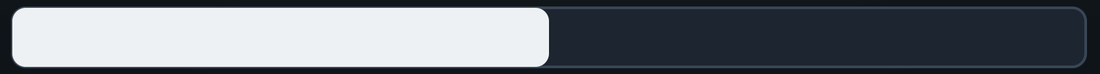

50%

CONSTRUCTED EXAMPLE · two-object example

What if the concepts are incomplete?

Sufficiency: the concepts say all the input says about the label

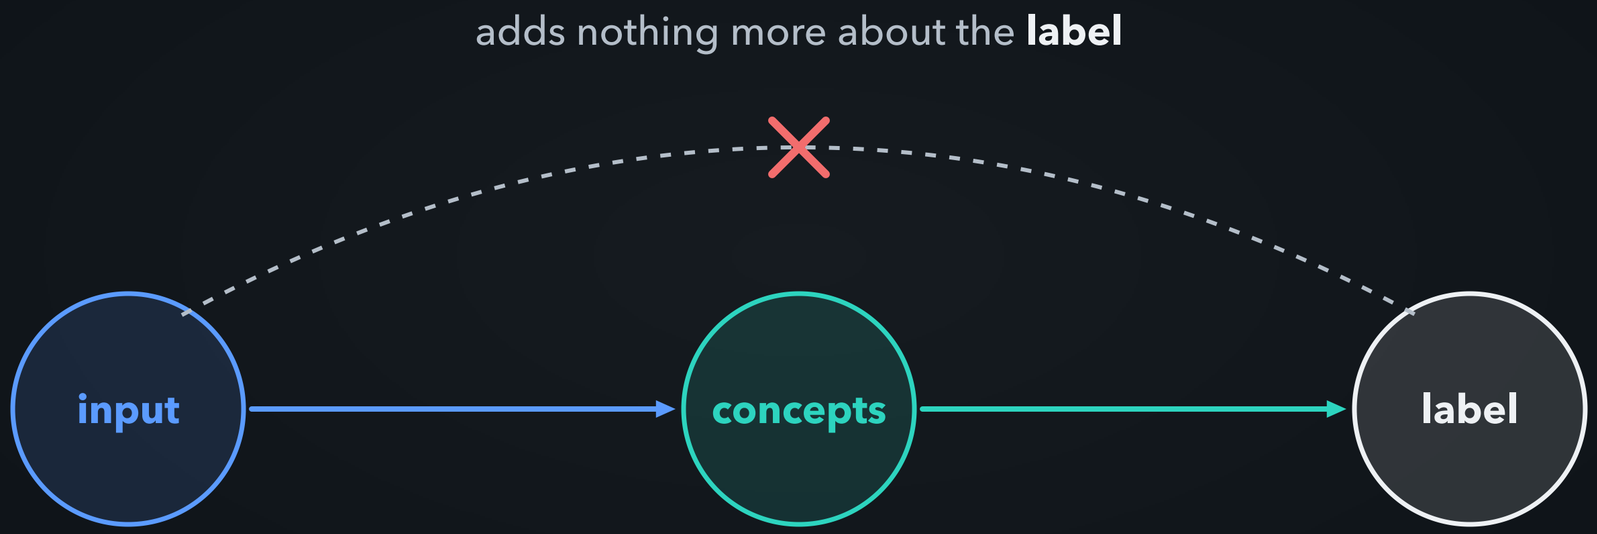

ideal: **sufficiency**

What if the concepts are incomplete?

Sufficiency in symbols

label independent of input given concepts

$$Y \not\perp\!\!\!\perp X \mid C$$

here it fails

the **label** depends on shape, which C lacks

What if the concepts are incomplete?

A strict CBM guarantees something weaker

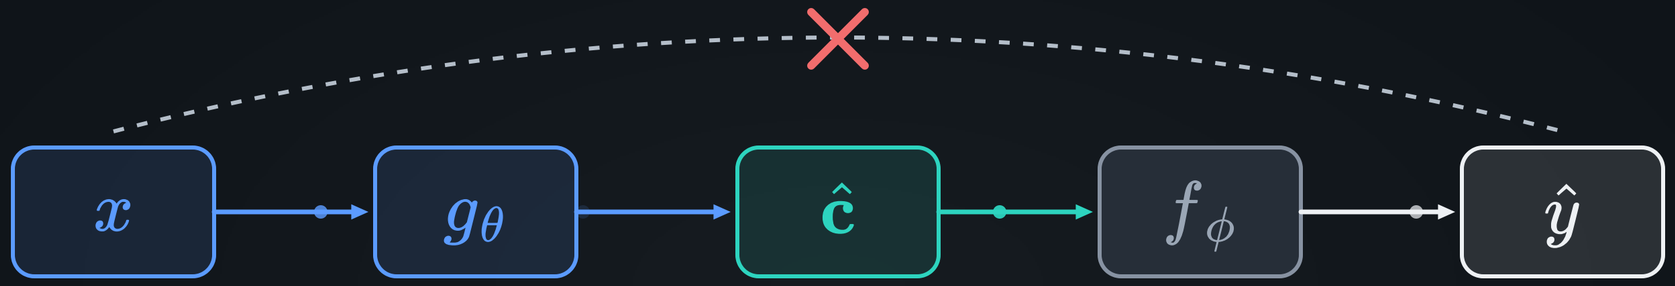

the **prediction** depends on the input only through the concepts

$$\hat{y} = f_\phi(\hat{\mathbf{c}}), \qquad \hat{\mathbf{c}} = g_\theta(x)$$

whether the true **label** does: a property of the vocabulary

What if the concepts are incomplete?

The fix is a richer vocabulary: add shape

vocabulary: colour, size, + shape

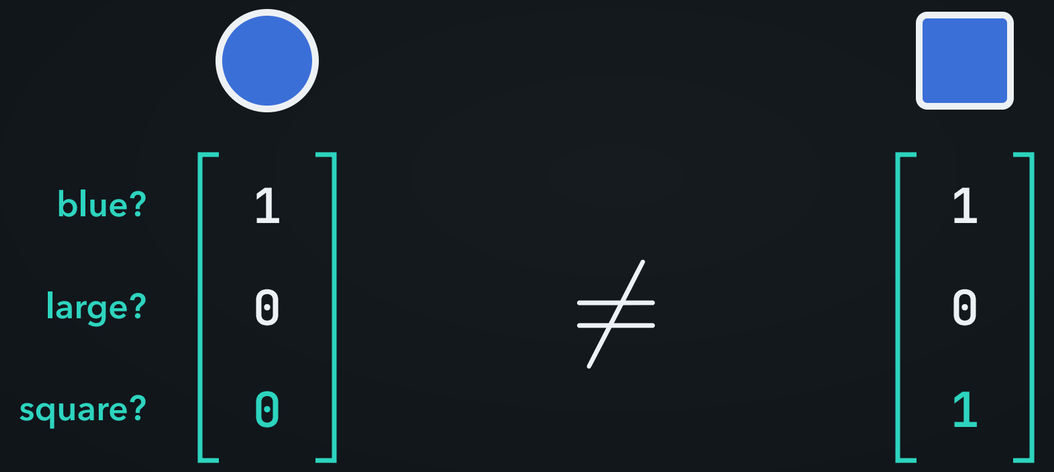

the two vectors now differ

CONSTRUCTED EXAMPLE · two-object example

PART 10

Modern concept bottleneck architectures

Follow the ladder of supervision

Modern concept bottleneck architectures

Labelling every concept on every image is the main cost

stripes	1	1	0	0	0	0
horse-like shape	1	1	1	1	0	0
long neck	0	0	0	0	1	1

36 Wikimedia Commons photographs (CC BY-SA / CC BY / public domain)

a person labels every concept on every image

the main cost, which newer methods try to reduce

Modern concept bottleneck architectures

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

One training record: an image, its concept labels and a class

Photo: Diego Delso, Wikimedia Commons, CC BY-SA 4.0

```
{
  "image": "zebra_12.jpg",
  "concepts": {
    "stripes": 1,
    "horse_like_shape": 1,
    "long_neck": 0
  },
  "label": "zebra"
}
```

CONSTRUCTED EXAMPLE · a teaching record

CUB birds, as used by Koh et al. 2020

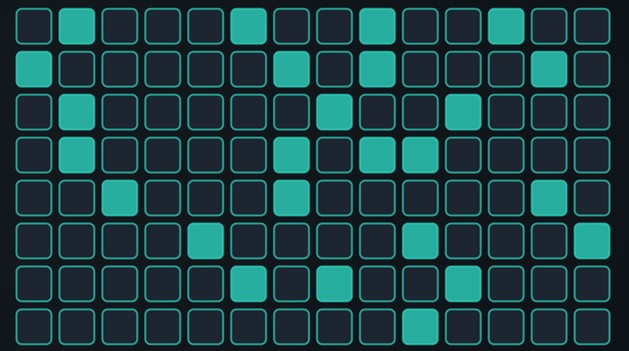

112 concept labels per photograph

DATASET · Herring Gull's labels in CUB

Modern concept bottleneck architectures

Rung one: images, concept labels and classes

rung 1 · task data

$$\{x,\ c,\ y\}$$

images, concept labels, **class labels**

CBM, concept embedding model (CEM)

rung 2

rung 3

less supervision

Modern concept bottleneck architectures

Rung two: keep class labels, drop per-image concept labels

rung 1 · task data

$\{x, c, y\}$

images, concept labels, **class labels**

CBM, concept embedding model (CEM)

drop

concept labels

rung 2 · task data

$\{x, y\}$

images and **class labels**

post-hoc CBM

rung 3

less supervision

Modern concept bottleneck architectures

Rung three: no training for this task

rung 1 · task data

$\{x, c, y\}$

images, concept labels, **class labels**

CBM, concept embedding model (CEM)

drop

concept labels

rung 2 · task data

$\{x, y\}$

images and **class labels**

post-hoc CBM

drop

task training

rung 3 · task data

$\varnothing$

no training for this task

classification by description

less supervision

Modern concept bottleneck architectures

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Rung two: class labels teach the task head

Example: CLIP phrases + a learned head

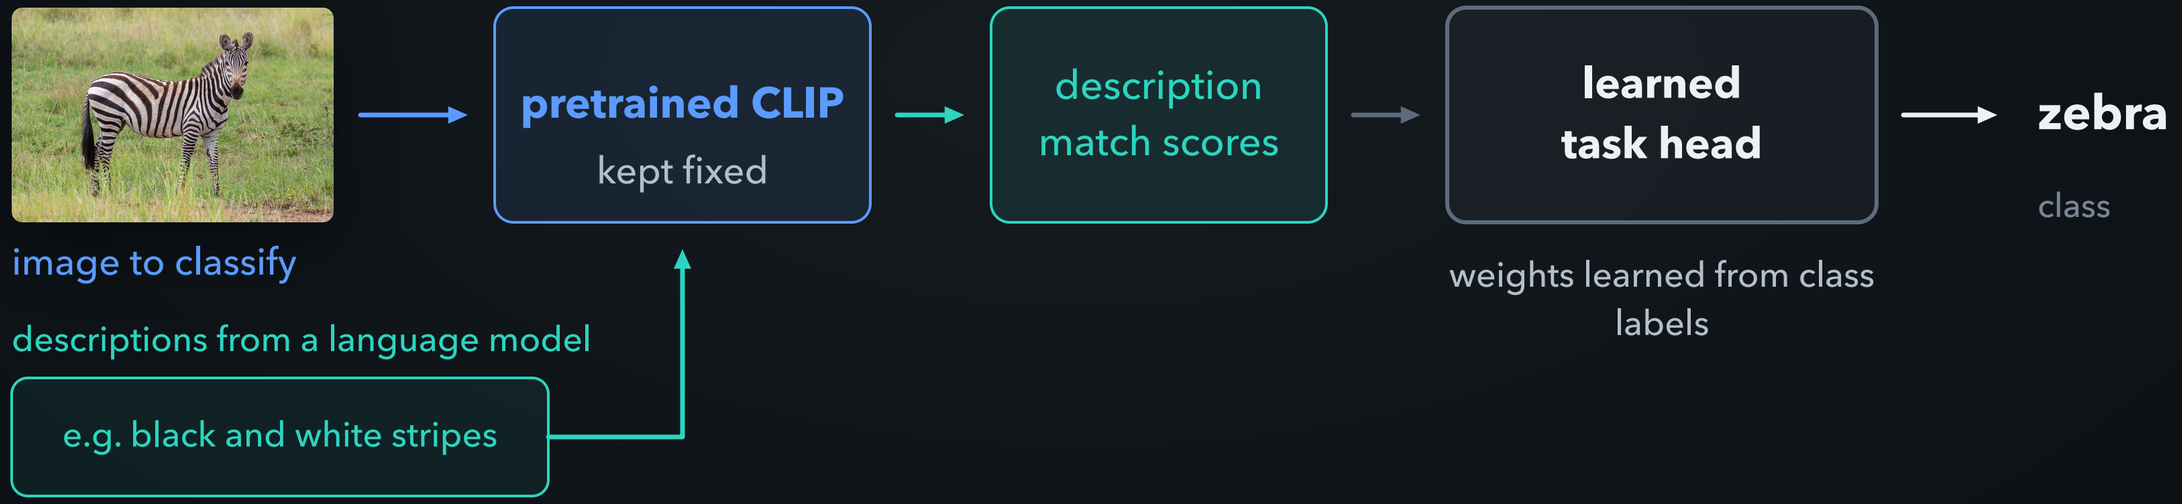

Rung two: **class labels** teach us how to combine the scores.

`SCHEMATIC · example path; no measured scores shown`

Photo: Diego Delso, Wikimedia Commons, CC BY-SA 4.0

Modern concept bottleneck architectures

Rung three: combine the match scores with a fixed rule

Classification by description

no task-labelled examples
no fitting

image to classify

pretrained CLIP
kept fixed

description
match scores

fixed
class averages

average the description
scores within each class

zebra

highest average

descriptions from a language model

e.g. black and white stripes

Rung three: combine the scores with a **fixed average**.

`SCHEMATIC · example path; no measured scores shown`

Photo: Diego Delso, Wikimedia Commons, CC BY-SA 4.0

Modern concept bottleneck architectures

Fewer concept labels, so check the concepts another way

	rung 1	rung 2	rung 3
concept labels for each task image		none	none
someone checks each concept against the photo		?	?

without per-image concept labels, check the concepts another way

Modern concept bottleneck architectures

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

For each variant, ask: is the concept really in the image?

Photo: Diego Delso, Wikimedia Commons, CC BY-SA 4.0

who checks that a named concept is really visible in the image?

we return to this question in the recap and evaluation

`black and white stripes`

Modern concept bottleneck architectures

A variant can change four things

what changes	plain CBM
concept source	people label every concept
detector	a trained concept predictor
representation	one number per concept
what reaches the head	only the concept values

Modern concept bottleneck architectures

One running example: photographs of zebras, horses and giraffes

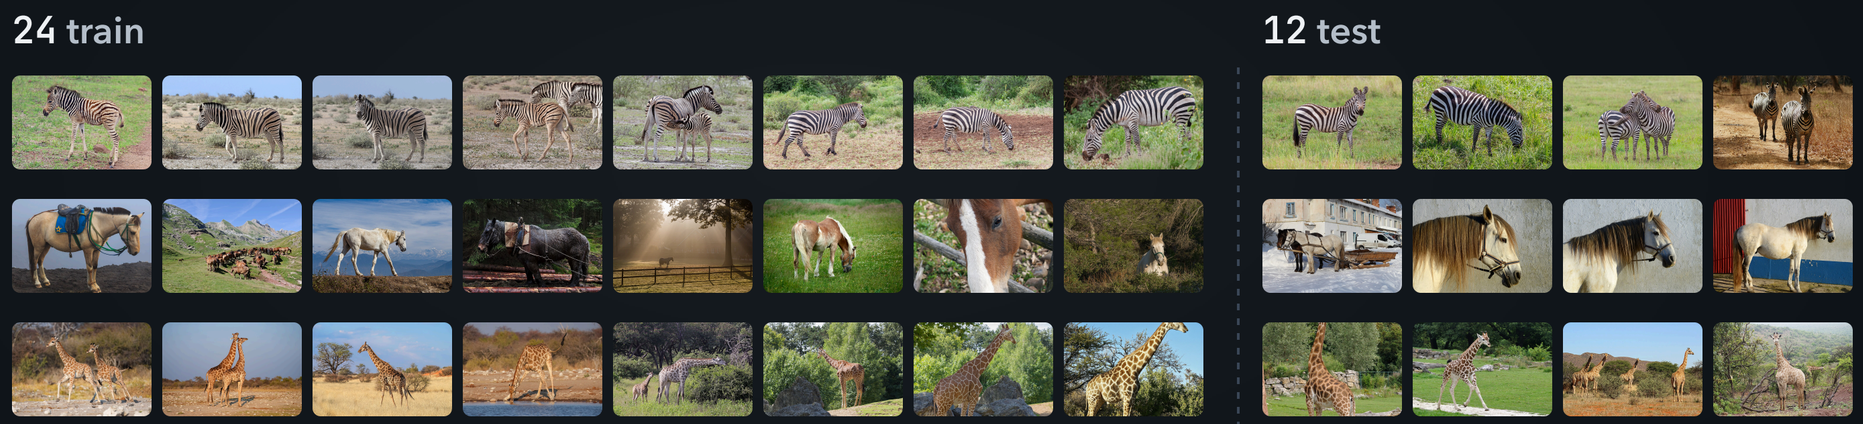

36 Wikimedia Commons photographs (CC BY-SA / CC BY / public domain)

the three classes from Part 4

Modern concept bottleneck architectures

For each variant: how the path from image to concepts to class changes

$$x \longrightarrow \hat{c} \longrightarrow \hat{y}$$

each variant, in symbols		
concept predictor	g_θ	· · ·
what reaches the head	$\hat{c}$	· · ·
task head	f_ϕ	· · ·
what is trained		· · ·
data needed		· · ·

PART 10 · RUNG ONE

Concept labels and class labels

Next: Concept Embedding Models (CEMs)

Training uses an image, its concept labels and its class label.

Concept and class labels | Class labels | No task fitting

Concept embedding models

One number per concept is a narrow channel

Concept Embedding Model (CEM)

Photo: Diego Delso, Wikimedia Commons, CC BY-SA 4.0

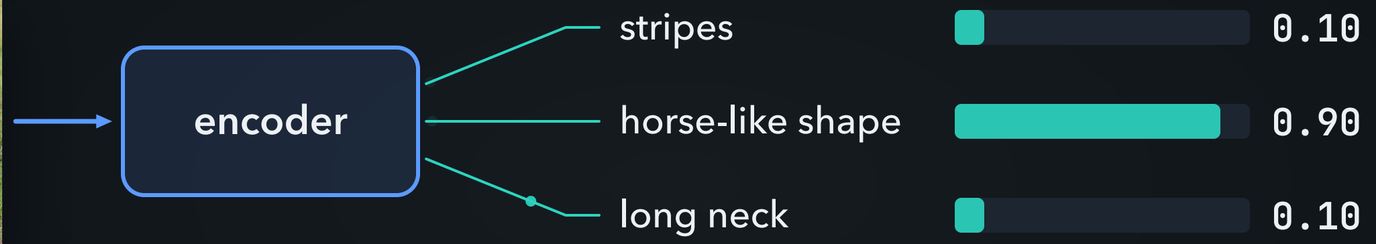

one number per concept: a narrow channel

Part 8: an incomplete vocabulary

`blue 1, large 0 for both`

same values, so best accuracy 0.50

ILLUSTRATIVE VALUES · hand-chosen values

Concept embedding models

What does a prototype mean here?

An embedding is a short list of learned numbers.

stripes present

$$\begin{bmatrix} 0.82 \\ 0.30 \end{bmatrix}$$

stripes absent

$$\begin{bmatrix} 0.22 \\ 0.74 \end{bmatrix}$$

Both vectors come from the current image.

These are not stored reference photos or class averages.

The coordinates are learned features, without individual names.

ILLUSTRATIVE VALUES · hand-chosen vector values

Concept embedding models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Learned layers turn the image into the two vectors

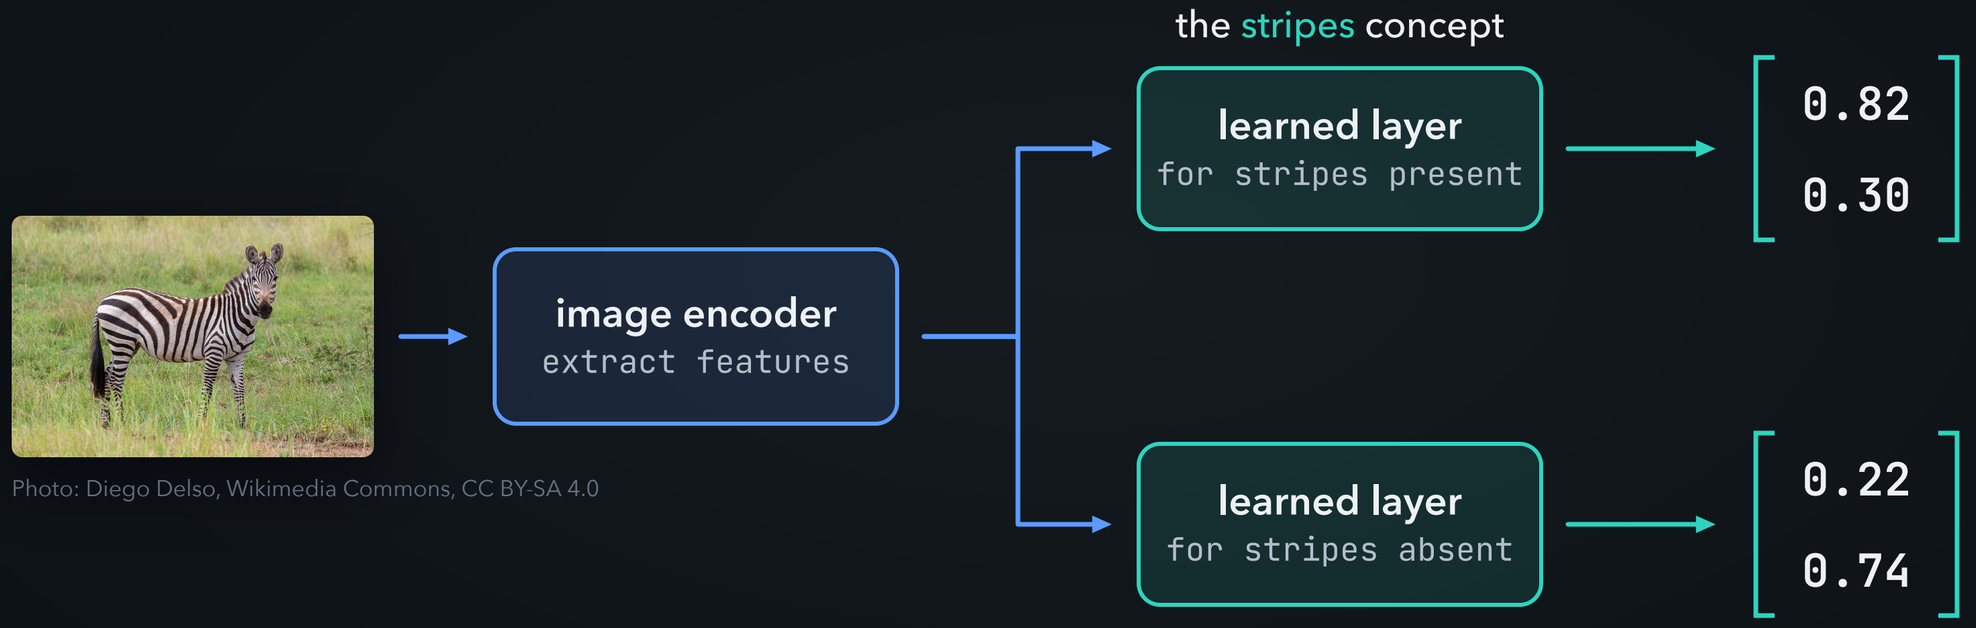

Photo: Diego Delso, Wikimedia Commons, CC BY-SA 4.0

Training learns the **weights** of these layers.

Keep the weights fixed: a new image can still produce different vectors.

ILLUSTRATIVE VALUES · hand-chosen outputs, not a trained run

Concept embedding models

Each concept gets two embeddings: active and inactive

encoder

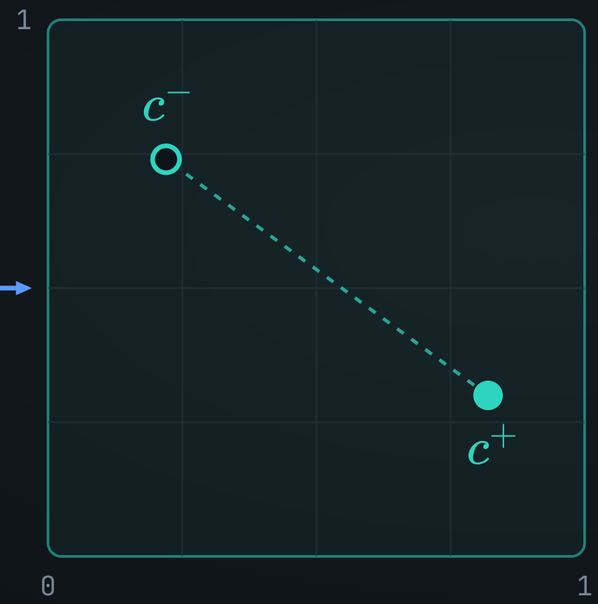

stripes: two embeddings in one plane

$$c^+ = \begin{bmatrix} 0.82 \\ 0.30 \end{bmatrix}$$ active

$$c^- = \begin{bmatrix} 0.22 \\ 0.74 \end{bmatrix}$$ inactive

one pair for every concept, computed from the image

embedding: a short list of numbers (here two)

ILLUSTRATIVE VALUES · hand-chosen 2-D embeddings

Concept embedding models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

A learned scorer predicts the probability p

encoder

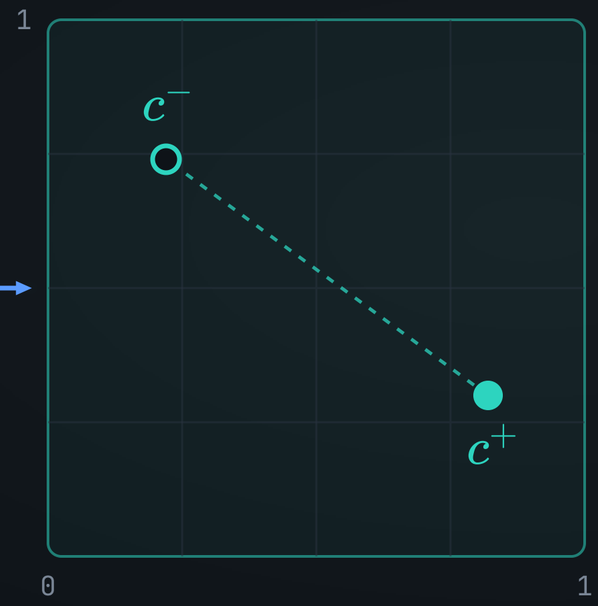

stripes: two embeddings in one plane

a learned scorer reads both vectors

and predicts the probability of stripes

$$p = 0.1$$

suppose it misses the stripes, as in Part 4

ILLUSTRATIVE VALUES · hand-chosen 2-D embeddings

Concept embedding models

The concept embedding mixes the two, weighted by p

encoder

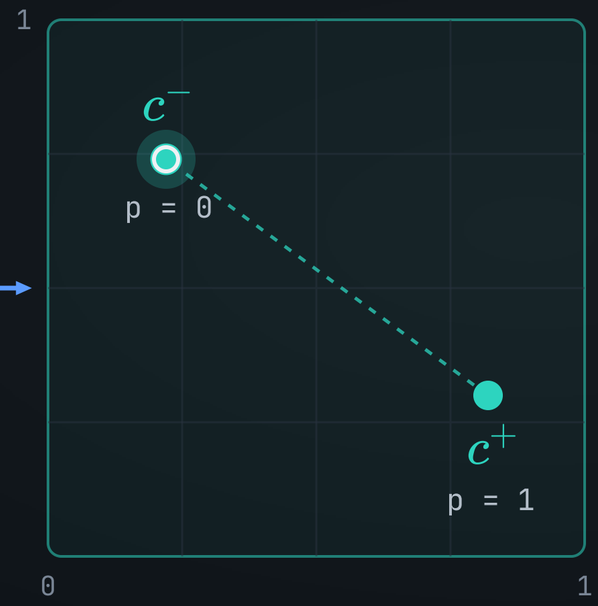

stripes: two embeddings in one plane

mixed = p × active + (1 − p) × inactive

$$\hat{c} = p\,c^+ + (1-p)\,c^-$$

the mixed point lies on the segment from inactive to active

ILLUSTRATIVE VALUES · hand-chosen 2-D embeddings

Concept embedding models

With p = 0.1, the mixed vector sits close to inactive

encoder

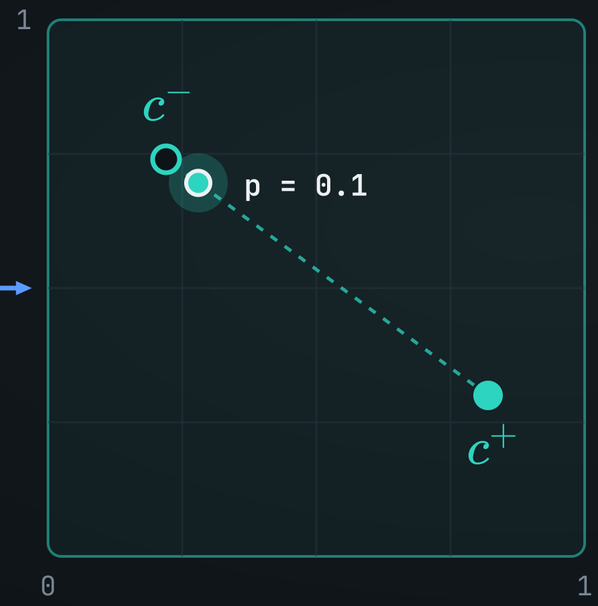

stripes: two embeddings in one plane

$\hat{c}$ for stripes, with p = 0.1:

first entry:

$$0.1 \times 0.82 + 0.9 \times 0.22 = 0.28$$

$$0.1 \begin{bmatrix} 0.82 \\ 0.30 \end{bmatrix} + 0.9 \begin{bmatrix} 0.22 \\ 0.74 \end{bmatrix} = \begin{bmatrix} 0.28 \\ 0.70 \end{bmatrix}$$

close to inactive

ILLUSTRATIVE VALUES · hand-chosen 2-D embeddings

Concept embedding models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

At test time, the head reads the mixed vectors

ILLUSTRATIVE VALUES · hand-chosen embeddings and head

stripes
p = 0.10
(0.280, 0.696)

horse-like shape
p = 0.90
(0.705, 0.440)

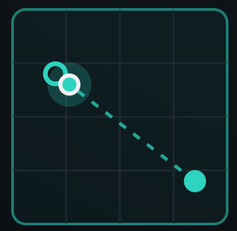

long neck
p = 0.10
(0.265, 0.650)

6 numbers

0.280
0.696
0.705
0.440
0.265
0.650

task head
then softmax

zebra 0.14
horse 0.72
giraffe 0.13

fixed weights; no true labels needed

Concept embedding models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

An optional correction changes p, with weights fixed

ILLUSTRATIVE VALUES · hand-chosen embeddings and head

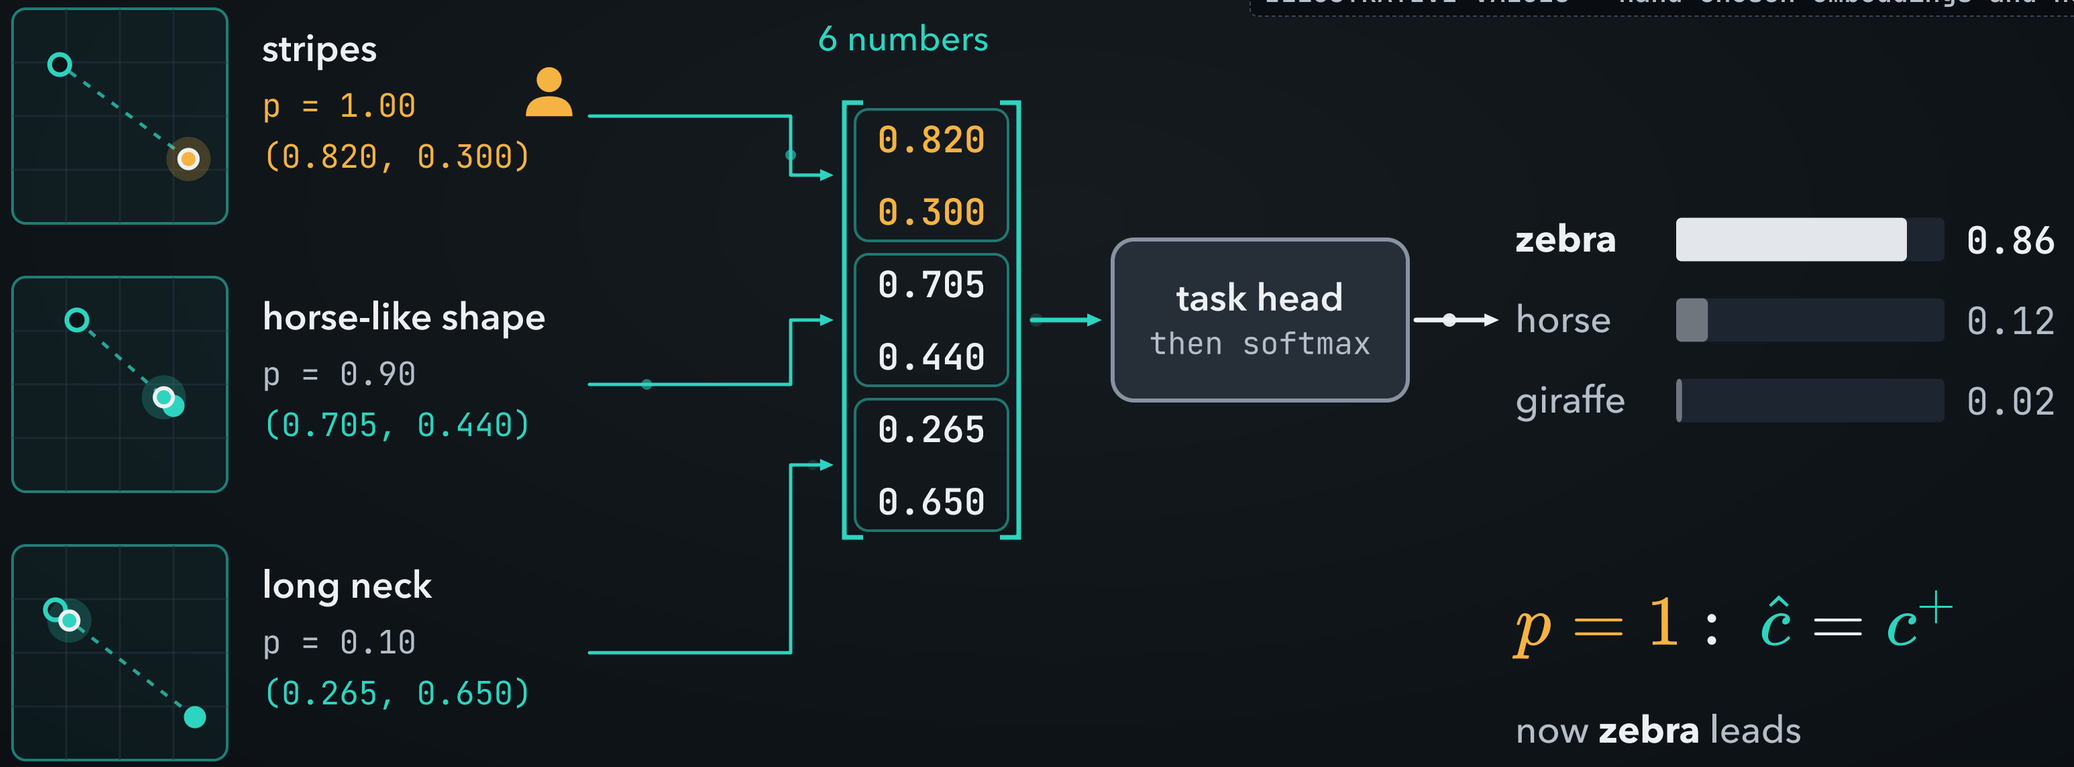

$p = 1 : \ \hat{c} = c^{+}$

now **zebra** leads

Concept embedding models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Training uses concept labels and class labels

Training example: photo + concept labels + **class label**

Photo: Diego Delso, Wikimedia Commons, CC BY-SA 4.0

encoder

vectors and p
learned layers + scorer

mixed vectors

task head
class prediction

class label: **zebra**

concept label: stripes present

concept loss
p versus true concept

task loss
compare with class: zebra

Backpropagation updates the encoder, vector layers, scorer and head.

Concept embedding models

Training learns the weights. Testing uses them

what changes	training	test
input	image, concept labels, class label	new image
network weights	updated using both losses	kept fixed
vectors and probabilities	computed from this image	computed from the new image

Test-time vectors are newly computed; the learned weights stay fixed.

Concept embedding models

Concept embedding models in symbols

Concept embedding models (Espinosa Zarlenga et al., NeurIPS 2022) x → ĉ → ŷ

concept predictor	$g_\theta : x \mapsto (p_j,\ c_j^+,\ c_j^-)$
what reaches the head	$\hat{c}_j = p_j\, c_j^+ + (1 - p_j)\, c_j^- \in \mathbb{R}^m$
task head	$\hat{y} = f_\phi(\hat{c}_1, \ldots, \hat{c}_k)$
what is trained	$g_\theta,\ f_\phi$ (joint training)
data needed	$\{(x,\ c,\ y)\}$ images, concept labels, class labels

j: one concept · k: number of concepts · m: numbers in each embedding

PART 10 · RUNG TWO

Use CLIP scores with a learned head

Post-hoc concept bottleneck: the CLIP phrase example

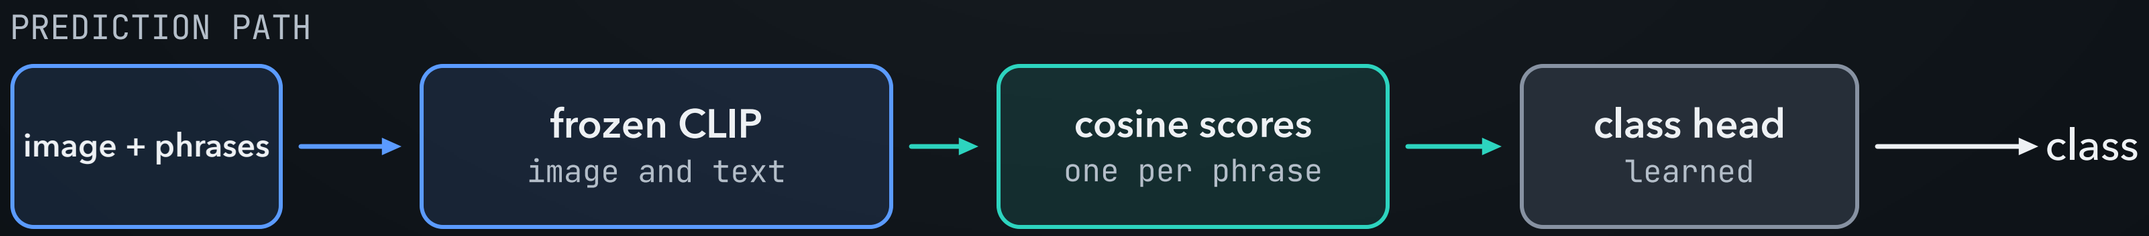

Fit the class head using **class labels**; keep CLIP fixed.

Next: follow the phrase bank, image scores and class prediction.

SCHEMATIC · predict the class through named concepts

Post-hoc concept bottleneck models

Start with animal classes; ask for visible concepts

Post-hoc Concept Bottleneck Model (PCBM): the CLIP phrase route

Animal classes: zebra, horse, giraffe

Ask a large language model (LLM)

"Suggest visible properties that distinguish these animals. Avoid class names."

→

black and white stripes

a long neck

a flowing mane

The LLM proposes words. CLIP will score the photos.

ILLUSTRATIVE VALUES · example prompt and possible suggestions; not a recorded LLM response

Post-hoc concept bottleneck models

Review the phrases, then fix one shared concept bank

Keep visible properties; reject class names, duplicates and weak matches.

“zebra” → reject: class name

“striped coat” → merge with “black and white stripes”

black and white stripes

a long neck

a flowing mane

brown patches on a pale coat

hooves

a saddle

green grass

a wooden fence

Every image is scored against this same list.

DATASET · saved bank: hand-chosen, not LLM-generated; review examples illustrative

Post-hoc concept bottleneck models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

CLIP compares a photo with a concept phrase

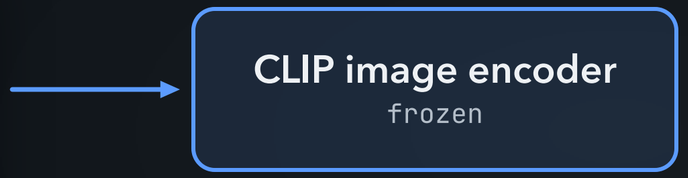

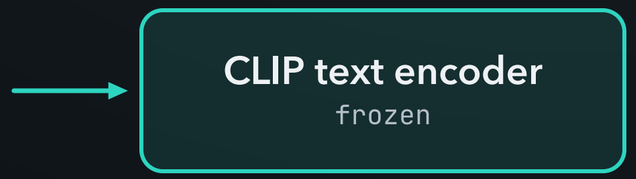

cosine(image vector, phrase vector) = 0.274

A match score, not a probability of stripes.

Photo: Diego Delso, Wikimedia Commons, CC BY-SA 4.0

COMPUTED · REAL MODEL · frozen CLIP ViT-B/32; saved photo and exact phrase

Post-hoc concept bottleneck models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

One image gives one score per concept

image encoder
512 features

cosine similarity
with each phrase

black and white stripes	0.274
long neck	0.243
flowing mane	0.234
brown patches on a pale coat	0.257
hooves	0.243
saddle	0.223
green grass	0.217
wooden fence	0.211

8 named scores; not probabilities

Next: standardise scores, then train the class head.

Photo: Diego Delso, Wikimedia Commons, CC BY-SA 4.0

COMPUTED · REAL MODEL · frozen CLIP ViT-B/32; saved phrase vectors

Post-hoc concept bottleneck models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Use the same concept prompts for every image

Three held-out photos

	raw cosine similarity: three prompts from the same bank		
	“black and white stripes”	**“a long neck”**	**“a flowing mane”**
zebra	0.274	0.243	0.234
horse	0.168	0.199	0.186
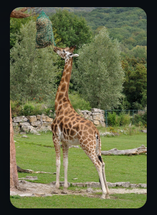 giraffe	0.208	0.301	0.218

Same encoder. Same prompts. One score per image and concept.

COMPUTED · REAL MODEL · frozen CLIP ViT-B/32; cosine scores, not probabilities

Commons photos: Diego Delso (CC BY-SA 4.0); George Chernilevsky (Public domain); Quartl (CC BY-SA 3.0)

Post-hoc concept bottleneck models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Each prompt gets its own training average and spread

Score all 24 training photos with the same prompts.

	“black and white stripes”	“a long neck”	“a flowing mane”
training average	0.215	0.239	0.227
training spread standard deviation	0.042	0.037	0.023

Calculate down each prompt column, across training images.

The three test photos do not set the scale.

COMPUTED · REAL MODEL · saved training split; each prompt has its own mean and spread

Post-hoc concept bottleneck models

Standardise each score: subtract the mean, divide by the spread

standardised score = (score − training mean) ÷ training spread

$$z = \frac{0.274 - 0.215}{0.042} \approx 1.41$$

COMPUTED · REAL MODEL · stripes scores, mean and spread over the 24 training photos

Post-hoc concept bottleneck models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

The same photos and prompts, now on a common scale

Three held-out photos

raw cosine → standardised score

	“black and white stripes”	“a long neck”	“a flowing mane”
zebra	0.274 → **+1.41**	0.243 → **+0.10**	0.234 → **+0.32**
horse	0.168 → **-1.10**	0.199 → **-1.08**	0.186 → **-1.80**
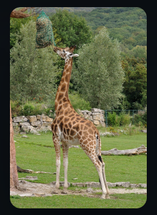 giraffe	0.208 → **-0.16**	0.301 → **+1.70**	0.218 → **-0.38**

0 = training average · positive = above · negative = below

Below average does not mean absent.

COMPUTED · REAL MODEL · training statistics reused; values rounded for display

Commons photos: Diego Delso (CC BY-SA 4.0); George Chernilevsky (Public domain); Quartl (CC BY-SA 3.0)

Post-hoc concept bottleneck models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

For this zebra, stripes scores highest

each bar: (score − training mean) ÷ training spread

stripes: 1.41 spreads above its average

COMPUTED · REAL MODEL · CLIP ViT-B/32 on 36 Commons photos, teaching example

Post-hoc concept bottleneck models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Train the task head using concept scores and class labels

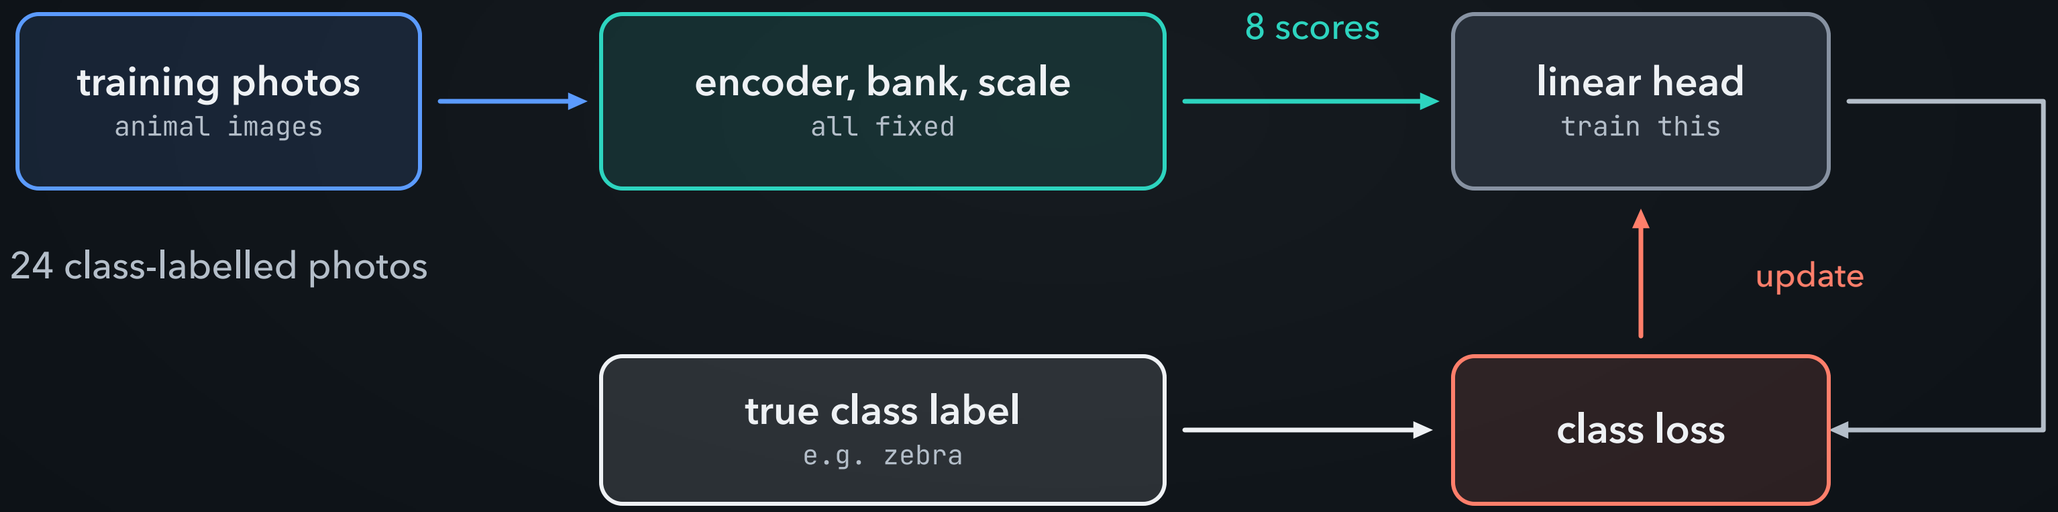

Class labels teach the head which concepts predict each animal.

SCHEMATIC · animal example: CLIP phrase bank, then a sparse task head

Post-hoc concept bottleneck models

Sparse: only 4 of 24 weights are non-zero

	zebra	horse	giraffe
black and white stripes	+1.49	0	0
long neck	0	-0.27	+1.02
flowing mane	0	+0.28	0
brown patches on a pale coat	0	0	0
hooves	0	0	0
saddle	0	0	0
green grass	0	0	0
wooden fence	0	0	0

8 concepts × 3 classes = 24 weights

4 / 24

weights are non-zero

L1 penalty: the size of every weight is added to the loss

COMPUTED · REAL MODEL · L1 logistic head fit on 24 photos

Post-hoc concept bottleneck models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

The zebra score: stripes score times weight, plus a constant

stripes score × weight + constant = **zebra score**

$$1.406 \times 1.486 - 0.054 \approx 2.035$$

class scores

shares, after the softmax

COMPUTED · REAL MODEL · CLIP ViT-B/32 on 36 Commons photos

Post-hoc concept bottleneck models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

At test time: one new image, with all weights fixed

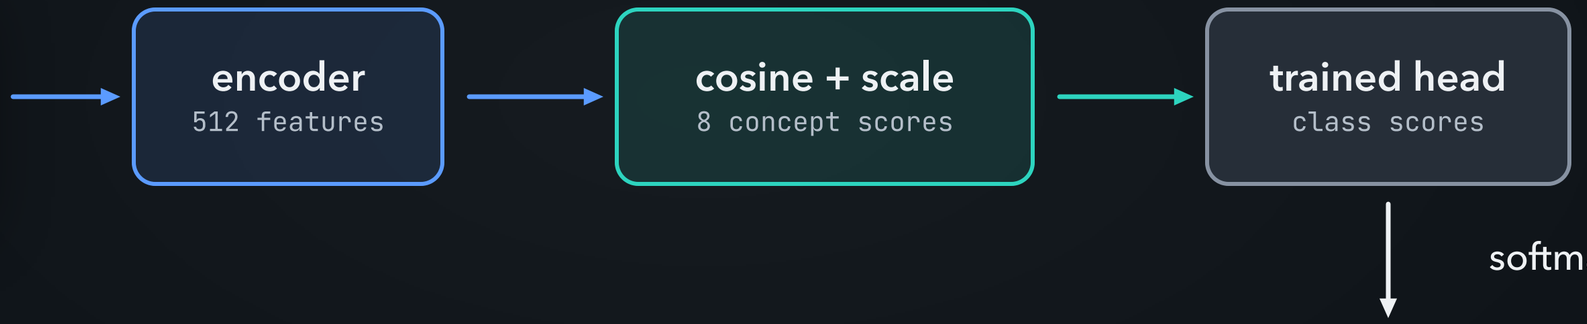

Input: one image.

No concept labels. No class label.

Everything is saved. Nothing is fitted again.

Photo: Diego Delso, Wikimedia Commons, CC BY-SA 4.0

COMPUTED · REAL MODEL · same held-out zebra; fixed encoder, bank, scale and head

Post-hoc concept bottleneck models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

All 12 held-out photos are classified correctly

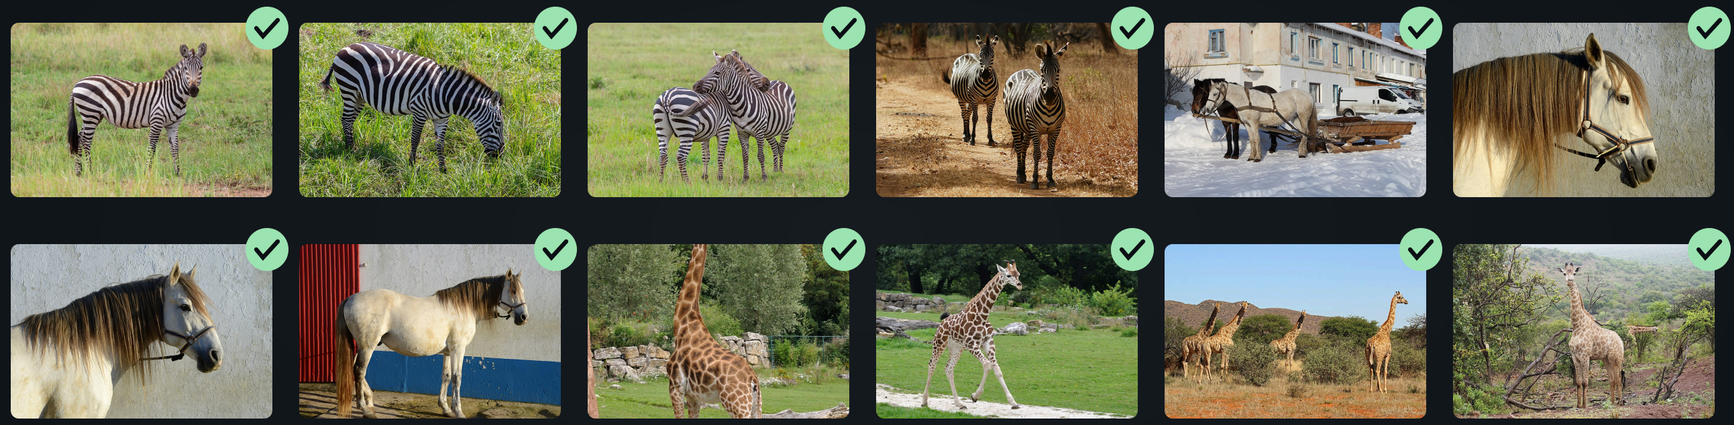

36 Wikimedia Commons photographs (CC BY-SA / CC BY / public domain)

12 of 12 held-out photos correct

a tiny teaching example, not a benchmark

Post-hoc concept bottleneck models

Edit the stripes score, and zebra falls

set to −0.66, its mean on non-zebra training photos

$$-0.659 \times 1.486 - 0.054 \approx -1.034$$

the new zebra logit; the photo and the head stay the same

COMPUTED · REAL MODEL · the same head, one score edited

Post-hoc concept bottleneck models

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

PCBM-h adds an extra path for evidence the concepts miss

Hybrid Post-hoc Concept Bottleneck Model (PCBM-h)

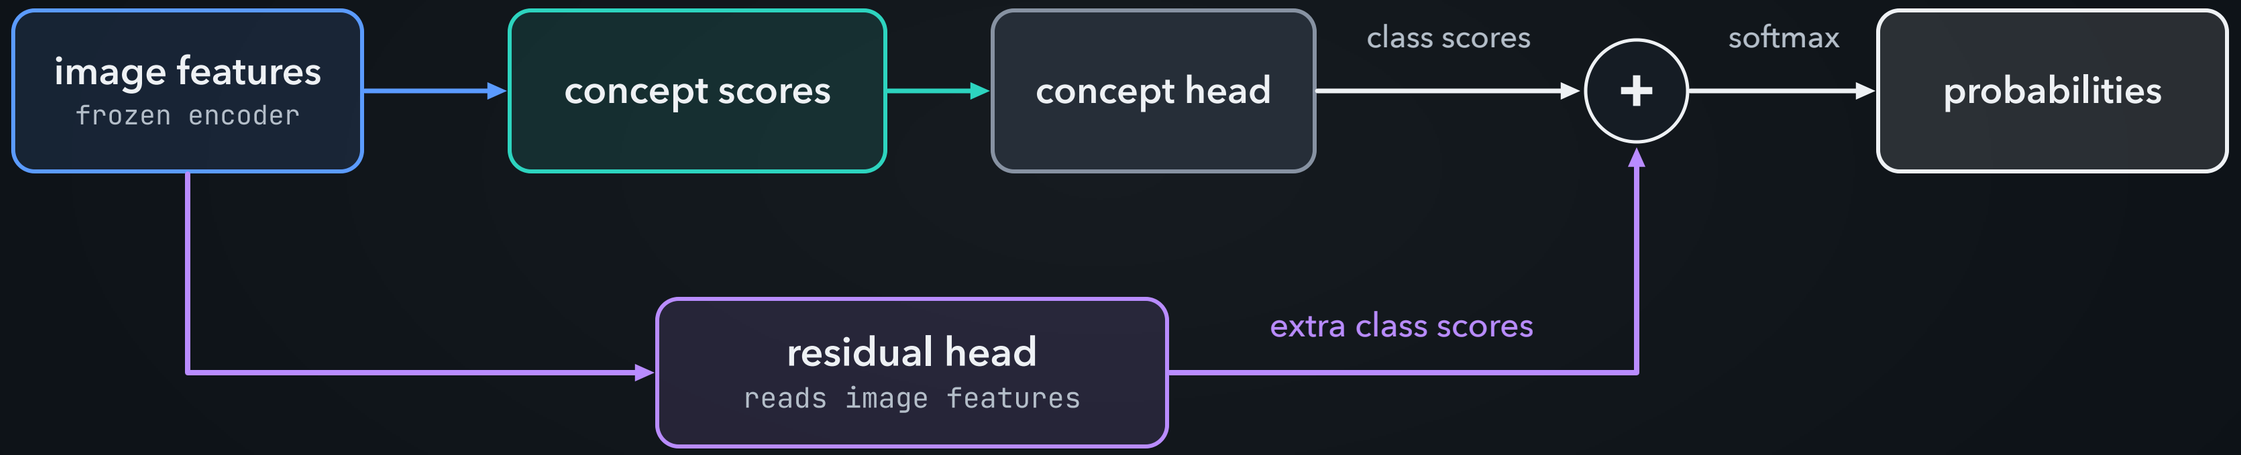

Worked photo: zebra class score, before softmax

$$2.03 + 0.31 \approx \mathbf{2.34}$$

Train the extra head with class labels; keep the concept model fixed.

COMPUTED · REAL MODEL · saved CLIP animal example; scores rounded for display

Post-hoc concept bottleneck models

x g_θ $\hat{c}$ f_ϕ $\hat{y}$

An edit leaves the extra path unchanged

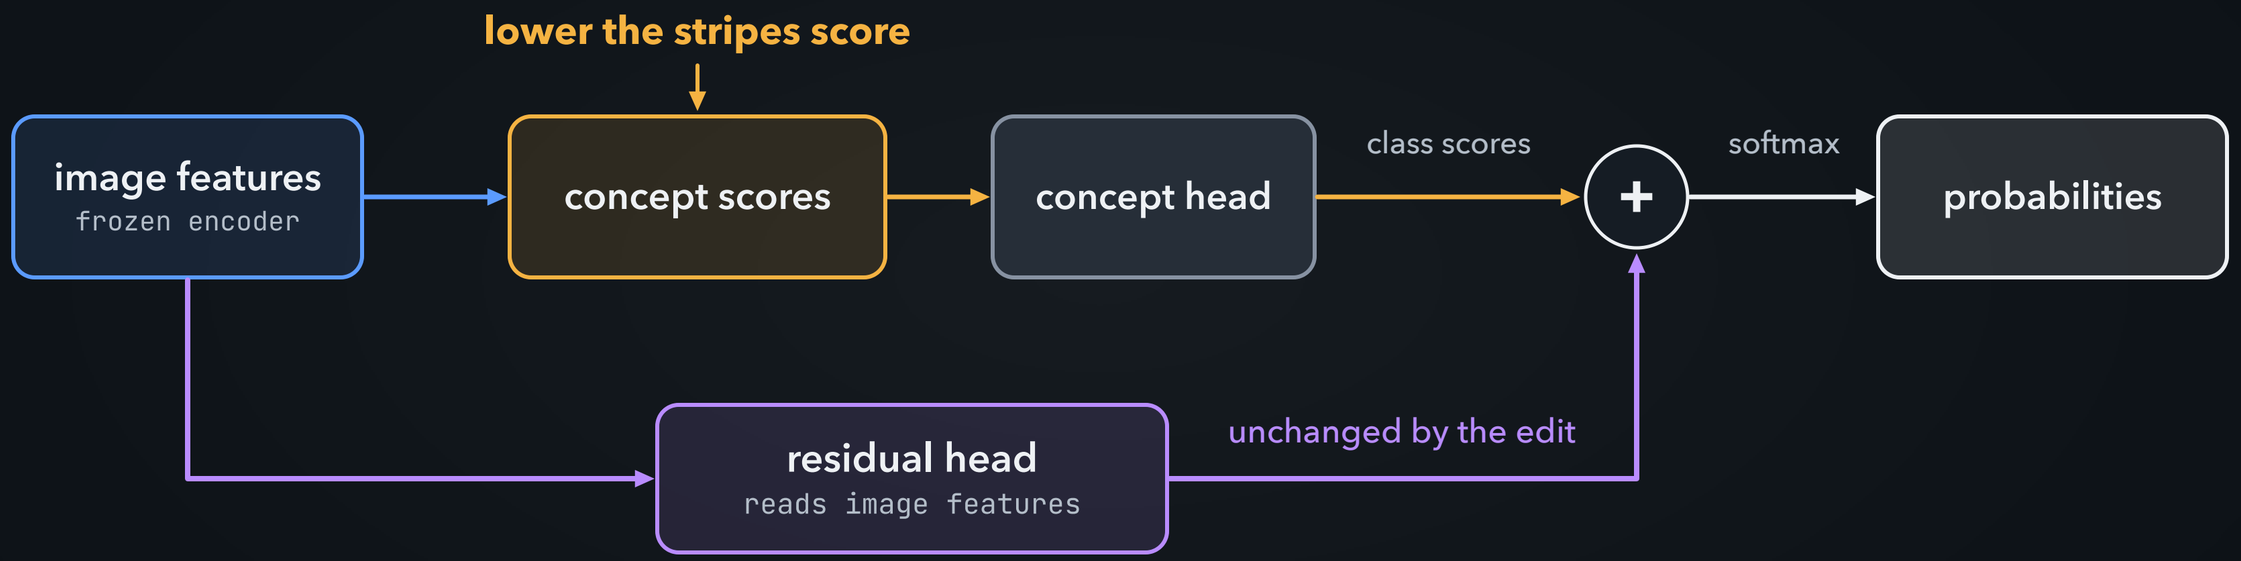

Zebra probability after the same stripes edit

concept model: **0.14**

with residual: **0.20**

The prediction can change, but some evidence bypasses the concepts.

COMPUTED · REAL MODEL · same held-out zebra photo; same edit; residual contribution fixed

Post-hoc concept bottleneck models

Compared with the plain CBM: the concepts come from phrases

what changes	plain CBM	post-hoc CBM
concept source	people label every concept	**a concept bank of phrases**
detector	a trained concept predictor	**frozen CLIP + cosine similarity**
representation	one number per concept	one number per concept
what reaches the head	only the concept values	**concept scores, plus a residual in PCBM-h**

3 of 4 changed; each concept is still one number

Post-hoc concept bottleneck models

Post-hoc concept bottleneck models in symbols

Post-hoc CBM (Yuksekgonul et al., ICLR 2023) x → ĉ → ŷ

concept predictor	$\hat{c}_j = \text{standardise}(\cos(e(x), v_j))$
what reaches the head	$\hat{c} = (\hat{c}_1, \ldots, \hat{c}_k)$ standardised cosine scores
task head	$\hat{y} = f_\phi(\hat{c})$ sparse linear $+ r(e(x))$ PCBM-h
what is trained	f_ϕ only (and r in PCBM-h)
data needed	$\{(x, y)\}$ + concept bank

e: image vector · v_j: phrase vector · r: residual

PART 10 · RUNG THREE

No fitting on the new task

Next: Classification by Description (CBD)

The pretrained model stays fixed; a fixed rule combines description scores.

Concept and class labels

Class labels

No task fitting

Classification by description

Rung three: no fitting on the new task

DATA EACH VARIANT NEEDS

$\{x, c, y\}$ CBM, CEM → $\{x, y\}$ PCBM → ∅ **by description**

Classification by description (CBD)

Menon and Vondrick, ICLR 2023

no task concept labels, no task **class labels**, no fitting

Classification by description

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Each class gets a checklist, and CLIP ticks it softly

photo

Herring Gull

☑ a white body with black wingtips	0.9
☑ a yellow beak and legs	0.8
☑ a red spot on its beak	0.7

average **0.80**

Ivory Gull

☑ a white bird	0.8
☐ a black bill	0.2
☐ pink legs	0.3

average **0.43**

Ring-billed Gull

☐ red eyes	0.2
☑ webbed feet	0.6
☑ a long, black tail	0.4

average **0.40**

best average: **wins**

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

ILLUSTRATIVE VALUES · real descriptors; tick strengths drawn for teaching

Classification by description

A language model describes every class

language model
writes descriptors

```
{ "Herring Gull": [
    "a large bird with a wingspan of up to 2.3 meters",
    "a white body with black wingtips",
    "a yellow beak and legs",
    "a red spot on its beak",
    "webbed feet"
]}
```

descriptor: a short phrase about what the **class** looks like

DATASET · Menon and Vondrick's released CUB descriptors

Classification by description

CLIP scores the photo against each descriptor

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

each score: a soft tick

- a large bird with a wingspan of up to 2.3 meters
- a white body with black wingtips
- a yellow beak and legs
- a red spot on its beak
- webbed feet

class score: the average

The **highest average** wins.

Classification by description

One descriptor score: a cosine between two vectors

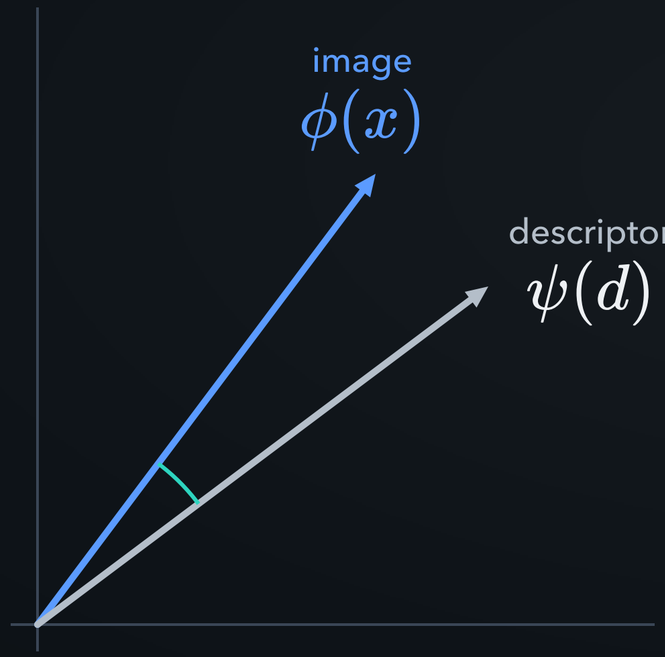

entries (0.6, 0.8) and (0.8, 0.6)

$$\hat{c}_d = \cos\left(\phi(x),\ \psi(d)\right)$$
$$= 0.6 \times 0.8 + 0.8 \times 0.6 = 0.96$$

both vectors have length 1, so the cosine is the sum of products

near 1: they point the same way; smaller: they point apart

ILLUSTRATIVE VALUES · two-entry vectors for teaching; CLIP's have 512 entries

Classification by description

The class score averages its descriptor scores

class score — average over the class's descriptors — descriptor score

$$\hat{y}_k = \frac{1}{|D_k|} \sum_{d \in D_k} \hat{c}_d$$

D_k: the list of descriptors of class k, and $|D_k|$ how many there are

Nothing is fitted to the target data.

Classification by description

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Every score comes with its reasons

bars start at 0.30

$$\hat{y}_{\text{Herring}} = \tfrac{1}{5}\big(0.324 + 0.339 + 0.339 + 0.340 + 0.328\big) = 0.334$$

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

COMPUTED · REAL MODEL · CLIP ViT-B/32 on our gull photo, released CUB descriptors

Classification by description

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Slaty-backed gull edges ahead

class scores: **Slaty-backed 0.337** · Herring 0.334 · Glaucous-winged 0.328 · Western 0.326 · California 0.324

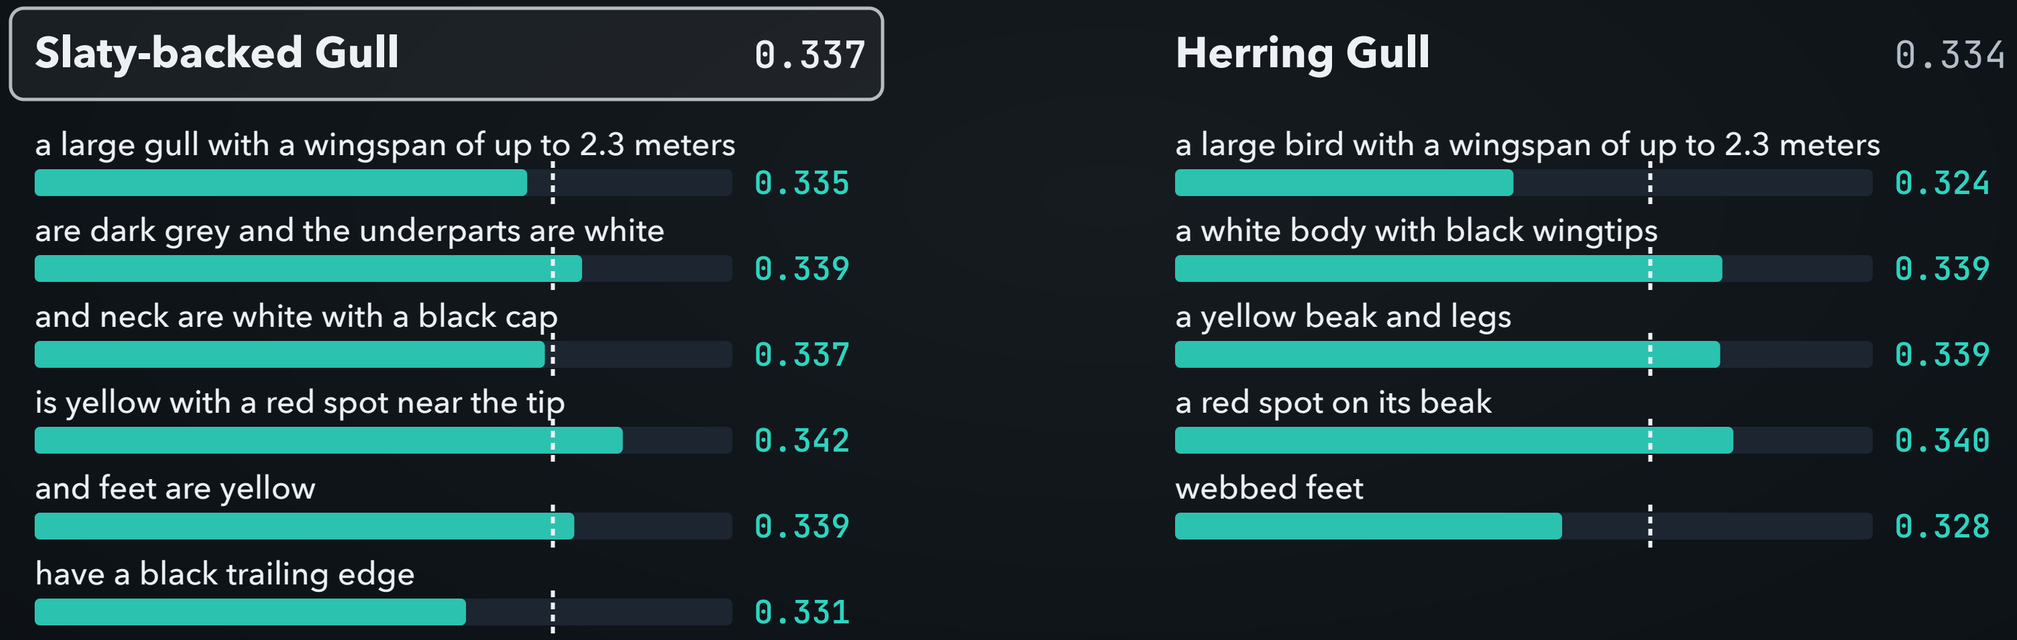

Slaty-backed Gull edges ahead by 0.003

bars start at 0.30; dashed line: the average

Our photo shows a **herring gull**: a near miss.

COMPUTED · REAL MODEL · CLIP ViT-B/32 on our gull photo, released CUB descriptors

Classification by description

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Attractive when there are no labels at all

no **labelled images**

classes we already score: Slaty-backed, Herring, Glaucous-winged, …

+ one more class

```
{ "Ivory Gull": [
    "a white bird",
    "black wingtips",
    "a black bill",
    "pink legs",
    "yellow eyes"
]}
```

its checklist is scored at once

no retraining

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

DATASET · Menon and Vondrick's released CUB descriptors

Classification by description

A descriptor can be wrong

The language model says:

```
{ "California Gull": [
    "a white bird with a black head and bill",
    ...
]}
```

It does not have a black head.

The wrong descriptor is still averaged into the score, one of 7.

DATASET · Menon and Vondrick's released CUB descriptors

Classification by description

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

A match score is not a fact check

Our herring gull photo

White head, yellow bill

From the California Gull checklist:

"a white bird with a black head and bill"

These words do not match this photo.

CLIP

Match score

0.325

Similarity, not a probability

California Gull

Class average

still includes this score

Check the words

Correct and visible in the photo?

Check the score

Does changing that feature change its score?

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

COMPUTED · REAL MODEL · CLIP ViT-B/32 on this photo; released CUB descriptor

Classification by description

From rung two to rung three: the head is no longer learned

	rung two: learned head		by description
the task head	a linear layer	→	**a fixed average**
what is trained	the head	→	**nothing**
data needed	images and **class labels**	→	**none**

Classification by description

Classification by description in symbols

Classification by description (Menon and Vondrick, ICLR 2023) x → ĉ → ŷ

concept predictor	$\text{CLIP: image encoder } \phi,\ \text{text encoder } \psi$
what reaches the head	$\hat{c}_d = \cos\big(\phi(x),\ \psi(d)\big),\quad d \in D_k$
task head	$\hat{y}_k = \frac{1}{\lvert D_k \rvert} \sum_{d \in D_k} \hat{c}_d \quad \text{(fixed)}$
what is trained	nothing
data needed	none: a language model writes the descriptors D_k

The cost: nothing guarantees that the descriptors are correct, or grounded.

Comparing the variants

One table: one row per variant, one column at a time

variant	data needed	concepts from	what the head reads	trained
CBM Koh et al. 2020				
CEM Espinosa Zarlenga et al. 2022				
PCBM Yuksekgonul et al. 2023				
CBD (by description) Menon and Vondrick 2023				

Comparing the variants

Column one: the supervision each variant needs

key: x an image, c its concept labels, y its class label

variant	data needed	concepts from	what the head reads	trained
CBM Koh et al. 2020	$\{(x, c, y)\}$			
CEM Espinosa Zarlenga et al. 2022	$\{(x, c, y)\}$			
PCBM Yuksekgonul et al. 2023				
CBD (by description) Menon and Vondrick 2023				

Comparing the variants

Higher rungs drop the concept labels

key: x an image, c its concept labels, y its class label

variant	data needed	concepts from	what the head reads	trained
CBM Koh et al. 2020	$\{(x, c, y)\}$			
CEM Espinosa Zarlenga et al. 2022	$\{(x, c, y)\}$			
PCBM Yuksekgonul et al. 2023	$\{(x, y)\}$ + a concept bank			
CBD (by description) Menon and Vondrick 2023	none for new task			

Comparing the variants

Column two: where the concepts come from

key: x an image, c its concept labels, y its class label

variant	data needed	concepts from	what the head reads	trained
CBM Koh et al. 2020	$\{(x, c, y)\}$	human labels		
CEM Espinosa Zarlenga et al. 2022	$\{(x, c, y)\}$	human labels		
PCBM Yuksekgonul et al. 2023	$\{(x, y)\}$ + a concept bank	a concept bank		
CBD (by description) Menon and Vondrick 2023	none for new task	a language model		

Comparing the variants

Column three: what the task head reads

key: x an image, c its concept labels, y its class label

variant	data needed	concepts from	what the head reads	trained
CBM Koh et al. 2020	$\{(x,\ c,\ y)\}$	human labels	$\hat{c}_j \in [0, 1]$ one number per concept	
CEM Espinosa Zarlenga et al. 2022	$\{(x,\ c,\ y)\}$	human labels	$\hat{c}_j \in \mathbb{R}^m$ a vector per concept	
PCBM Yuksekgonul et al. 2023	$\{(x,\ y)\}$ + a concept bank	a concept bank	scaled cosine scores phrase example	
CBD (by description) Menon and Vondrick 2023	none for new task	a language model	$\cos(\phi(x),\ \psi(d))$ raw CLIP scores	

Comparing the variants

Column four: what is trained

key: x an image, c its concept labels, y its class label

variant	data needed	concepts from	what the head reads	trained
CBM Koh et al. 2020	$\{(x,\ c,\ y)\}$	human labels	$\hat{c}_j \in [0,1]$ one number per concept	$g_\theta,\ f_\phi$ predictor + head
CEM Espinosa Zarlenga et al. 2022	$\{(x,\ c,\ y)\}$	human labels	$\hat{c}_j \in \mathbb{R}^m$ a vector per concept	$g_\theta,\ f_\phi$ predictor + head
PCBM Yuksekgonul et al. 2023	$\{(x,\ y)\}$ + a concept bank	a concept bank	scaled cosine scores phrase example	f_ϕ head only
CBD (by description) Menon and Vondrick 2023	none for new task	a language model	$\cos(\phi(x),\ \psi(d))$ raw CLIP scores	no task fitting

Comparing the variants

High class accuracy does not prove meaningful concepts

Control experiment: keep image knowledge, remove concept learning.

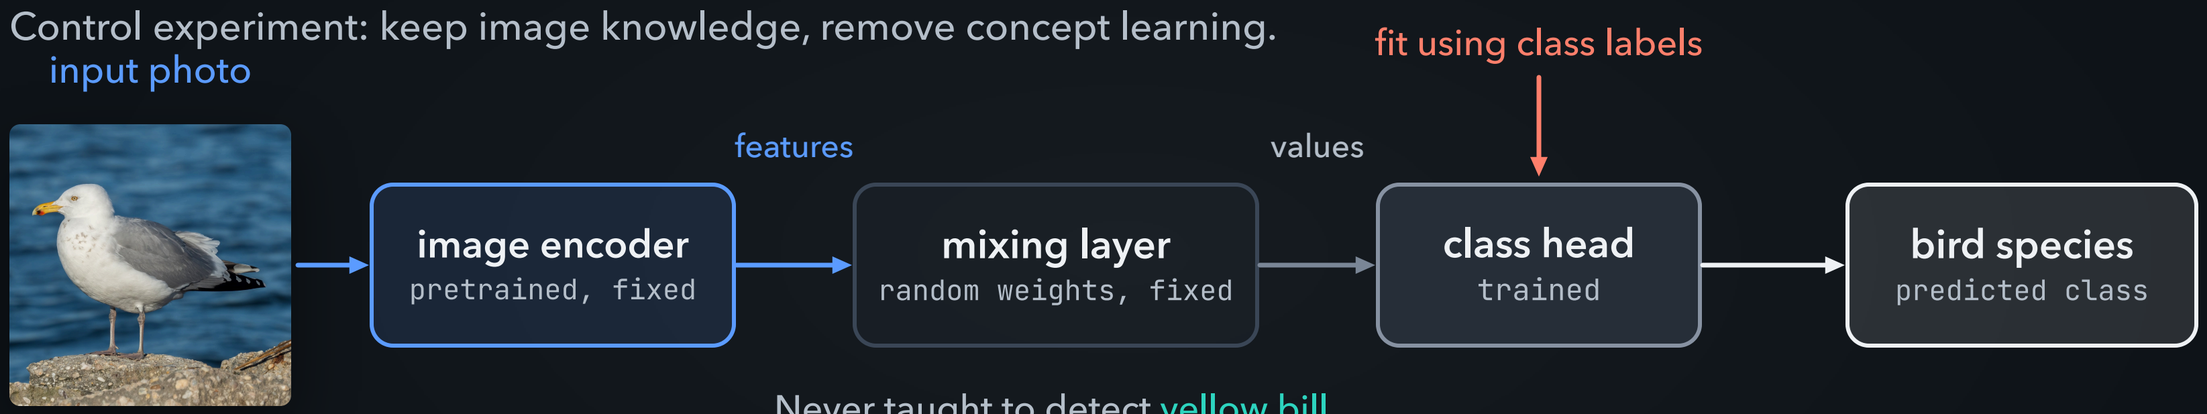

Never taught to detect yellow bill.

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

With enough mixed features, the head can still **classify well**.

Class accuracy does not establish what each intermediate value means.

Random: how the weights are chosen. Fixed: whether training changes them.

REPORTED · Srivastava et al. 2024, Section 4 / Figure 3; bird pipeline schematic

Comparing the variants

Random weights: draw the mixing coefficients by chance

Zoom in: the layer mixes features from the image encoder.

Draw the weights once by chance, without learning them from labels.

image features		weights		contributions
2	×	0.4	=	0.8
1	×	-0.2	=	-0.2

add the contributions

0.8 - 0.2 = 0.6

one mixed feature

Nothing taught this value to mean yellow bill.

ILLUSTRATIVE VALUES · chosen features and possible random weights; arithmetic computed

Comparing the variants

Fixed means: training does not update these weights

During training of this baseline:

image encoder	random mixing layer	class head
keep pretrained weights	keep the one draw	adjust using class labels

The same mixing weights can give different outputs:

photo	image features	same weights	mixed output
A	`[2, 1]`	`[0.4, -0.2]`	0.6
B	`[1, 2]`	`[0.4, -0.2]`	0

Check the concept meaning as well as the class accuracy.

At test time, the trained class head stays fixed too.

`ILLUSTRATIVE VALUES · same toy weights and two chosen feature vectors`

Comparing the variants

Who supplies the concepts, and who checks them?

variant	who supplies the concepts	who checks them
CBM, CEM	people label them	**people**
PCBM	a concept bank	whoever built the bank
CBD	a language model	**review needed**

Comparing the variants

Each rung up the ladder removes a cost and a check

$\{x,\ c,\ y\}$ CBM, CEM

$\{x,\ y\}$ PCBM

$\varnothing$ by description

Labour saved on labels has to be spent later, on checking.

SCHEMATIC · bar levels drawn for teaching

PART 11

Evaluating a concept bottleneck

Check the concepts, predictions and interventions

How should we evaluate a CBM?

Accuracy alone hides most of what a concept bottleneck promises

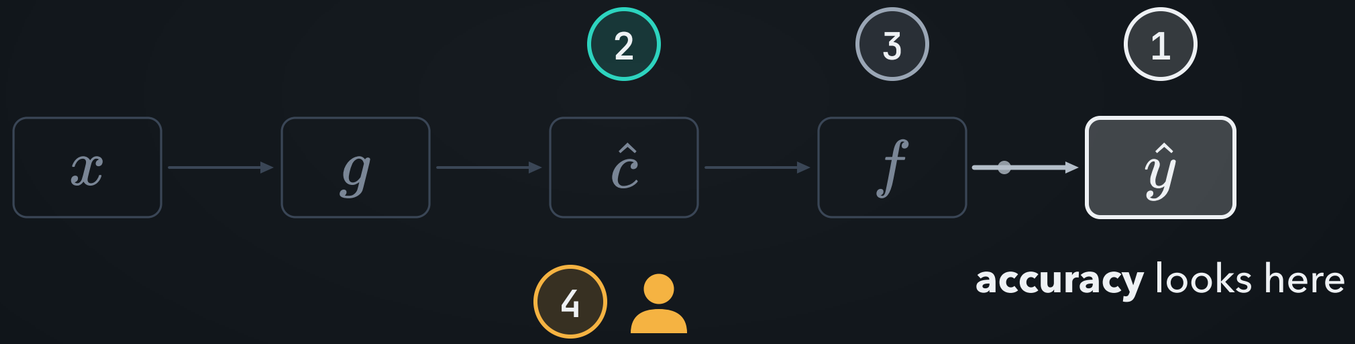

four questions, each about a different part of the model

How should we evaluate a CBM?

Task accuracy: are the predictions correct?

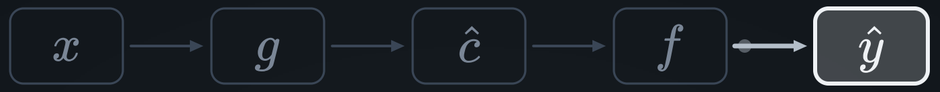

In everyday terms: is the final **answer** right?

MEASURE

the share of held-out test images classified correctly

held-out: images the model never saw during training

How should we evaluate a CBM?

Task accuracy counts the right answers

$$\text{task accuracy} = \frac{1}{N}\sum_{i=1}^{N} \mathbf{1}\left[\hat{y}_i = y_i\right]$$

N test images; $\mathbf{1}[\cdot]$ is 1 when its statement is true, 0 otherwise

prediction $\hat{y}_i$	true label y_i	$\mathbf{1}[\hat{y}_i = y_i]$
small red circle	small red circle	1
large blue square	large blue square	1
large red square	small red square	0

$$\text{task accuracy} = \frac{1+1+0}{3} = \frac{2}{3}$$

ILLUSTRATIVE VALUES · three predictions for teaching

How should we evaluate a CBM?

Concept accuracy: are the concept values right?

In everyday terms: did the model read the evidence correctly?

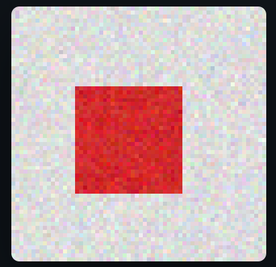

test image 10
small red square

	score $\hat{c}$	$\text{round}(\hat{c})$	label c	match
red	1.00	1	1	1
blue	0.00	0	0	1
circle	0.04	0	0	1
square	0.99	1	1	1
large	0.78	1	0	0

4 / 5

values right on this image

$$\text{concept acc.} = \frac{1}{Nk} \sum_{i,j} \mathbf{1}\left[\text{round}(\hat{c}_{ij}) = c_{ij}\right]$$

MEASURED · saved model, test image 10

How should we evaluate a CBM?

Grounding: does each concept rest on the right evidence?

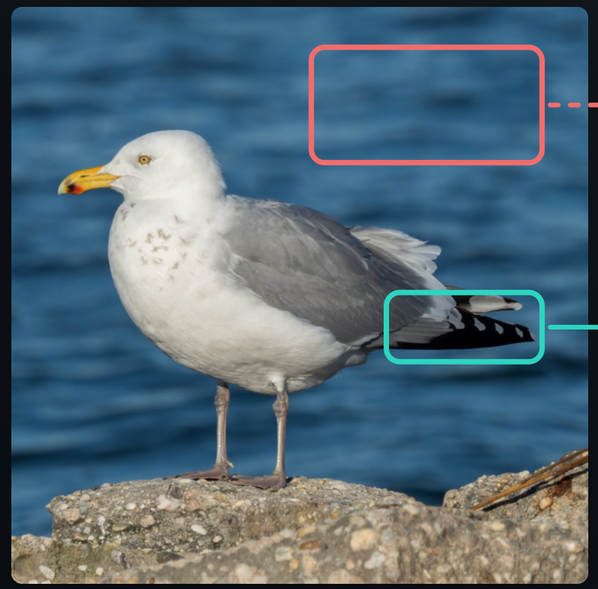

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

not this: water

wingtip black

intended evidence: wingtip

grounding: check where the detector looks

SCHEMATIC · regions drawn by hand

How should we evaluate a CBM?

Faithfulness: does the output really depend on the concepts?

Faithful: the explanation tells the truth about how the **answer** was made

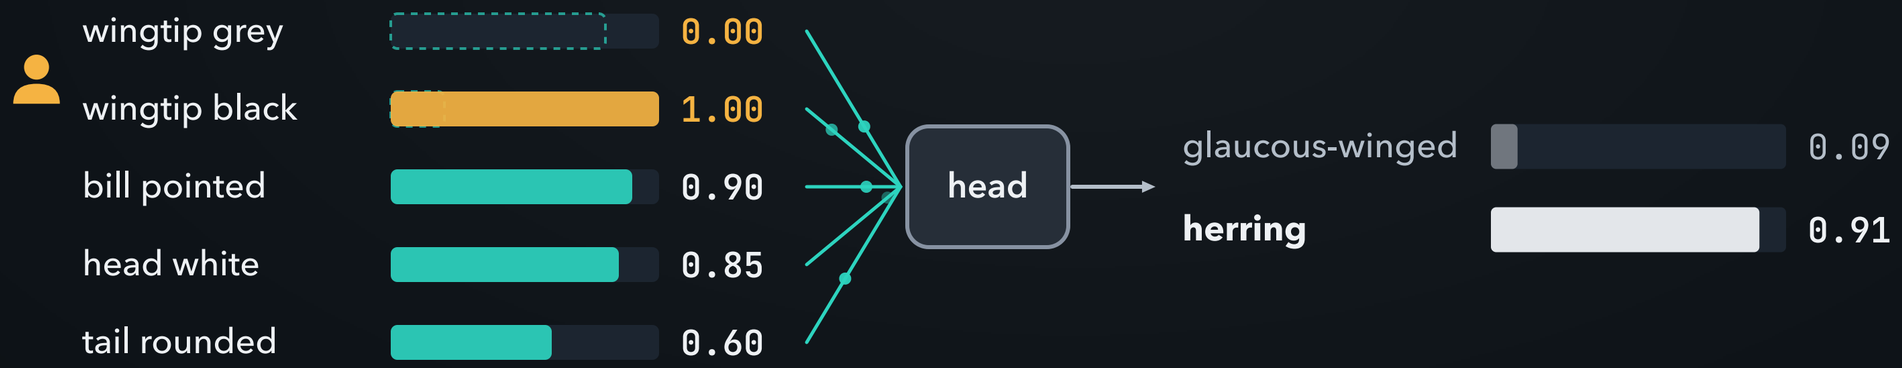

change a concept: the **output** should move as the head says

ILLUSTRATIVE VALUES · the gull head from Part 2

How should we evaluate a CBM?

Intervention utility: do corrections help, and at what cost?

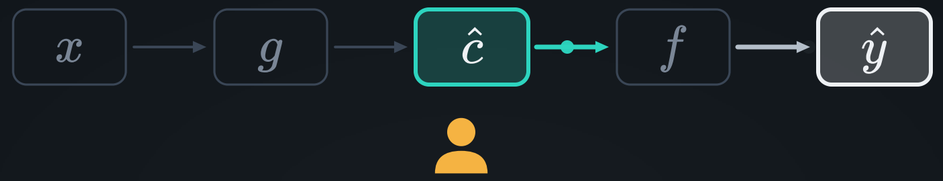

when a person fixes a concept, does the **answer** improve, and at what cost in time?

rises early: better

SCHEMATIC · shape of a curve, no data

How should we evaluate a CBM?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

A controlled test: synthetic shapes with known concepts

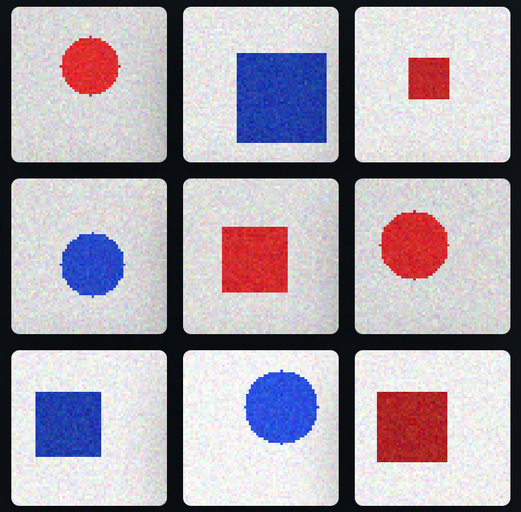

real test inputs, 64 × 64 px

red or blue

circle or square

small or large

→ **8 classes**

2 × 2 × 2

5 concepts

each 1 (yes) or 0 (no)

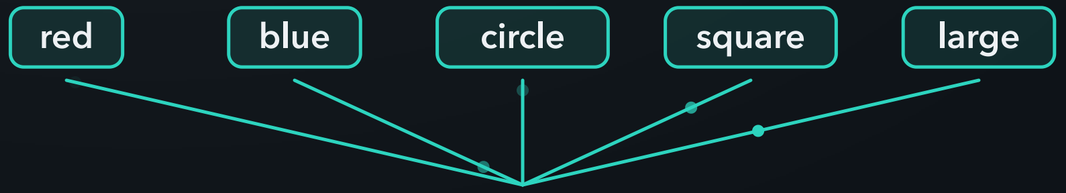

the concepts fully determine the **class**

MEASURED · one seed · 512 synthetic test images · oracle corrections

How should we evaluate a CBM?

x g_θ $\hat{c}$ f_ϕ $\hat{y}$

Three models share one small CNN

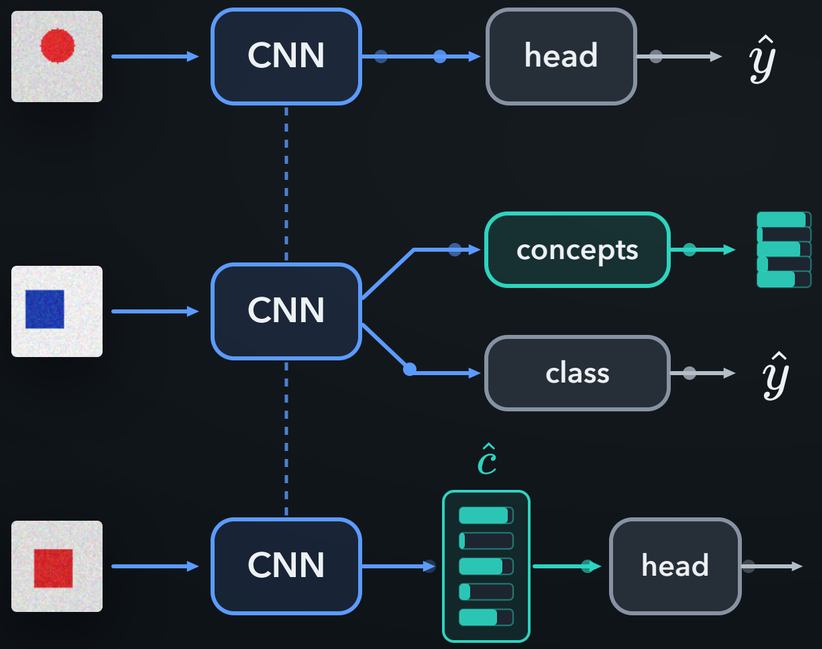

same small CNN, same data, 60 epochs

epoch: one pass through the 768 training images

the class head sits beside the concepts

the head reads only the concepts

How should we evaluate a CBM?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Task accuracy: the three models are close

each bar: the share of the 512 test images classified correctly

on this question alone, the three look about the same

MEASURED · one seed · 512 synthetic test images · oracle corrections

How should we evaluate a CBM?

Concept accuracy: most concept values are right

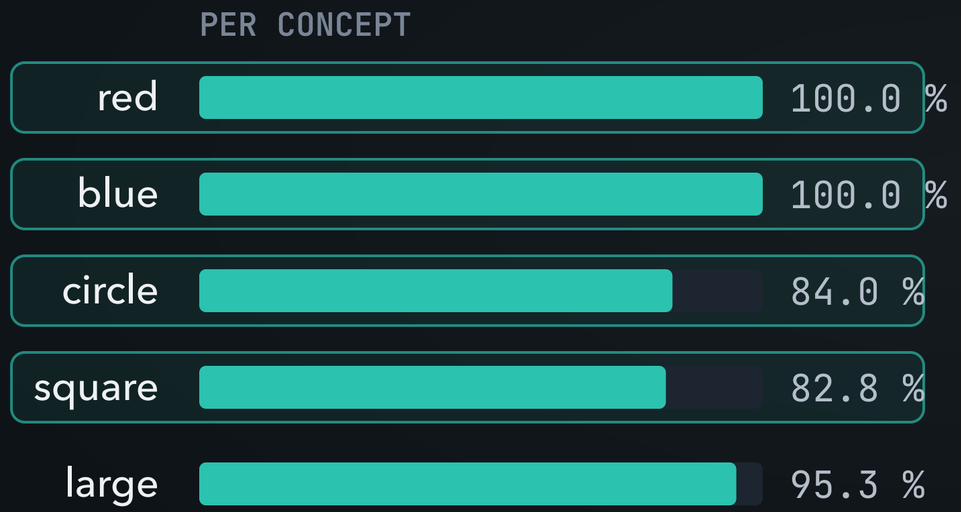

92.4 %

of concept values correct, strict CBM

all 512 images × 5 concepts

MEASURED · one seed · 512 synthetic test images · oracle corrections

How should we evaluate a CBM?

One edit to one concept changes the prediction

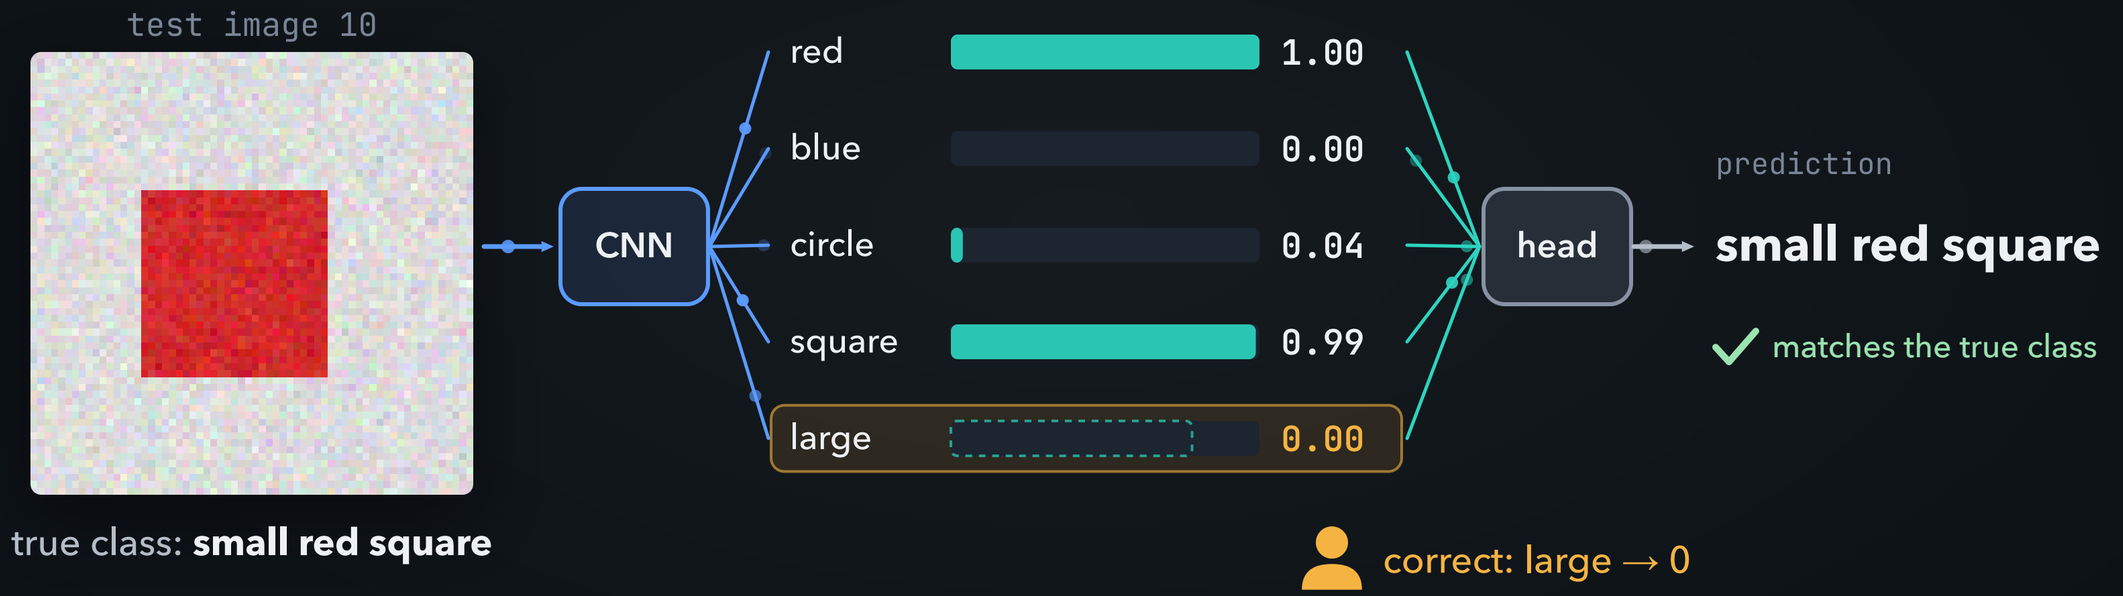

$$\hat{\mathbf{c}} = (1.00,\ 0.00,\ 0.04,\ 0.99,\ 0.78) \rightarrow (1.00,\ 0.00,\ 0.04,\ 0.99,\ 0)$$

MEASURED · saved model, test image 10

How should we evaluate a CBM?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Oracle corrections: fix the worst concept first

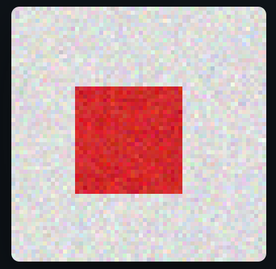

test image 10

	score $\hat{c}$	label c	error $\lvert\hat{c} - c\rvert$	
red	1.00	1		0.00
blue	0.00	0		0.00
circle	0.04	0		0.04
square	0.99	1		0.01
large	0.78	0		0.78

fix this one first

oracle: the true labels play a perfect expert

$$\lvert 0.78 - 0 \rvert = 0.78$$

MEASURED · saved model, test image 10

How should we evaluate a CBM?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Each step right is one more **correction** per image

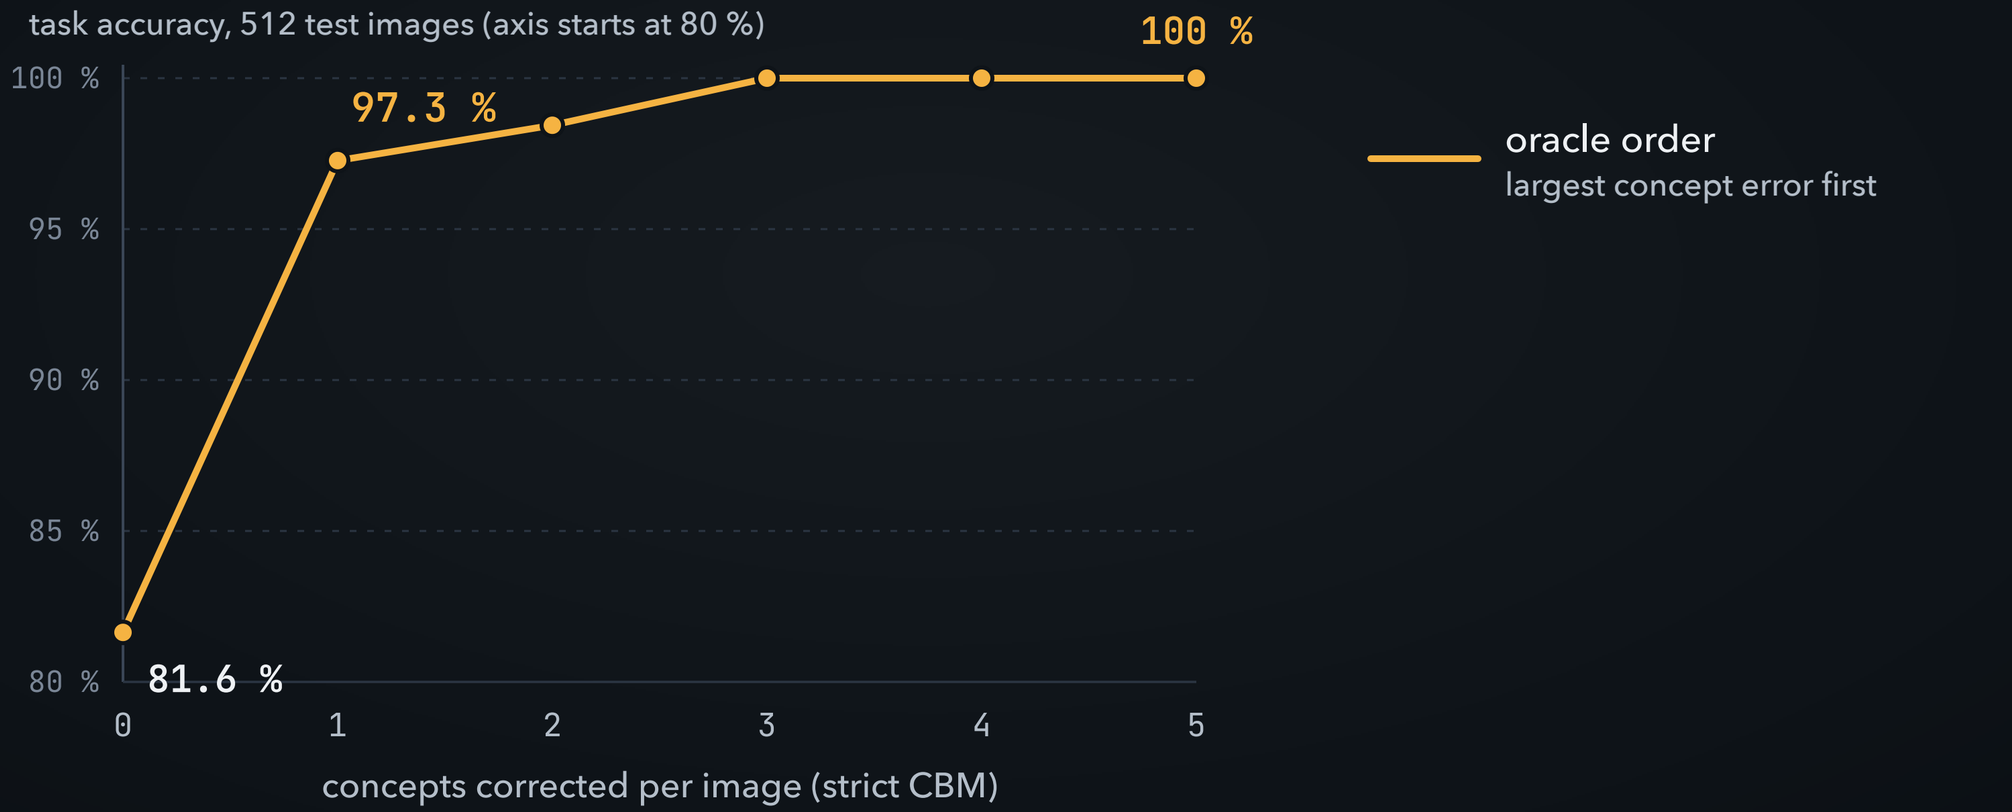

MEASURED · one seed · 512 synthetic test images · oracle corrections

How should we evaluate a CBM?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

Random corrections help more slowly

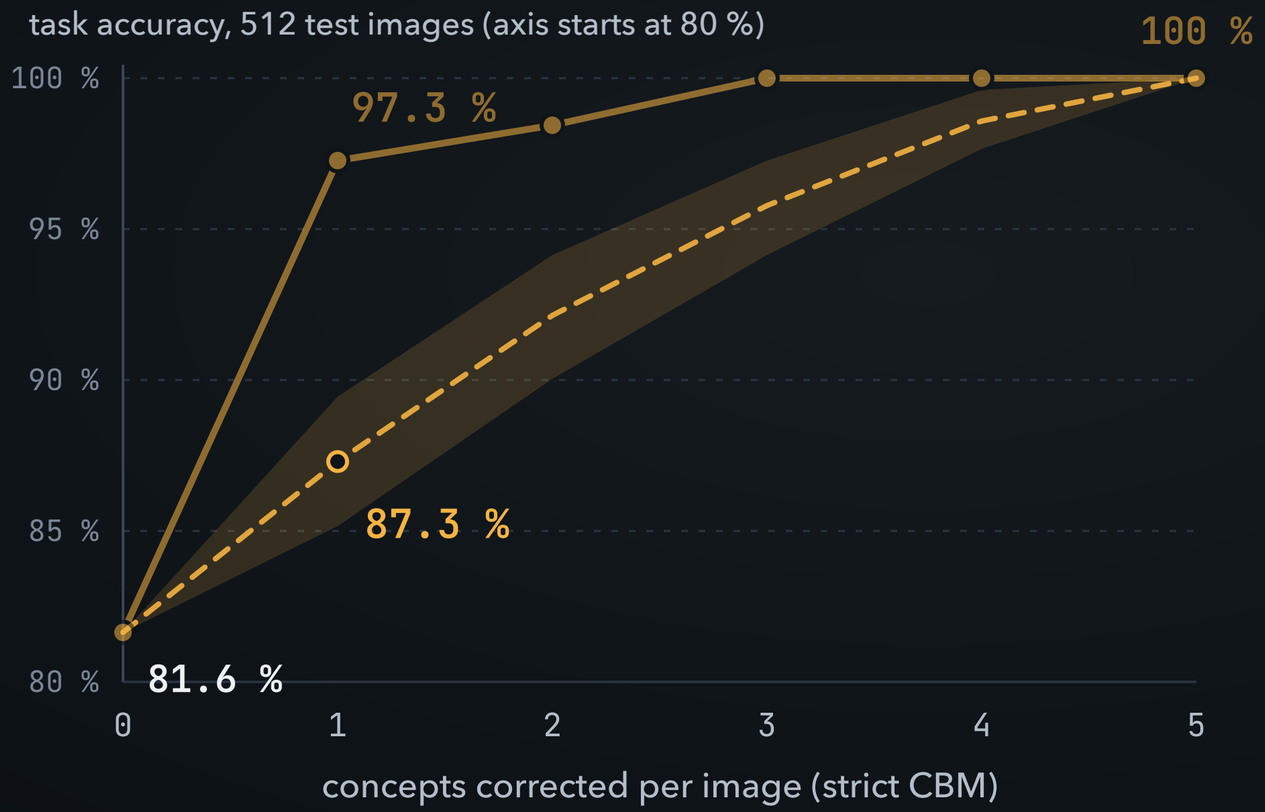

oracle order
largest concept error first

random order
mean over 32 random orders

lowest to highest
of the 32 runs

a random pick often fixes a concept that was already right

MEASURED · one seed · 512 synthetic test images · oracle corrections

How should we evaluate a CBM?

x — g_θ — $\hat{c}$ — f_ϕ — $\hat{y}$

In the multi-task model, a concept edit changes nothing

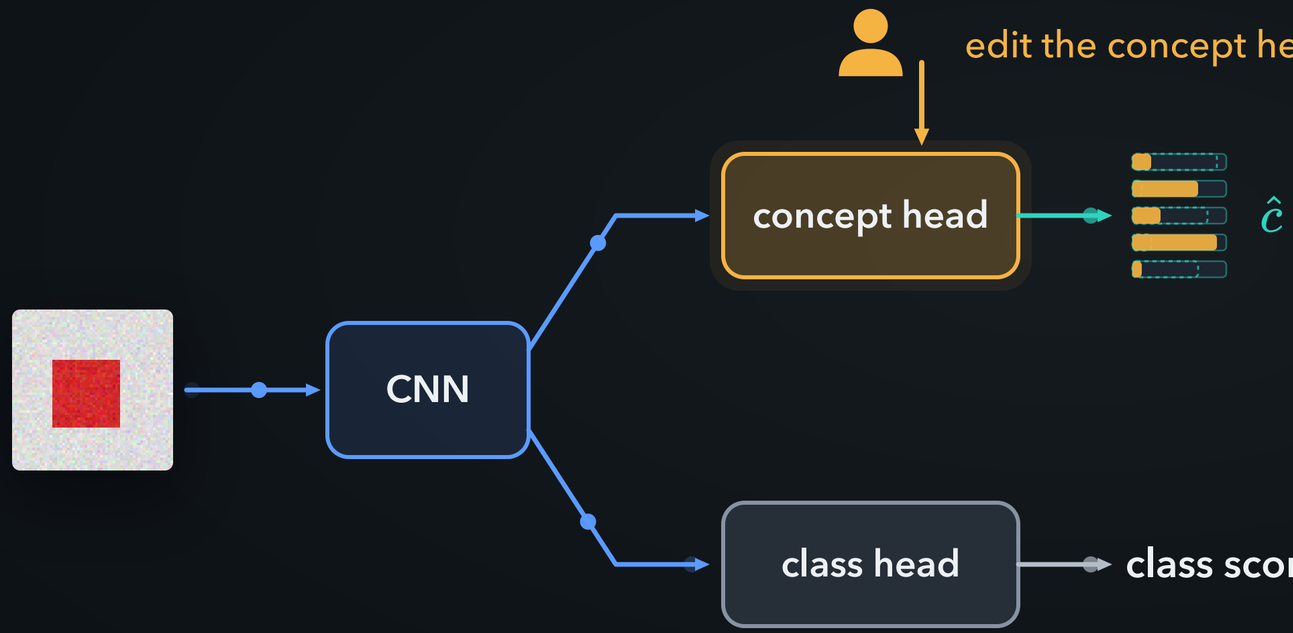

✓ unchanged (verified exactly)

the class head never reads the concepts

MEASURED · one seed · 512 synthetic test images · oracle corrections

How should we evaluate a CBM?

Read these numbers with care

1. one training seed: a single training run
2. 512 synthetic test images
3. oracle corrections from the true labels
4. no human study: a perfect, effortless expert

---

real experts make mistakes, and their time is costly

How should we evaluate a CBM?

Similar accuracy, very different intervention utility

task accuracy	78.7 % · 82.0 % · 81.6 % black box · multi-task · strict CBM
concept accuracy	92.4 % of concept values (strict CBM) grounding: not measured here
faithfulness	strict CBM: the head reads only the concepts multi-task: the class head ignores them
intervention utility	one correction: 81.6 → 97.3 % (strict CBM), no change (multi-task) human cost: not measured here

`MEASURED · one seed · 512 synthetic test images · oracle corrections`

How should we evaluate a CBM?

Why similar accuracy can hide very different intervention utility

accuracy marks only the final answer, not the route to it

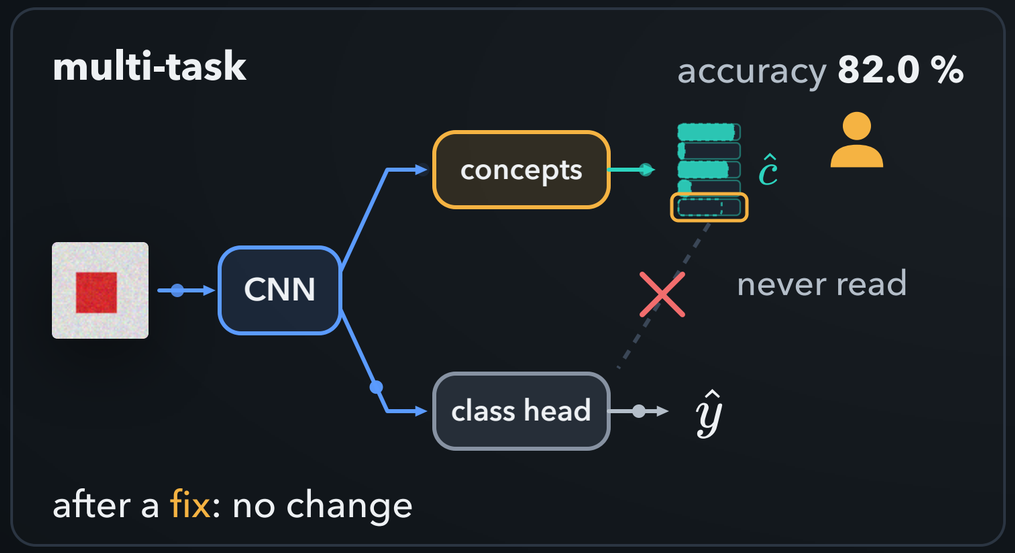

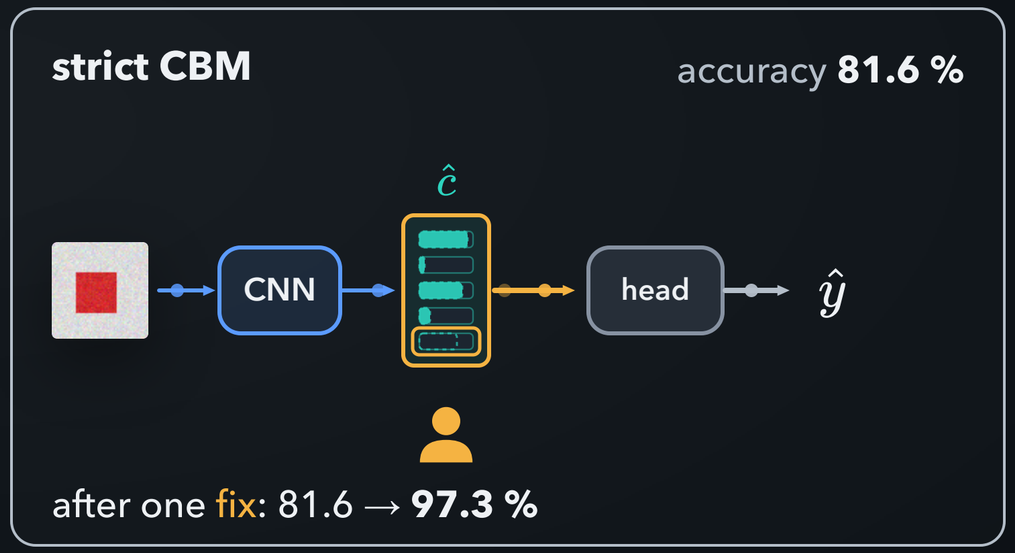

a fix helps only if the head reads the concepts

MEASURED · one seed · 512 synthetic test images · oracle corrections

PART 12

Beyond images

The same idea with different inputs

Beyond images

Where else do concept bottlenecks fit?

SO FAR: PHOTOGRAPHS OF BIRDS

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

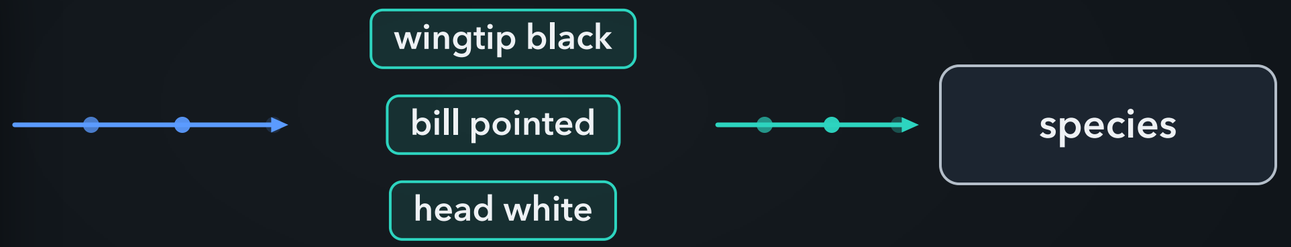

photo x concepts $\hat{\mathbf{c}}$ **species** $\hat{y}$

they fit wherever experts reason with named findings

Beyond images

Medicine: a knee X-ray → clinical findings → a severity grade

input x

named concepts $\hat{\mathbf{c}}$
each finding is a small grade, 0 = none

joint space narrowing	1
bone spurs	2
sclerosis	1
7 more findings	⋮

task $\hat{y}$

osteoarthritis grade

$$\hat{y} \in \{0, 1, 2, 3\}$$

$$\hat{y} = 2$$

0 none or doubtful · 3 severe

Knee X-ray: Koh et al. 2020, Fig. 1 (CC BY 4.0), Osteoarthritis Initiative

Koh et al., ICML 2020 · Osteoarthritis Initiative

ILLUSTRATIVE VALUES · values for teaching

Beyond images

A radiologist corrects a finding, and the grade is recomputed

input x

Knee X-ray: Koh et al. 2020, Fig. 1 (CC BY 4.0), Osteoarthritis Initiative

named concepts $\hat{\mathbf{c}}$

each finding is a small grade, 0 = none

joint space narrowing	2
bone spurs	2
sclerosis	1
7 more findings	⋮

radiologist corrects a finding

task $\hat{y}$

osteoarthritis grade

$$\hat{y} \in \{0, 1, 2, 3\}$$

$$\hat{y} = 3$$

0 none or doubtful · 3 severe

Why it fits: radiologists already write their reports in these terms

Koh et al., ICML 2020 · Osteoarthritis Initiative

ILLUSTRATIVE VALUES · values for teaching

Beyond images

Language: a review → four aspects → a star rating

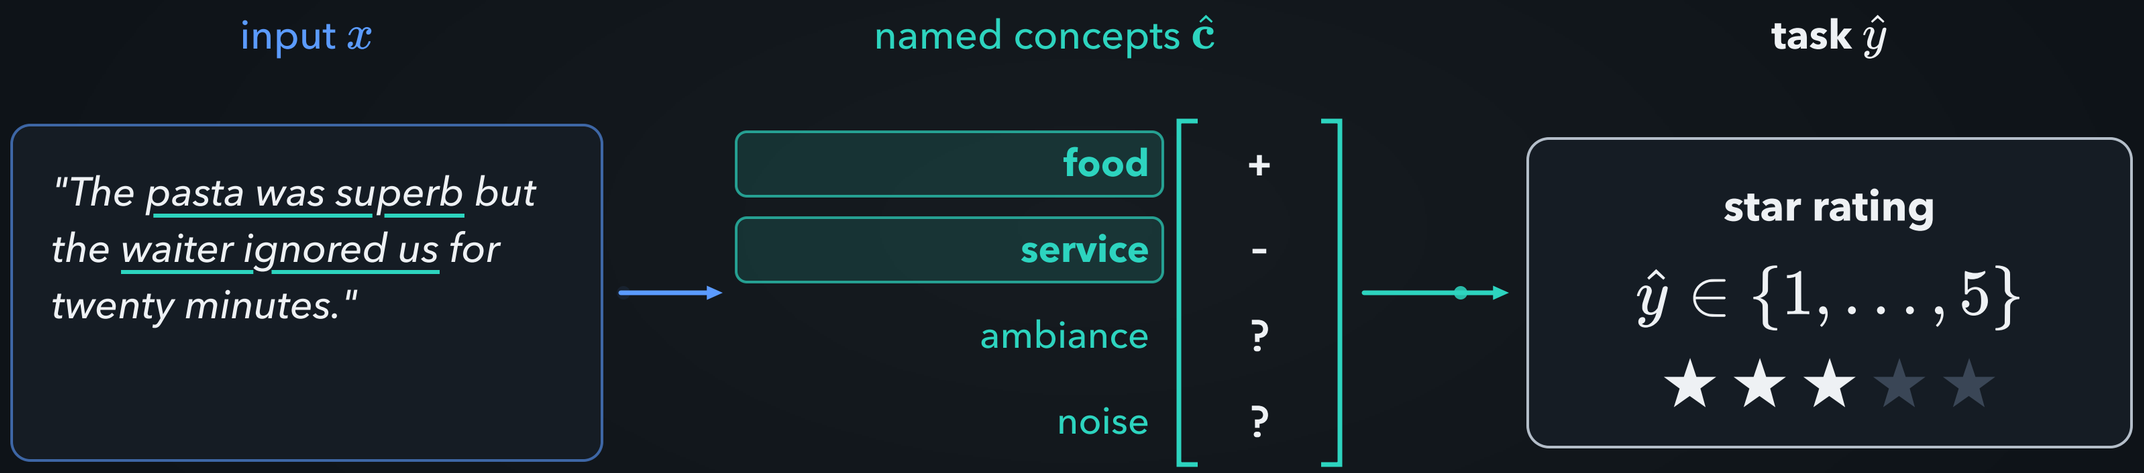

+ praised · – criticised · ? not mentioned

Why it fits: reviewers already write about the food and the service

CEBaB, Abraham et al., NeurIPS 2022 · text bottleneck models, Ludan et al. 2023

ILLUSTRATIVE VALUES · review and values for teaching

Beyond images

Intensive care: vital signs → summary concepts → drugs needed soon?

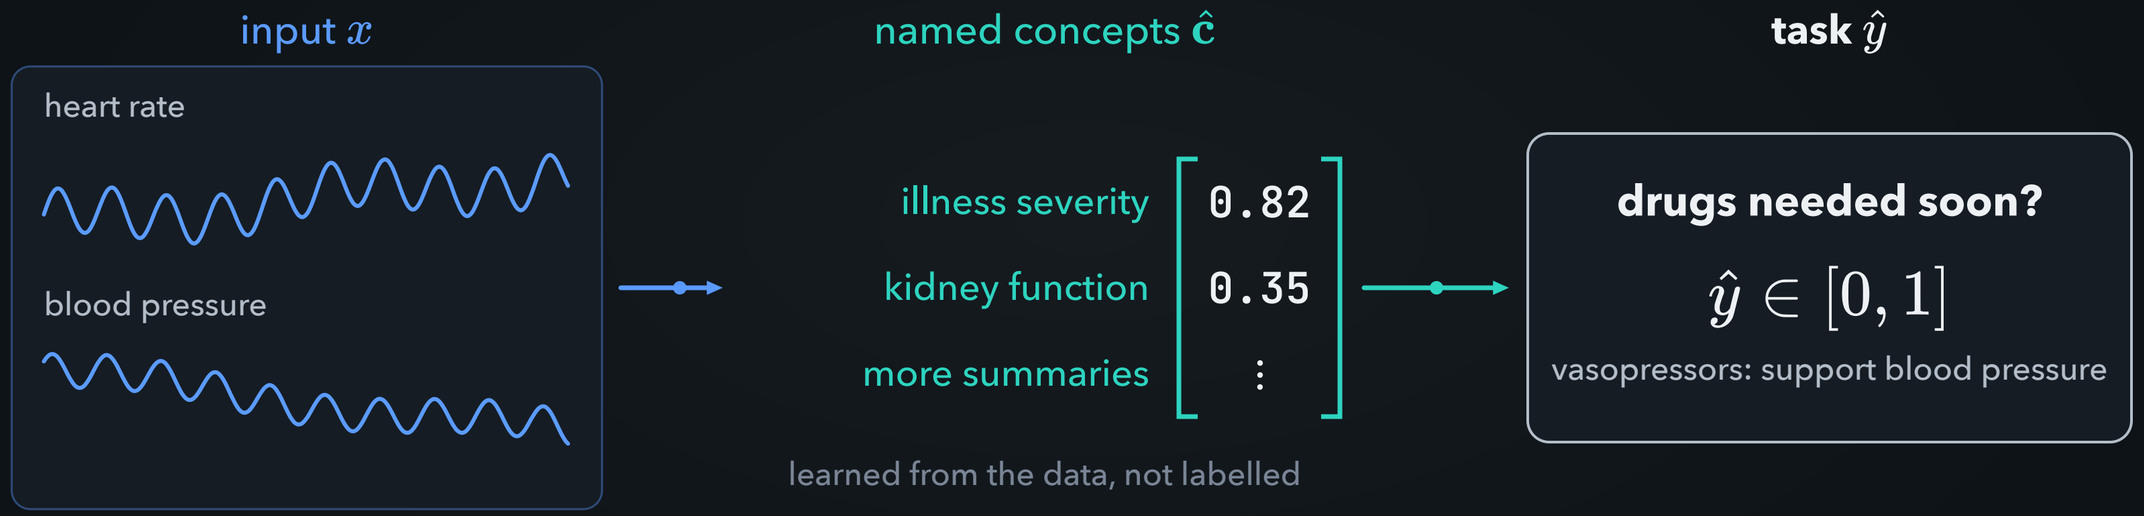

Why it fits: doctors already sum up a patient in terms like illness severity and kidney function

Wu et al., MLHC 2022 · MIMIC-III, 15,552 patients

ILLUSTRATIVE VALUES · schematic traces, values for teaching

Beyond images

Wearables: sensor signal → simple activities → complex activity

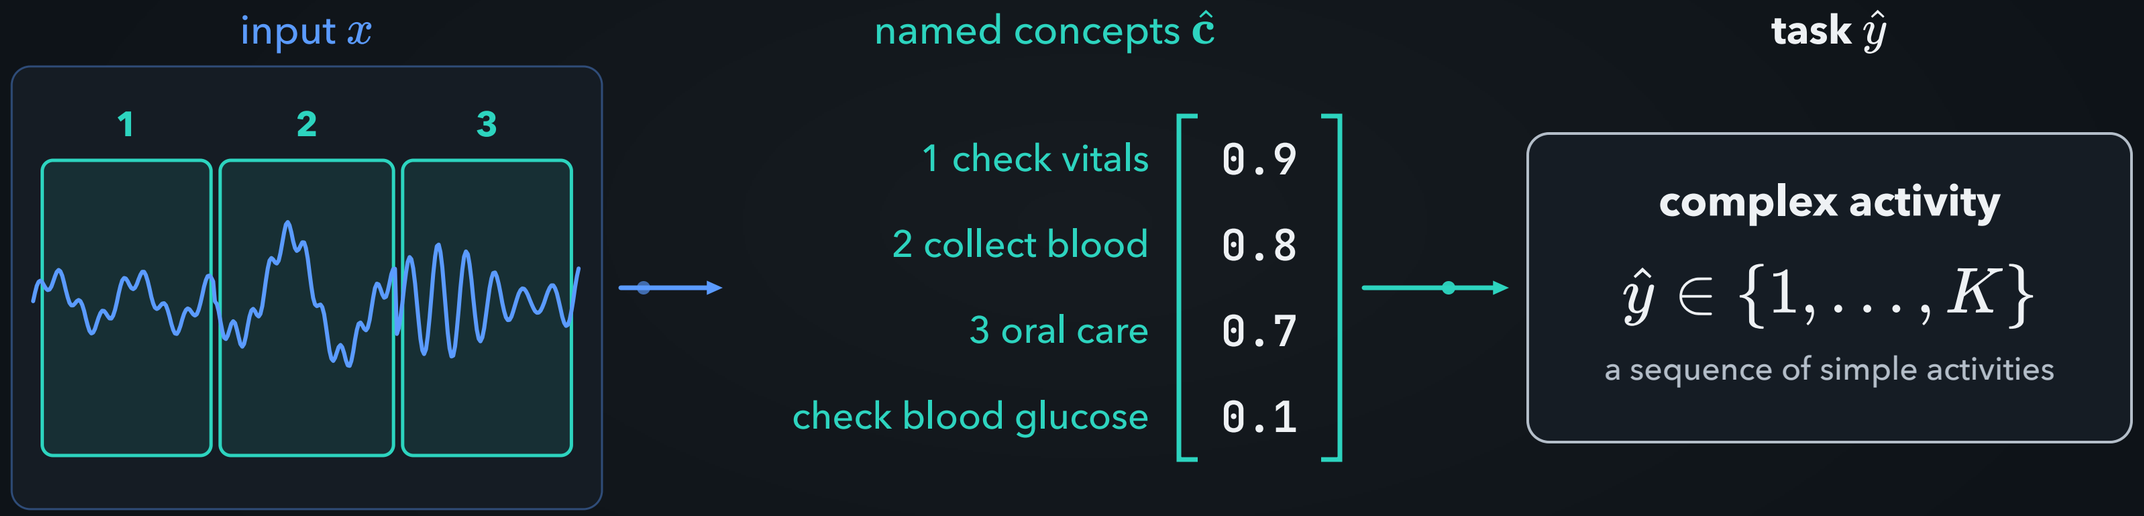

Why it fits: people already describe a **complex activity** as its simple steps

X-CHAR, Jeyakumar et al., IMWUT 2023 · concepts from its nurse data

ILLUSTRATIVE VALUES · schematic trace, values for teaching

Beyond images

Sleep: an EEG → the patterns scorers use → a sleep stage

input x

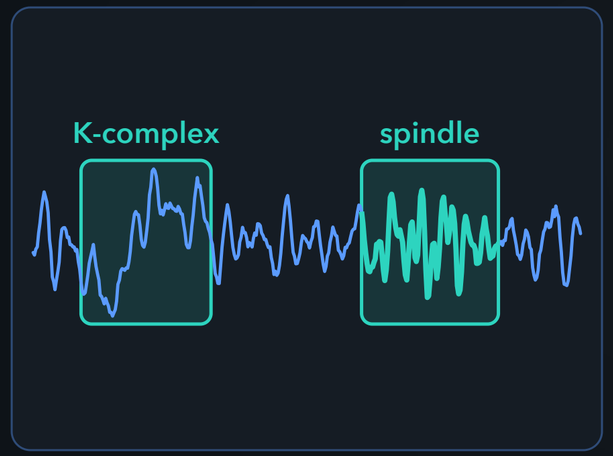

named concepts $\hat{\mathbf{c}}$

spindle	1.00
K-complex	0.80
slow waves	0.15

task $\hat{y}$

sleep stage

$\hat{y} \in \{\mathrm{W}, \mathrm{N1}, \mathrm{N2}, \mathrm{N3}, \mathrm{R}\}$

N2 candidate: **0.86**

W wake · N1, N2, N3: stages one to three · R: REM

Why it fits: scorers already mark these patterns by hand

patterns from the AASM scoring manual (Berry et al.)

ILLUSTRATIVE VALUES · schematic EEG, the sleep head from Part 6

Beyond images

One recipe in every domain

Medicine
knee X-ray → grade

Language
review → stars

Intensive care
vital signs → drug need

Wearables
sensor signal → activity

Sleep
EEG → sleep stage

THE RECIPE

name the findings an expert checks

$$c_1, \ldots, c_k$$

predict them from the raw signal

$$\hat{\mathbf{c}} = g_\theta(x)$$

let the **decision** use them alone

$$\hat{y} = f_\phi(\hat{\mathbf{c}})$$

PART 13

What can still go wrong?

An editable interface needs careful checks

An interpretable interface is not the end of the story

Follow the path from photo to concepts to answer: what can fail?

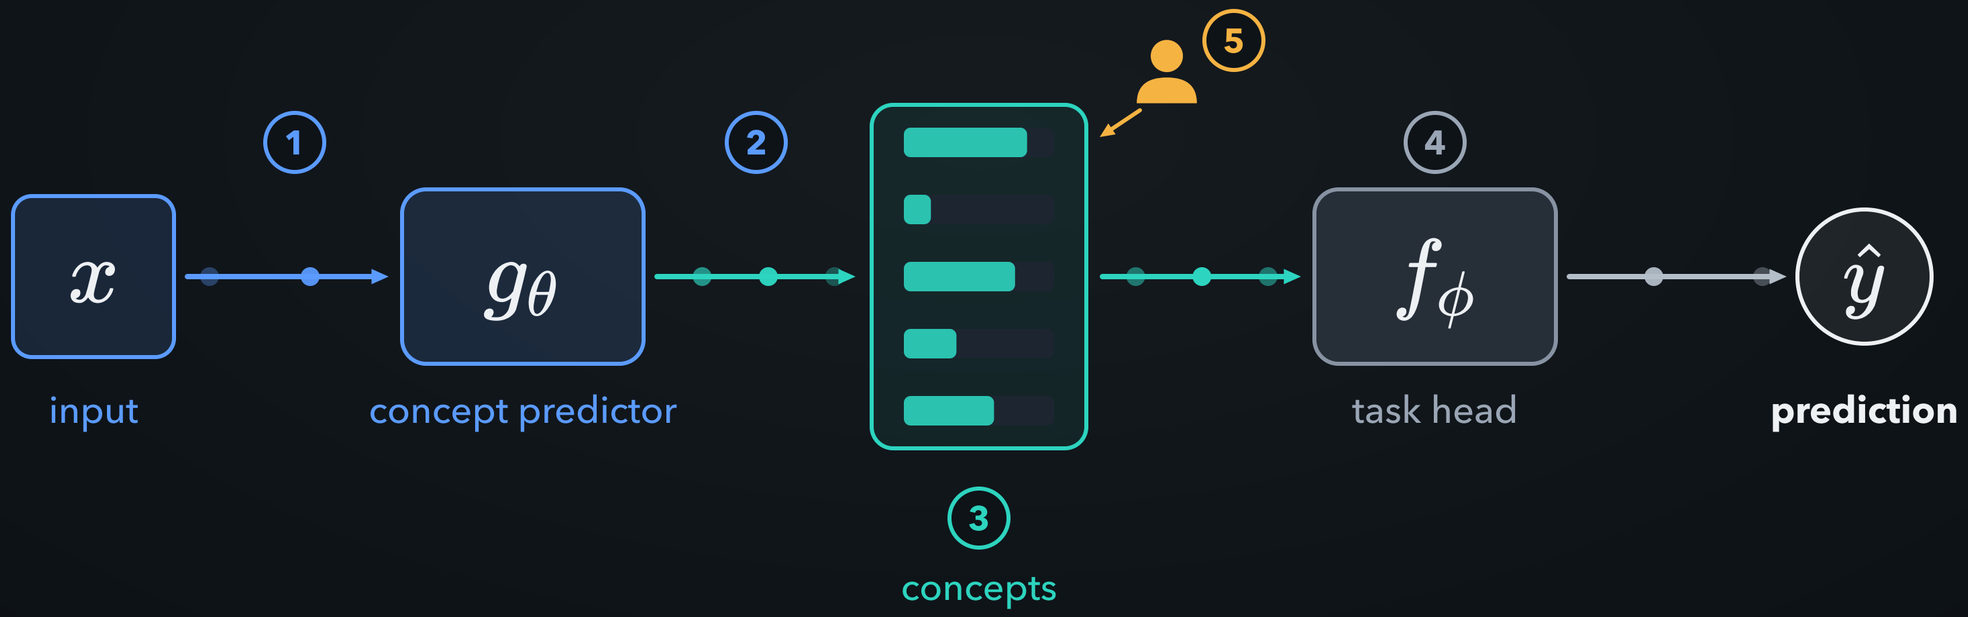

An interpretable interface is not the end of the story

1 · Encoder: distribution shift

new inputs unlike the training inputs

$$p_{\text{new}}(x) \neq p_{\text{train}}(x)$$

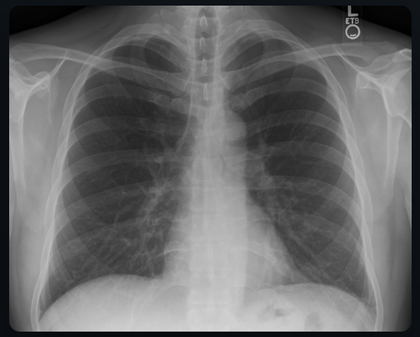

Chest X-ray: Stillwaterising, Wikimedia Commons, CC0

An interpretable interface is not the end of the story

A new hospital: the same model scores lower

new inputs unlike the training inputs

$$p_{\text{new}}(x) \neq p_{\text{train}}(x)$$

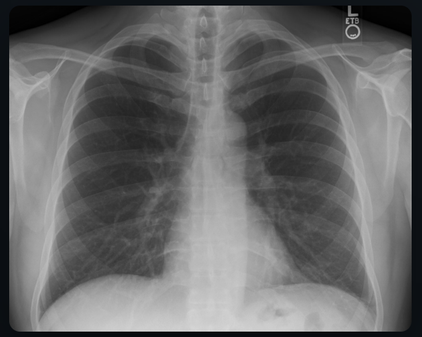

Chest X-ray: Stillwaterising, Wikimedia Commons, CC0

pneumonia detection, AUC (1 = perfect ranking, 0.5 = guessing)

same hospitals	0.931
new hospital	0.815

In practice: fine where it was built, worse at the new site

REPORTED · Zech et al., PLOS Medicine 2018 (Part 1)

An interpretable interface is not the end of the story

2 · Concept detector: wrong grounding

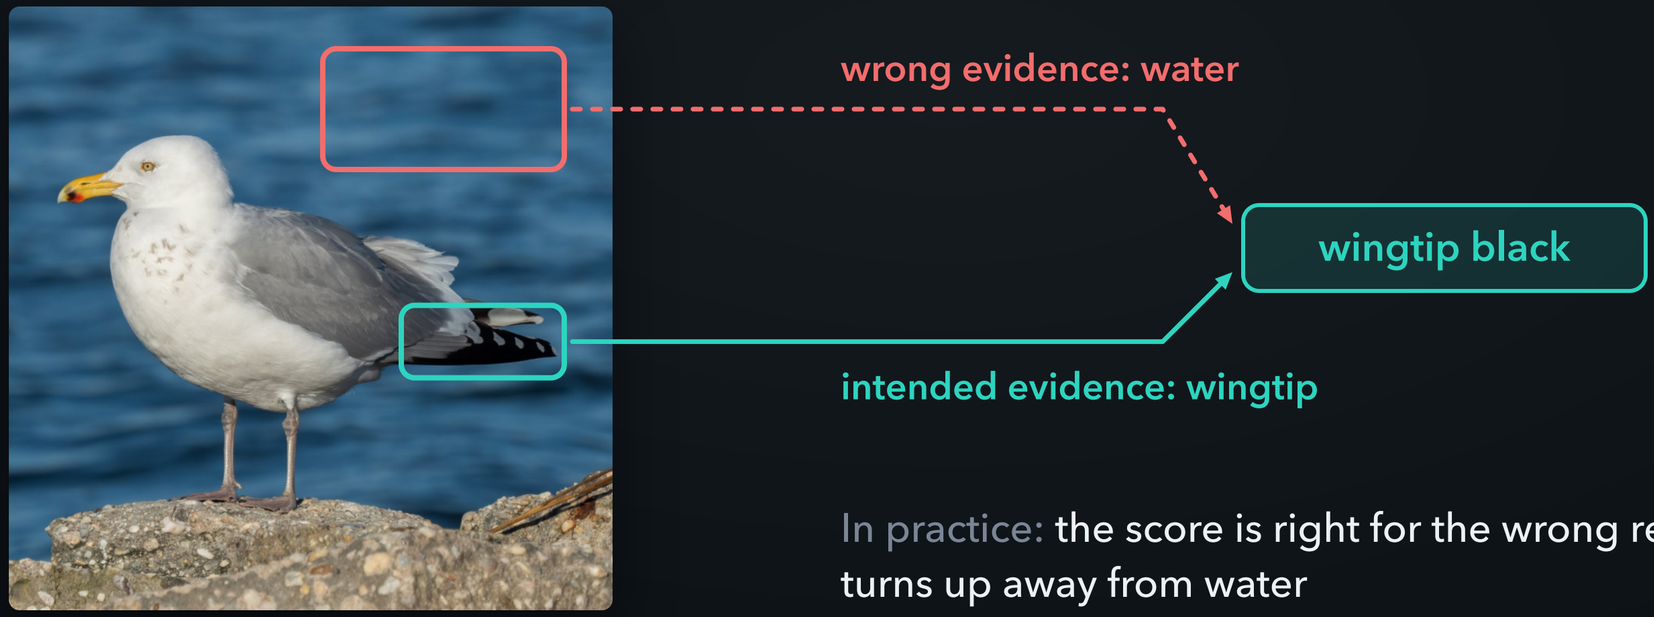

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

In practice: the score is right for the wrong reason, until a gull turns up away from water

SCHEMATIC · regions drawn by hand

An interpretable interface is not the end of the story

3 · Vocabulary: missing information

the same concept vector, different **classes**

$$\mathbf{c}(x_1) = \mathbf{c}(x_2), \quad y_1 \neq y_2$$

shape never reaches the head: best accuracy **50%**

In practice: two cases an expert tells apart look identical to the head

CONSTRUCTED EXAMPLE · the blue circle and square of Part 8

An interpretable interface is not the end of the story

4 · Task head: the wrong relationship

$$\text{score} = 1.2 - 2\,\text{duplicate charge} + 0.6\,\text{refund requested}$$

a correct edit meets a wrong rule: billing falls from 0.80 to 0.45

In practice: the expert's fix is right, yet the decision moves the wrong way

CONSTRUCTED EXAMPLE · head built faulty on purpose, Part 7

An interpretable interface is not the end of the story

5 · Intervention: wrong, or too costly

size: large 0.00

Part 11 took corrections from true labels: an expert who never errs

In practice: real experts misjudge too, and every check costs their time

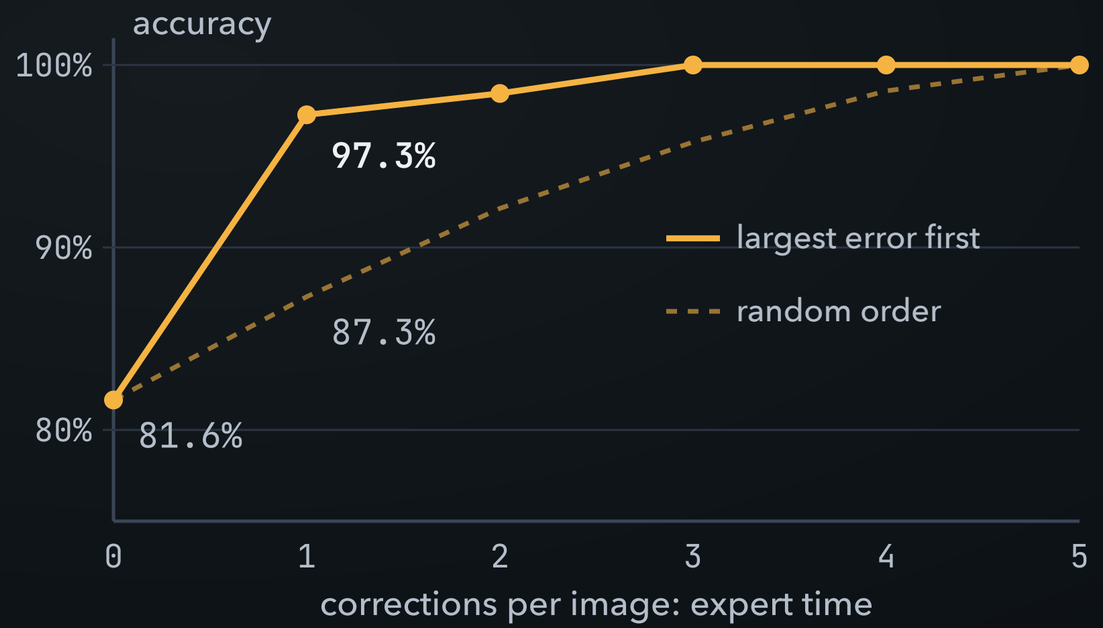

MEASURED · Part 11: one seed, 512 synthetic test images, oracle corrections

An interpretable interface is not the end of the story

Five questions for any concept bottleneck

1. ☐ Are the concepts meaningful?
2. ☐ Are they grounded correctly?
3. ☐ Are they sufficient for the task?
4. ☐ Does the model actually use them?
5. ☐ Do corrections improve decisions?

An interpretable interface is not the end of the story

A 2026 survey maps the field onto the same parts

SURVEY · TMLR 2026

What's in the Bottle? A Survey and Roadmap of Concept Bottleneck Models

P. Knab, D. Steinmann, C. Bartelt, K. Kersting, B. Schiele, T. Seidl, U. Schlegel, W. Stammer

input | concepts | output | training

open challenges include grounding, completeness, leakage and intervention reliability

RECAP

Putting the pieces together

From the image to an editable explanation

Final replay

One last time: a black box maps the photo to a species

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

ILLUSTRATIVE VALUES · hand-written concept values and head

Final replay

Open it: named concepts sit on the path to the prediction

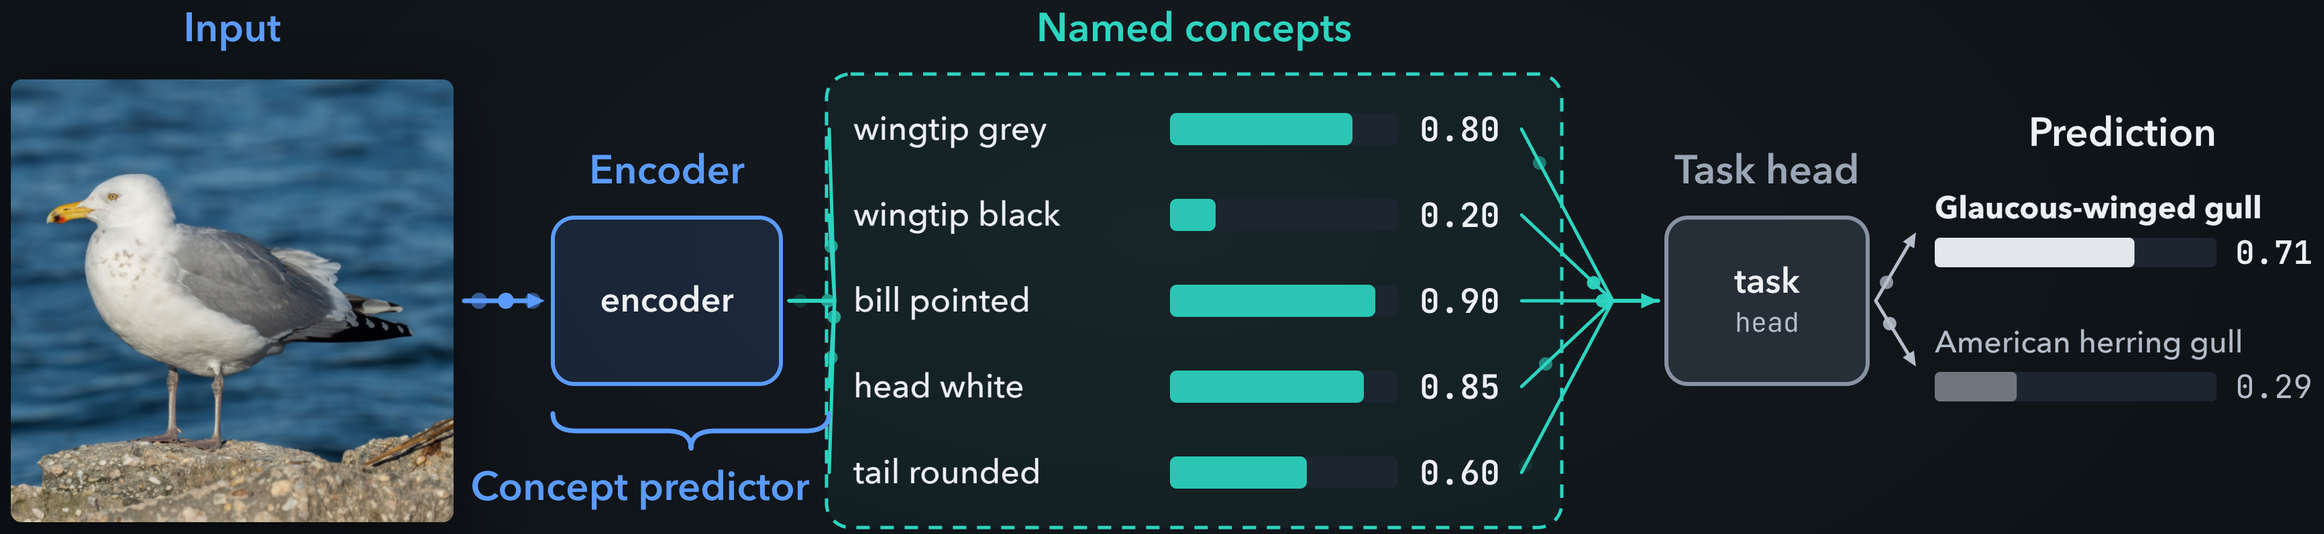

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

ILLUSTRATIVE VALUES · hand-written concept values and head

Final replay

A person corrects one concept, and the prediction follows

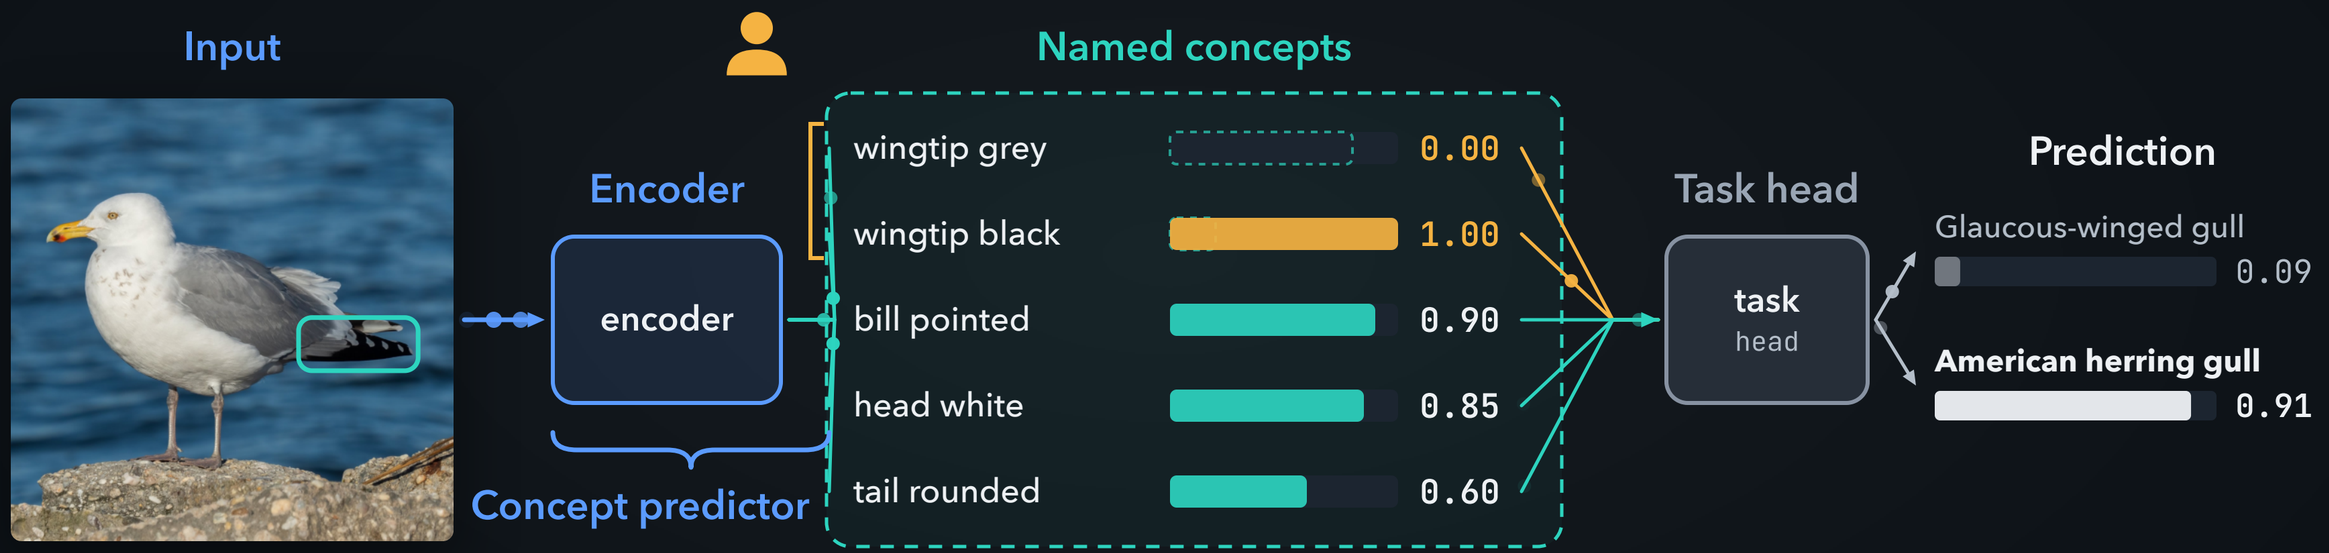

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

ILLUSTRATIVE VALUES · hand-written concept values and head

Final replay

The interface is only as good as its parts

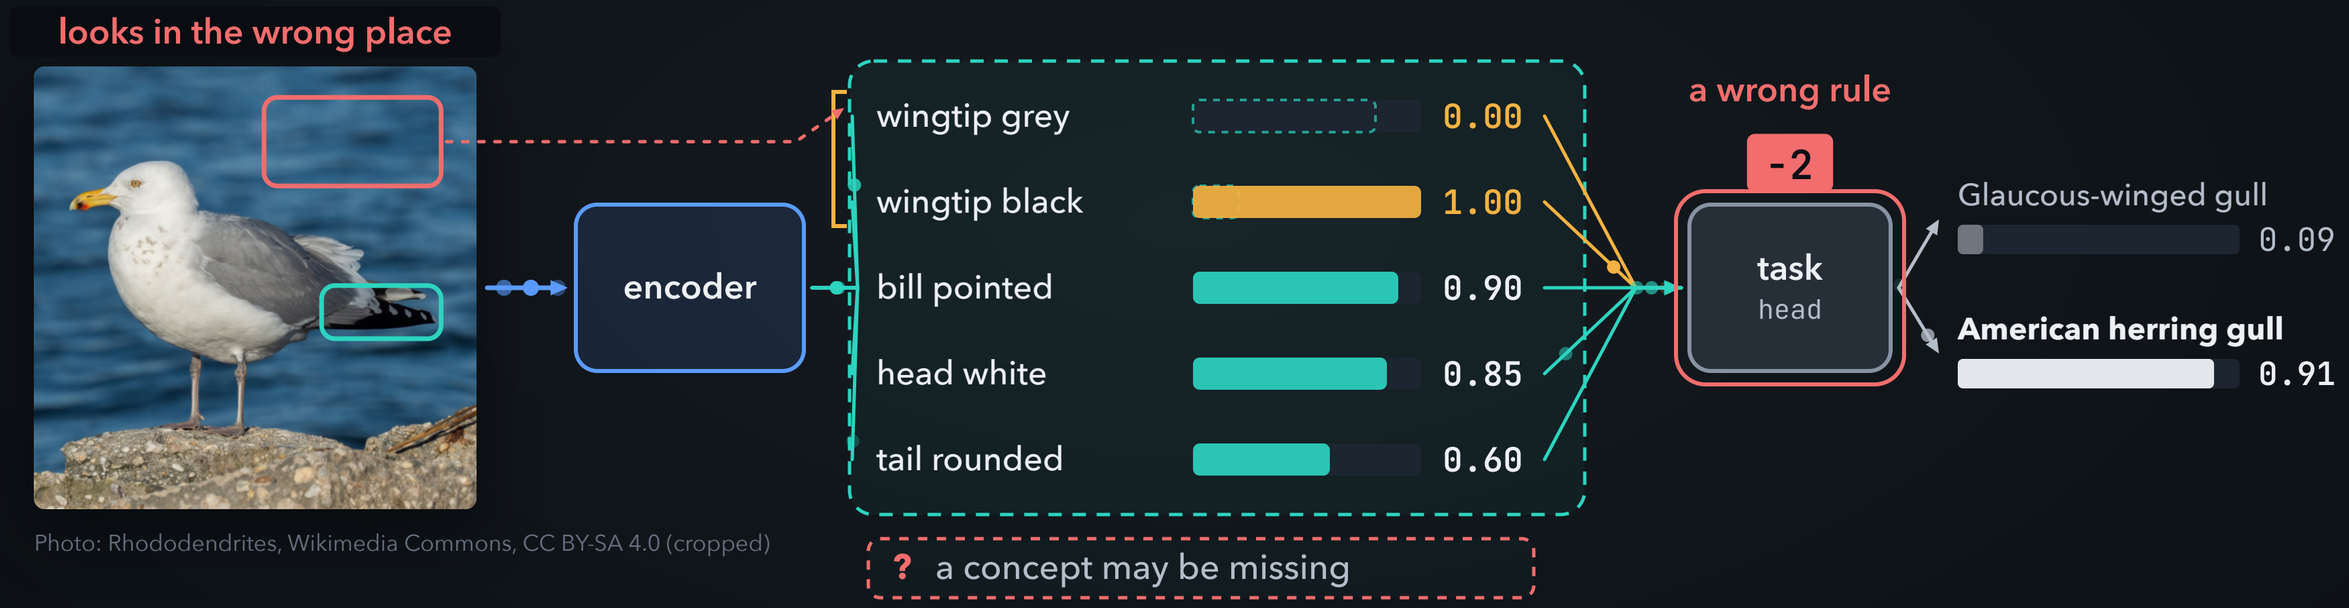

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

ILLUSTRATIVE VALUES · hand-written concept values and head

Final replay

Concepts that can be inspected, evaluated and **corrected**

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

inspected evaluated **corrected**

`ILLUSTRATIVE VALUES · hand-written concept values and head`

Final replay

The interface must carry what the prediction depends on

The challenge is to make sure the interface faithfully represents the information on which the prediction depends.

put simply: whatever the **answer** uses should be visible in the concepts

Final replay

What to remember

1. Post-hoc methods explain afterwards; a concept bottleneck makes the **answer** pass through named concepts
2. Fixing a concept helps only if the head reads it the right way
3. A soft score can carry more than its name says: extra information leaks
4. Newer variants need fewer labels, so check each concept looks at the right evidence

APPENDIX

Testing concept use in a black box

Testing with Concept Activation Vectors (TCAV)

Appendix: testing with concept activation vectors

TCAV asks about a named concept: does the model use grey wings when it says herring gull?

Photo: Rhododendrites, Wikimedia Commons, CC BY-SA 4.0 (cropped)

one more post-hoc method: it asks about concepts directly

TCAV: Kim et al., ICML 2018

Appendix: testing with concept activation vectors

Step 1: collect photos with grey wings and photos without

with grey wings

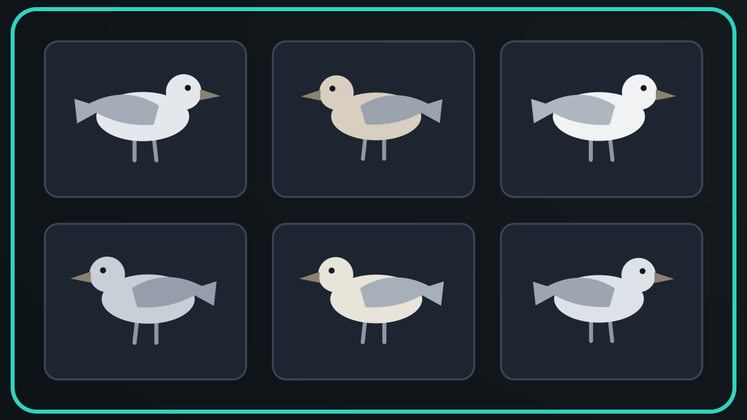

without

150 photos in each group

none is a **herring gull**

SCHEMATIC · drawings stand in for labelled bird photos (CUB, Part 2)

Appendix: testing with concept activation vectors

Step 2: the concept direction separates them inside the model

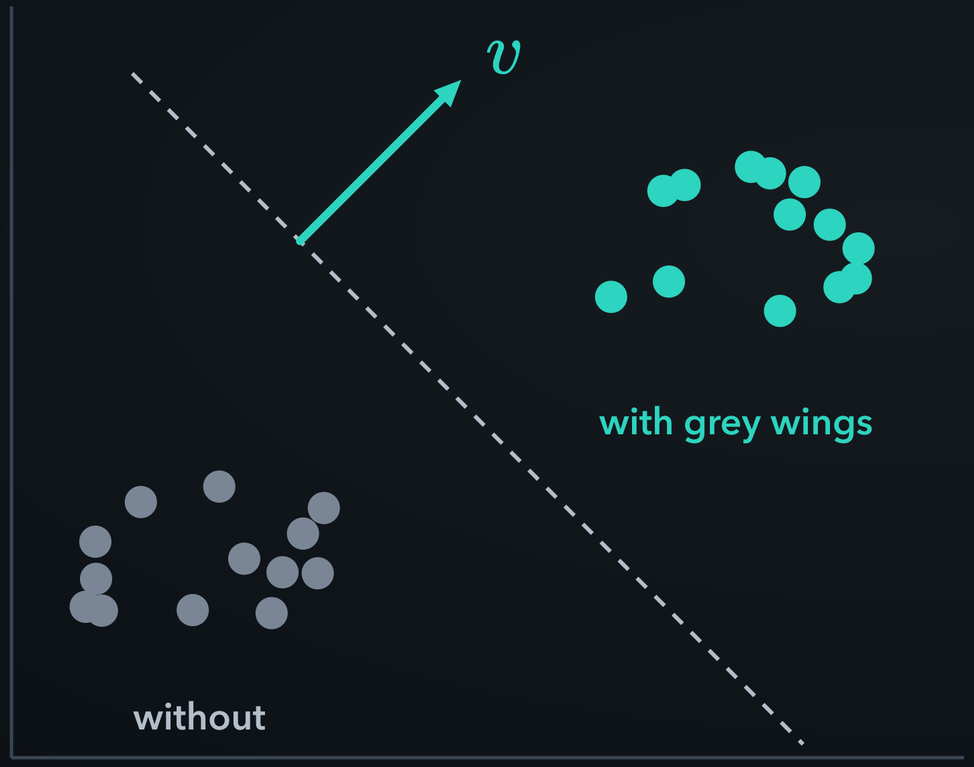

SCHEMATIC · 2 of 512 dimensions drawn

activations: the CLIP features

$$h(x) \in \mathbb{R}^{512}$$

v, the concept activation vector:
moving along it means more grey wing

Appendix: testing with concept activation vectors

Step 3: nudge each herring gull photo along that direction

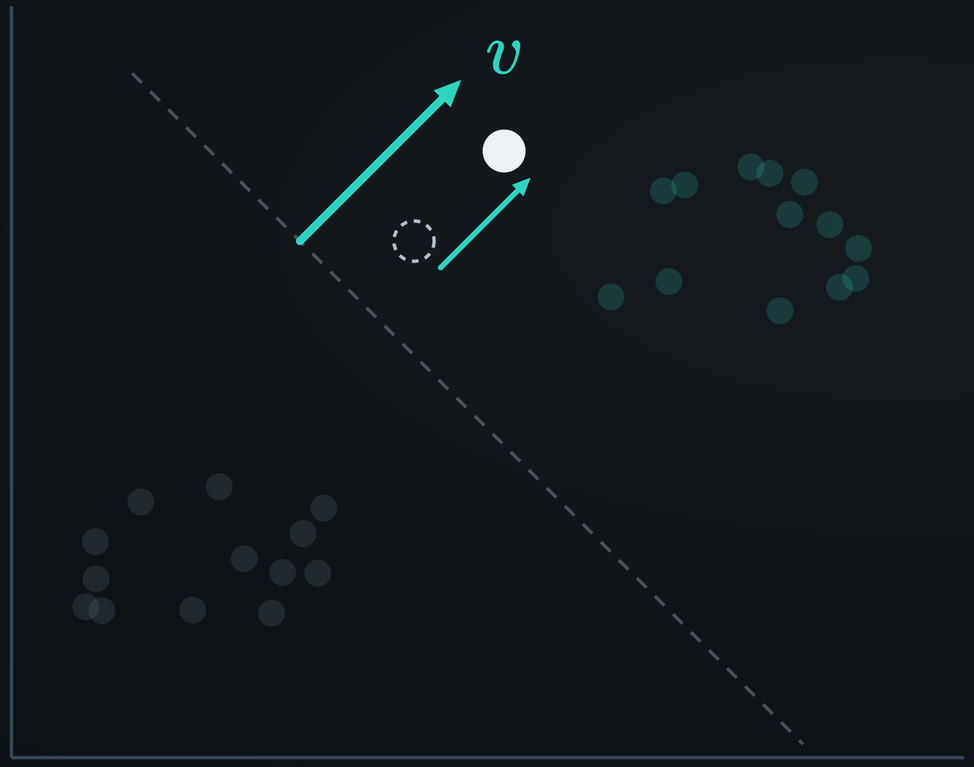

SCHEMATIC · one photo, 2 of 512 dimensions drawn

herring gull score

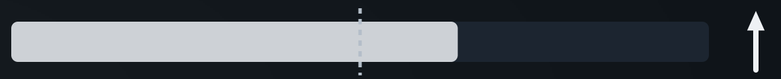

does the **score** go up?

$$\nabla_h \hat{y}_{\text{herring gull}} \cdot v > 0 \; ?$$

Appendix: testing with concept activation vectors

The TCAV score: the share of photos whose score rises

the share of the **60 herring gull photos** whose **score** goes up

$$\mathrm{TCAV} = \frac{\#\{\, \nabla_h \hat{y} \cdot v > 0 \,\}}{60}$$

grey wings: between **0.93** and **1.00** in all 10 repeats

grey wings: one score per repeat

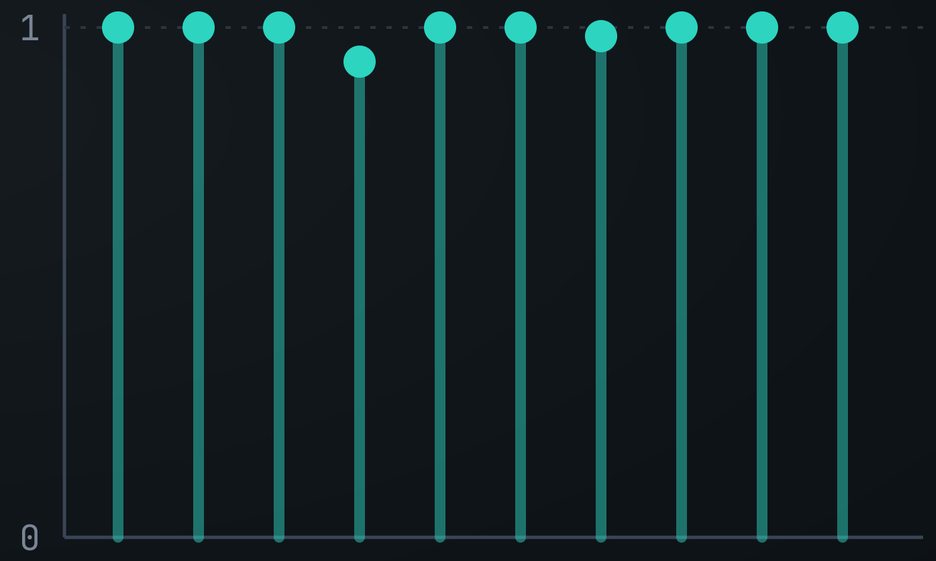

repeats 1 to 10, new example photos each time

COMPUTED · REAL MODEL · CLIP ViT-B/32 zero-shot herring gull score, CUB photos; Kim et al.'s method

Appendix: testing with concept activation vectors

Compare with random directions

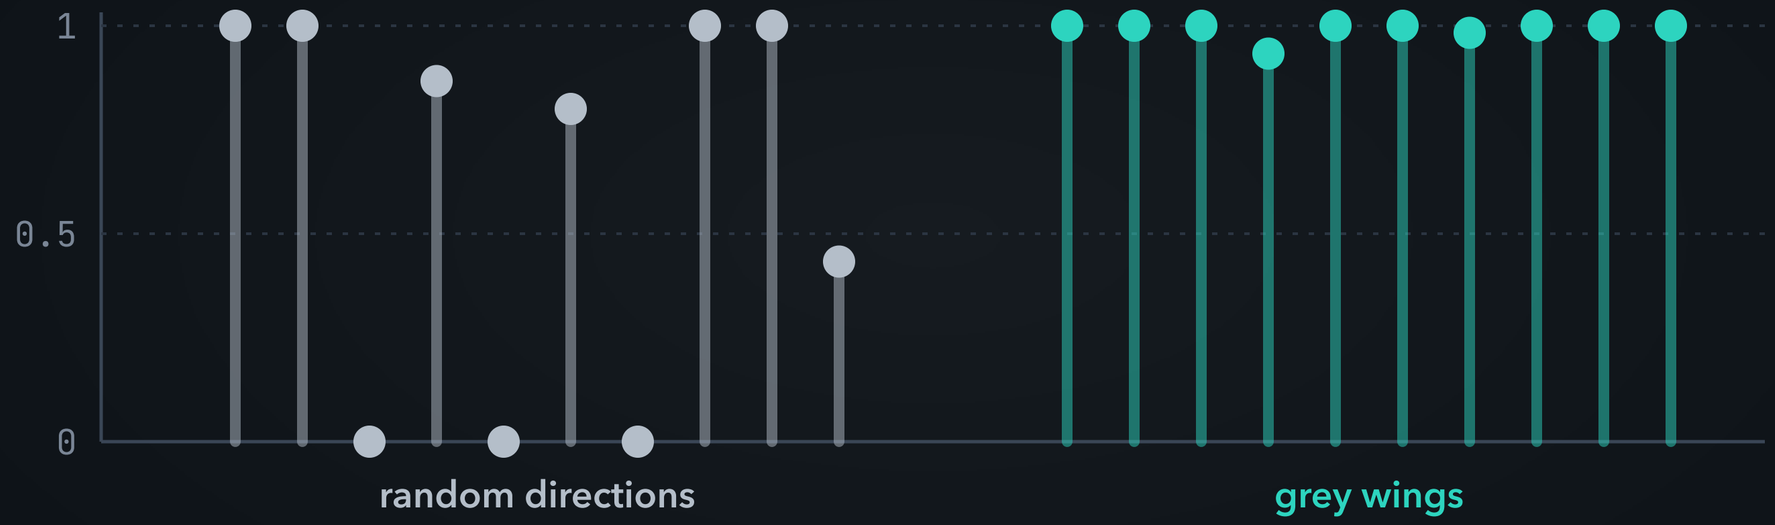

a concept counts only if its score is steady and differs from random

it says what the model responds to, not how to fix it

COMPUTED · REAL MODEL · CLIP ViT-B/32 zero-shot herring gull score, CUB photos; Kim et al.'s method

Concept Bottleneck Models

Predicting through interpretable concepts.

Narration: synthetic voice (Gemini TTS). Photographs: Rhododendrites, Yathin S Krishnappa, Gzen92, Dietmar Rabich, Hans Hillewaert, Stillwaterising, and the authors of the 36 Part 10 photographs, via Wikimedia Commons (CC BY-SA, CC BY, CC0, public domain); knee X-ray from Koh et al. 2020 (CC BY 4.0). Satellite imagery: Sentinel-2, SentinelKilnDB. Data: CUB-200-2011. Models: CLIP ViT-B/32, ResNet-50; concept lists: Gemini.
Sources: Koh et al. 2020; Wah et al. 2011; Simonyan et al. 2014; Kim et al. 2018; Adebayo et al. 2018; Zech et al. 2018; Rudin 2019; Chen et al. 2019; Zarlenga et al. 2022; Yuksekgonul et al. 2023; Menon and Vondrick 2023; Srivastava et al. 2024; Mahinpei et al. 2021; Havasi et al. 2022; Knab et al. 2026.